

How to Tango with Django

Release 1

Leif Azzopardi and David Maxwell

September 24, 2015

1 Overview	1
1.1 Why Work with this Book?	1
1.2 What You will Learn	2
1.3 Technologies and Services	2
1.4 Rango: Initial Design and Specification	3
1.5 Exercises	3
1.6 N-Tier Architecture	4
1.7 Wireframes	4
1.8 Pages and URL Mappings	5
1.9 Entity-Relationship Diagram	6
1.10 Summary	6
2 Getting Ready to Tango	7
2.1 Using the Terminal	7
2.2 Installing the Software	10
2.3 Integrated Development Environment	14
2.4 Exercises	15
3 Django Basics	17
3.1 Testing your Setup	17
3.2 Creating your Django Project	17
3.3 Creating a Django Application	19
3.4 Creating a View	20
3.5 Mapping URLs	21
3.6 Basic Workflows	23
3.7 Exercises	23
4 Templates and Static Media	25
4.1 Using Templates	25
4.2 Serving Static Media	29
4.3 Static Media Files and Templates	30
4.4 The Static Media Server	32
4.5 Basic Workflow	33
4.6 Exercises	34
5 Models and Databases	35
5.1 Rango's Requirements	35
5.2 Telling Django About Your Database	35

5.3	Creating Models	36
5.4	Creating and Synchronising the Database	37
5.5	Generated SQL	38
5.6	Django Models and the Django Shell	39
5.7	Configuring the Admin Interface	39
5.8	Creating a Population Script	41
5.9	Basic Workflows	43
5.10	Exercises	44
6	Models, Templates and Views	47
6.1	Basic Workflow: Data Driven Pages	47
6.2	Showing Categories on Rango's Homepage	47
6.3	Creating a Details Page	49
6.4	Exercises	53
7	Fun with Forms	55
7.1	Basic Workflow	55
7.2	Page and Category Forms	56
7.3	Exercises	61
8	User Authentication	63
8.1	Setting up Authentication	63
8.2	The User Model	64
8.3	Additional User Attributes	64
8.4	Creating a <i>User Registration</i> View and Template	66
8.5	Adding Login Functionality	70
8.6	Restricting Access	73
8.7	Logging Out	75
8.8	Exercises	76
9	Working with Templates	77
9.1	Reoccurring HTML and The Base Template	77
9.2	Template Blocks	78
9.3	Template Inheritance	80
9.4	Exercises	81
10	Cookies and Sessions	83
10.1	Cookies, Cookies Everywhere!	83
10.2	Sessions and the Stateless Protocol	84
10.3	Setting up Sessions in Django	85
10.4	A Cookie Tasting Session	86
10.5	Client Side Cookies: A Site Counter Example	87
10.6	Session Data	88
10.7	Browser-Length and Persistent Sessions	90
10.8	Basic Considerations and Workflow	90
10.9	Exercises	91
11	Bootstrapping Rango	93
11.1	Setting up The Base Template	93
11.2	Quick Style Change	96
11.3	The Login Page	102
11.4	The Registration Template	104
11.5	The End Result	105
12	Adding External Search Functionality	107

12.1	The Bing Search API	107
12.2	Adding Search Functionality	108
12.3	Putting Search into Rango	110
12.4	Exercises	112
13	Making Rango Tango! Exercises	113
13.1	Providing Categories on Every Page	114
13.2	Searching Within a Category Page	114
13.3	View Profile	114
13.4	Track Page Click Throughs	115
14	Doing the Tango with Rango!	117
14.1	List Categories on Each Page	117
14.2	Searching Within a Category Page	118
14.3	View Profile	120
14.4	Track Page Click Throughs	121
15	AJAX, Django and JQuery	123
15.1	Add a “Like Button”	123
15.2	Adding Inline Category Suggestions	126
15.3	Exercises	128
16	Deploying Your Project	131
16.1	Creating a PythonAnywhere Account	131
16.2	The PythonAnywhere Web Interface	131
16.3	Creating a Virtual Environment	132
16.4	Cloning your Git Repository	133
16.5	Setting up your Web Application	134
16.6	Log Files	136
16.7	Exercises	136
17	Summary	137
18	A Git Crash Course	139
18.1	Why Use Version Control?	139
18.2	Git on Windows	139
18.3	The Git System	140
18.4	Setting up Git	140
18.5	Basic Commands and Workflow	142
18.6	Recovering from Mistakes	145
18.7	Exercises	146
19	A CSS Crash Course	149
19.1	Including Stylesheets	149
19.2	Basic CSS Selectors	151
19.3	Element Selectors	151
19.4	Fonts	152
19.5	Colours and Backgrounds	153
19.6	Containers, Block-Level and Inline Elements	155
19.7	Basic Positioning	156
19.8	The Box Model	161
19.9	Styling Lists	161
19.10	Styling Links	162
19.11	The Cascade	164
19.12	Additional Reading	164

20 Javascript and JQuery	167
21 Test Driven Development	169
21.1 Exercises	169
22 Acknowledgements	171
23 Indices and tables	173

Overview

The aim of this book is to provide you with a practical guide to web development using *Django*. The book is designed primarily for students, providing a walkthrough of the steps involved in getting your first web applications up and running, as well as deploying them to a web server.

This book seeks to complement the [official Django Tutorials](#) and many of the other excellent tutorials available online. By putting everything together in one place, this book fills in many of the gaps in the official Django documentation providing an example-based design driven approach to learning the Django framework. Furthermore, this book provides an introduction to many of the aspects required to master web application development.

1.1 Why Work with this Book?

This book will save you time. On many occasions we've seen clever students get stuck, spending hours trying to fight with Django and other aspects of web development. More often than not, the problem was usually because a key piece of information was not provided, or something was not made clear. While the occasional blip might set you back 10-15 minutes, sometimes they can take hours to resolve. We've tried to remove as many of these hurdles as possible. This will mean you can get on with developing your application, and not have to sit there scratching your head.

This book will lower the learning curve. Web application frameworks can save you a lot of hassle and lot of time. Well, that is if you know how to use them in the first place! Often the learning curve is steep. This book tries to get you going - and going fast. By showing you how to put together a web application with all the bells and whistle from the onset, the book shortens the learning curve.

This book will improve your workflow. Using web application frameworks requires you to pick up and run with a particular design pattern - so you only have to fill in certain pieces in certain places. After working with many students, we heard lots of complaints about using web application frameworks - specifically about how they take control away from them (i.e. inversion of control). To help you, we've created a number of workflows to focus your development process so that you can regain that sense of control and build your web application in a disciplined manner.

This book is not designed to be read. Whatever you do, do not read this book! It is a hands-on guide to building web applications in Django. Reading is not doing. To increase the value you gain from this experience, go through and develop the application. When you code up the application, *do not just cut and paste the code*. Type it in, think about what it does, then read the explanations we have provided to describe what is going on. If you still do not understand, then check out the Django documentation, go to [Stack Overflow](#) or other helpful websites and fill in this gap in your knowledge. If you think it is worth mentioning, please get in touch with us so that we can improve the book - we already have a number of contributors and we will happily acknowledge your contribution!

1.2 What You will Learn

In this book, we will be taking an example-based approach (or inquiry-based learning). The book will show you how to design a web application called *Rango* (see the Design Brief in Section 1.4.1 below). Along the way, we'll show you how to perform the following tasks.

- Setup a development environment - including how to use the terminal, the Pip installer, how to work with Git, etc.
- Setup a Django project and create a basic Django application.
- Configure the Django project to serve static media and other media files.
- Work with Django's *Model-View-Template* design pattern.
- Create database models and use the object relational mapping functionality provided by Django.
- Create forms that can utilise your database models to create dynamically generated webpages.
- Use the User Authentication services provided by Django.
- Incorporate external services into the application.
- Include *Cascading Styling Sheets (CSS)* and *JavaScript* within a web application.
- Design and apply CSS to improve the look and feel the web application.
- Work with cookies and sessions with Django.
- Include more advanced functionality like *AJAX* into your application.
- Deploy your application to a web server using *PythonAnywhere*.

At the end of each chapter, we have included a number of exercises designed to push you harder and to see if you can apply what you have learned. The later chapters of the book provide a number of open development exercises along with coded solutions and explanations. Finally, all the code is available from *GitHub* at https://github.com/leifos/tango_with_django.

To see a fully-functional version of the application, you can also visit the *How to Tango with Django* website at <http://www.tangowithdjango.com/rango/>.

1.3 Technologies and Services

Through the course of this book, we will used various technologies and external services, including:

- Python, <http://www.python.org>
- Pip, <http://www.pip-installer.org>
- Django, <https://www.djangoproject.com>
- Git, <http://git-scm.com>
- GitHub, <https://github.com>
- HTML, <http://www.w3.org/html/>
- CSS, <http://www.w3.org/Style/CSS/>
- Javascript
- JQuery, <http://jquery.com>
- Twitter Bootstrap, <http://getbootstrap.com/>

- Bing Search API via Azure Datamarket, <http://datamarket.azure.com>
- PythonAnywhere, <https://www.pythonanywhere.com>

We've selected these technologies and services as they are either fundamental to web development, and/or enable us to provide examples on how to integrate your web application with CSS toolkits (like *Twitter Bootstrap*), external services like (those provided by *Microsoft Azure*) and deploy your application quickly and easily (with PythonAnywhere).

1.4 Rango: Initial Design and Specification

As previously mentioned, the focus of this book will be to develop an application called *Rango*. As we develop this application, it will cover the core components that need to be developed when building any web application.

1.4.1 Design Brief

Your client would like you to create a website called *Rango* that lets users browse through user-defined categories to access various web pages. In Spanish, the word rango is used to mean “*a league ranked by quality*” or “*a position in a social hierarchy*” (see <https://www.vocabulary.com/dictionary/es/rango>).

- **For the main page of the site, they would like visitors to be able to see:**
 - the 5 most viewed pages;
 - the five most rango'ed categories; and
 - some way for visitors to browse or search through categories.
- **When a user views a category page, they would like it to display:**
 - the category name, the number of visits, the number of likes;
 - along with the list of associated pages in that category (showing the page's title and linking to its url); and.
 - some search functionality (via Bing's Search API) to find other pages that can be linked to this category.
- For a particular category, the client would like the name of the category to be recorded, the number of times each category page has been visited, and how many users have clicked a “like” button (i.e. the page gets rango'ed, and voted up the social hierarchy).
- Each category should be accessible via a readable URL - for example, /rango/books-about-django/.
- Only registered users will be able to search and add pages to categories. And so, visitors to the site should be able to register for an account.

At first glance, the application to develop seems reasonably straightforward. In essence, it is just a list of categories which link to pages, right? However, there are a number of complexities and challenges that need to be addressed. First, let's try and build up a better picture of what needs to be developed by laying down some high-level designs.

1.5 Exercises

Before going any further, think about these specifications and draw up the following design artefacts.

- An N-Tier or System Architecture diagram.
- Wireframes of the Main Page and the Category Page.

- The URL Mappings.
- An Entity-Relationship diagram to describe the data model that we'll be implementing.

1.6 N-Tier Architecture

The high-level architecture for most web applications is a *3-Tier architecture*. Rango will be a variant on this architecture as it interfaces with an external service.

Fig. 1.1: Overview of the system architecture for Rango. Note the inclusion of an external Search *Application Programming Interface (API)*.

Since we are building a web application with Django, we will use the following technologies for the following tiers.

- The *client* will be a web browser (i.e Chrome, Firefox, Safari, etc.) which will render HTML/CSS pages.
- The *middleware* will be a Django application, and will be dispatched through Django's built-in development web server while we develop.
- The *database* will be the Python-based *SQLite3* Database engine.
- The *search API* will be the *Bing Search API*.

For the most part, this book will focus on developing the middleware, though it should be quite evident from Figure 1.1 that we will have to interface with all the other components.

1.7 Wireframes

Wireframes are great way to provide clients with some idea of what the application should look like when complete. They save a lot of time, and can vary from hand drawn sketches to exact mockups depending on the tools that you have available. For Rango, we'd like to make the index page of the site look like the screen shot shown in Figure 1.2. Our category page is shown in Figure 1.3.

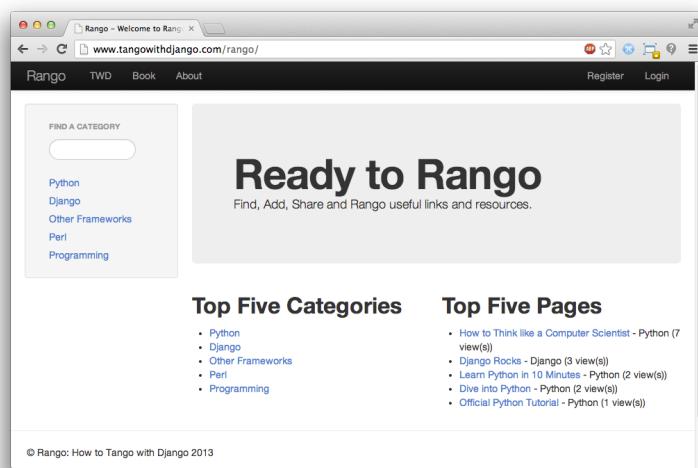


Fig. 1.2: The index page with the categories bar on the left, also showing the top five pages and top five categories.

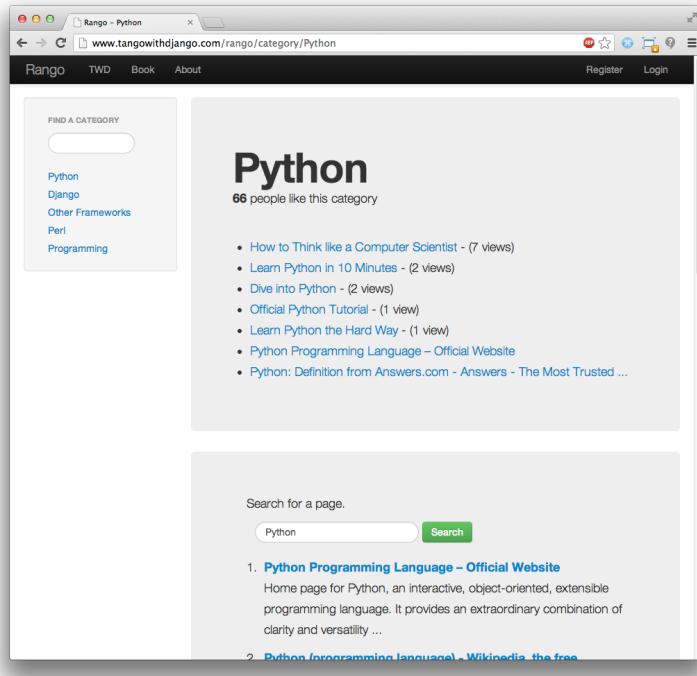


Fig. 1.3: The category page showing the pages in the category (along with the number of views). Below, a search for *Python* has been conducted, with the results shown underneath.

1.8 Pages and URL Mappings

From the specification, we have already identified two pages that our application will present to the user at different points in time. To access each of these pages we will need to describe in some fashion the URL mappings. Think of a URL mapping as the text a user will have to enter into a browser's address bar to reach the given page. The basic URL mappings for Rango are shown below.

- `/rango/` will point to the main (or index) page view.
- `/rango/about/` will point to an about page view.
- `/rango/category/<category_name>/` will point to the category page view for `<category_name>`, where the category name is one of:
 - games;
 - python recipes; or
 - code and compilers.
- `/rango/etc/`, where `etc` could be replaced with a URL for any later function we wish to implement.

As we build our application, we will probably need to create other URL mappings. However, the ones listed above will get us started. We will also at some point have to transform category names in a valid URL string, as well as handle scenarios where the supplied category name does not exist.

As we progress through the book, we will flesh out how to construct these pages using the Django framework and use its Model-View-Template design pattern. However, now that we have a gist of the URL mappings and what the pages are going to look like, we need to define the data model that will house the data for our web application.

1.9 Entity-Relationship Diagram

Given the specification, it should be clear that we have at least two entities: a *category* and a *page*. It should also be clear that a *category* can house many *pages*. We can formulate the following ER Diagram to describe this simple data model.

Fig. 1.4: The Entity Relationship Diagram of Rango's two main entities.

Note that this specification is vague. One page may be in one or many categories. So we could model the relationship as a many-to-many. This approach however introduces a number of complexities, so we will make the simplifying assumption that *one category contains many pages, but one page is assigned to one category*. This does not preclude that the same page can be assigned to different categories - but the page would have to be entered twice, which may not be ideal.

It's good practice to note down any working assumptions like this. You never know when they may come back to haunt you! By noting them down, this means you can communicate it with your development team and make sure that the assumption is sensible and that they are happy to proceed under such an assumption.

The resulting tables are shown below, where `Str` denotes a string or `char` field, `Int` denotes an integer field, `URL` denotes a URL field and `FK` denotes a Foreign Key.

We will also have a `User` table - which we have not shown here, but shall introduce later in the book. In the following chapters will we see how to instantiate these data models in Django and how to use Django's Object Relational Mapping to connect to the database.

1.10 Summary

These high level design and specifications will serve as a useful reference point when building our web application. While we will be focusing on using specific technologies, these steps are common to most database driven web sites. It's a good idea to become familiar and comfortable with producing such specifications and designs.

If you already have Python 2.7.5 and Django 1.5.4 installed, you have a good working knowledge of the command line, configured your paths, then you can skip straight to the [Django Basics](#) chapter. Otherwise, get started with Chapter 2.

1.10.1 Working with The Official Django Tutorials

We suggest undertaking the [Official Django Tutorials](#) as part of the exercises associated with each of this book's chapters. You can find a mapping between the tutorial exercises and book chapters below. The tutorial exercises will help reinforce your understanding of the Django framework, and also help you build up your skills.

Getting Ready to Tango

Let's get set up! To tango with Django, you'll need to ensure that you have everything you need installed on your computer and that you have a sound understanding of your development environment. This chapter walks you through what you need and what you need to know.

For this tutorial, you'll require the following key pieces of software.

- Python version 2.7.5
- Django version 1.5.4

As Django is a web application framework written in the Python programming language, you will be required to have a working knowledge of Python. If you haven't used Python before or you simply wish to brush up on your skills, then we highly recommend that you check out and work through one or more of the following guides.

- **A quick tutorial** - Learn Python in 10 Minutes by Stavros, <http://www.korokithakis.net/tutorials/python/>.
- **The Official Python Tutorial** at <http://docs.python.org/2/tutorial/>.
- **A brilliant book:** Think Python: How to Think like a Computer Scientist by Allen B. Bowney, available online at <http://www.greenteapress.com/thinkpython/>.
- **An amazing online course:** Learn to Program, by Jennifer Campbell and Paul Gries at <https://www.coursera.org/course/programming1>.

2.1 Using the Terminal

In order to set up your environment learning how to use the *Command Line Interpreter (CLI)* provided by your Operating System is really important. Through the course of this tutorial, you will be interacting with the CLI routinely. If you are already familiar with using the command line interface you can skip directly to *Installing the Software* section.

UNIX-based operating systems all use a similar-looking terminal. Descendants, derivatives and clones of UNIX include Apple's OS X and the many available Linux distributions available today. All of these operating systems contain a core set of commands which help you navigate through your filesystem and launch programs, all without the need of any graphical interface. This section provides the key commands you should familiarise yourself with.

Note: This tutorial is focused towards users of UNIX-based or UNIX-derived operating systems. While Python and Django can run in a Windows-based environment, many of the commands that we use in this book are for UNIX-based terminals. These commands can however be replicated in Windows by using the graphical user interface, using the relevant command in a Windows Command Prompt, or using Windows PowerShell which provides an CLI like similar to a UNIX terminal.

How to Tango with Django, Release 1

Upon launching a new terminal instance, you'll typically be presented with something like:

```
sibu:~ leif$
```

This is called the *prompt*, and indicates when the system is waiting to execute your every command. The prompt you see varies depending on the operating system you are using, but all look generally very similar. In the example above, there are three key pieces of information to observe:

- your username and computer name (username of `leif` and computer name of `sibu`);
- your *current working directory* (the tilde, or `~`); and
- the privilege of your user account (the dollar sign, or `$`).

The dollar sign (`$`) typically indicates that the user is a standard user account. Conversely, a hash symbol (#) may be used to signify the user logged in has [root privileges](#). Whatever symbol is present is used to signify that the computer is awaiting your input.

Open up a terminal window and see what your prompt looks like.

When you are using the terminal, it is important to know where you are in the file system. To find out where you are, you can issue the command `pwd`. This will display your present working directory. For example, check the example terminal interactions below.

```
Last login: Mon Sep 23 11:35:44 on ttys003
sibu:~ leif$ pwd
/Users/leif
sibu:~ leif$
```

You can see that the present working directory in this example is: `/Users/leif`.

You'll also note that the prompt indicates that my present working directory is `~`. This is because the tilde (~) represents your *home directory*. The base directory in any UNIX-based file system is the *root directory*. The path of the root directory is denoted by a single forward slash (/).

If you are not in your home directory you can change directory (`cd`) to your home directory by issuing the following command.

```
$ cd ~
```

Let's create a directory called `code`. To do thus, use the make directory command (`mkdir`), as shown below.

```
$ mkdir code
```

To move to the newly-created `code` directory, enter `cd code`. If you now check your current working directory, you'll notice that you will be in `~/code/`. This may also be reflected by your prompt. Note in the example below that the current working directory is printed after the `sibu` computer name.

Note: Whenever we refer to <workspace>, we'll be referring to your `code` directory.

```
sibu:~ leif$ mkdir code
sibu:~ leif$ cd code
sibu:code leif$
sibu:code leif$ pwd
/Users/leif/code
```

To list the files that are in a directory, you can issue the command `ls`. You can also see hidden files or directories - if you have any - you can issue the command `ls -a`, where `a` stands for *all*. If you `cd` back to your home directory (`cd ~`) and then issue `ls`, you'll see that you have something called `code` in your home directory.

To find out a bit more about what is in your directory, issue `ls -l`. This will provide a more detailed *listing* of your files and whether it is a directory or not (denoted by a `d` at the start of the line).

```
sibu:~ leif$ cd ~
sibu:~ leif$ ls -l
drwxr-xr-x 36 leif staff 1224 23 Sep 10:42 code
```

The output also contains information on the [permissions associated to the directory](#), who created it (`leif`), the group (`staff`), the size, the date/time the file was modified at, and, of course, the name.

You may also find it useful to be able to edit files within your terminal. There are many editors which you can use - some of which may already be installed on your computer. The `nano` editor for example is a straightforward editor - unlike `vi` which can take some time to learn. Below are a list of commonly-used UNIX commands that you will find useful.

2.1.1 Core Commands

All UNIX-based operating systems come with a series of built-in commands - with most focusing exclusively on file management. The commands you will use most frequently are listed below, each with a short explanation on what they do and how to use them.

- `pwd`: Prints your current *working directory* to the terminal. The full path of where you are presently is displayed.
- `ls`: Prints a list of files in the current working directory to the terminal. By default, you do not see the sizes of files - this can be achieved by appending `-lh` to `ls`, giving the command `ls -lh`.
- `cd`: In conjunction with a path, allows you to *change* your current working *directory*. For example, the command `cd /home/leif/` changes the current working directory to `/home/leif/`. You can also move up a directory level without having to provide the [absolute path](#) by using two dots, e.g. `cd ...`
- `cp`: Copies files and/or directories. You must provide the *source* and the *target*. For example, to make a copy of the file `input.py` in the same directory, you could issue the command `cp input.py input_backup.py`.
- `mv`: Moves files/directories. Like `cp`, you must provide the *source* and *target*. This command is also used to rename files. For example, to rename `numbers.txt` to `letters.txt`, issue the command `mv numbers.txt letters.txt`. To move a file to a different directory, you would supply either an absolute or relative path as part of the target - like `mv numbers.txt /home/david/numbers.txt`.
- `mkdir`: Creates a directory in your current working directory. You need to supply a name for the new directory after the `mkdir` command. For example, if your current working directory was `/home/david/` and you ran `mkdir music`, you would then have a directory `/home/david/music/`. You will need to then `cd` into the newly created directory to access it.
- `rm`: Shorthand for *remove*, this command removes or deletes files from your filesystem. You must supply the filename(s) you wish to remove. Upon issuing a `rm` command, you will be prompted if you wish to delete the file(s) selected. You can also remove directories [using the recursive switch](#). Be careful with this command - recovering deleted files is very difficult, if not impossible!
- `rmdir`: An alternative command to remove directories from your filesystem. Provide a directory that you wish to remove. Again, be careful: you will not be prompted to confirm your intentions.
- `sudo`: A program which allows you to run commands with the security privileges of another user. Typically, the program is used to run other programs as `root` - the [superuser](#) of any UNIX-based or UNIX-derived operating system.

Note: This is only a brief list of commands. Check out ubuntu's documentation on [Using the Terminal](#) for a more detailed overview, or the [Cheat Sheet](#) by FOSSwire for a quick reference guide.

2.2 Installing the Software

Now that you have a decent understanding of how to interact with the terminal, you can begin to install the software required for this tutorial.

2.2.1 Installing Python

So, how do you go about installing Python 2.7.5 on your computer? You may already have Python installed on your computer - and if you are using a Linux distribution or OS X, you will definitely have it installed. Some of your operating system's functionality [is implemented in Python](#), hence the need for an interpreter!

Unfortunately, nearly all modern operating systems utilise a version of Python that is older than what we require for this tutorial. There's many different ways in which you can install Python, and many of them are sadly rather tricky to accomplish. We demonstrate the most commonly used approaches, and provide links to additional reading for more information.

Warning: This section will detail how to run Python 2.7.5 *alongside* your current Python installation. It is regarded as poor practice to remove your operating system's default Python installation and replace it with a newer version. Doing so could render aspects of your operating system's functionality broken!

Apple OS X

The most simple way to get Python 2.7.5 installed on your Mac is to download and run the simple installer provided on the official Python website. You can download the installer by visiting the webpage at <http://www.python.org/getit/releases/2.7.5/>.

Warning: Ensure that you download the .dmg file that is relevant to your particular OS X installation!

1. Once you have downloaded the .dmg file, double-click it in the Finder.
2. The file mounts as a separate disk and a new Finder window is presented to you.
3. Double-click the file Python.mpkg. This will start the Python installer.
4. Continue through the various screens to the point where you are ready to install the software. You may have to provide your password to confirm that you wish to install the software.
5. Upon completion, close the installer and eject the Python disk. You can now delete the downloaded .dmg file.

You should now have an updated version of Python installed, ready for Django! Easy, huh?

Linux Distributions

Unfortunately, there are many different ways in which you can download, install and run an updated version of Python on your Linux distribution. To make matters worse, methodologies vary from distribution to distribution. For example, the instructions for installing Python on [Fedora](#) may differ from those to install it on an [Ubuntu](#) installation.

However, not all hope is lost. An awesome tool (or a *Python environment manager*) called [pythonbrew](#) can help us address this difficulty. It provides an easy way to install and manage different versions of Python, meaning you can leave your operating system's default Python installation alone. Hurrah!

Taken from the instructions provided from [the pythonbrew GitHub page](#) and [this Stack Overflow question and answer page](#), the following steps will install Python 2.7.5 on your Linux distribution.

1. Open a new terminal instance.

2. Run the command `curl -kL http://xrl.us/pythonbrewinstall | bash`. This will download the installer and run it within your terminal for you. This installs pythonbrew into the directory `~/.pythonbrew`. Remember, the tilde (~) represents your home directory!
3. You then need to edit the file `~/.bashrc`. In a text editor (such as gedit, nano, vi or emacs), add the following to a new line at the end of `~/.bashrc`:
`[[-s $HOME/.pythonbrew/etc/bashrc]] && source $HOME/.pythonbrew/etc/bashrc`
4. Once you have saved the updated `~/.bashrc` file, close your terminal and open a new one. This allows the changes you make to take effect.
5. Run the command `pythonbrew install 2.7.5` to install Python 2.7.5.
6. You then have to *switch* Python 2.7.5 to the *active* Python installation. Do this by running the command `pythonbrew switch 2.7.5`.
7. Python 2.7.5 should now be installed and ready to go.

Note: Directories and files beginning with a period or dot can be considered the equivalent of *hidden files* in Windows. **Dot files** are not normally visible to directory-browsing tools, and are commonly used for configuration files. You can use the `ls` command to view hidden files by adding the `-a` switch to the end of the command, giving the command `ls -a`.

Windows

By default, Microsoft Windows comes with no installations of Python. This means that you do not have to worry about leaving existing versions be; installing from scratch should work just fine. You can download a 64-bit or 32-bit version of Python from [the official Python website](#). If you aren't sure which one to download, you can determine if your computer is 32-bit or 64-bit by looking at the instructions provided [on the Microsoft website](#).

1. When the installer is downloaded, open the file from the location to which you downloaded it.
2. Follow the on-screen prompts to install Python.
3. Close the installer once completed, and delete the downloaded file.

Once the installer is complete, you should have a working version of Python ready to go. By default, Python 2.7.5 is installed to the folder `C:\Python27`. We recommend that you leave the path as it is.

Upon the completion of the installation, open a Command Prompt and enter the command `python`. If you see the Python prompt, installation was successful. However, in certain circumstances, the installer may not set your Windows installation's PATH environment variable correctly. This will result in the `python` command not being found. Under Windows 7, you can rectify this by performing the following:

1. Click the *Start* button, right click *My Computer* and select *Properties*.
2. Click the *Advanced* tab.
3. Click the *Environment Variables* button.
4. In the *System variables* list, find the variable called *Path*, click it, then click the *Edit* button.
5. At the end of the line, enter `;C:\python27;C:\python27\scripts`. Don't forget the semicolon - and certainly *do not* add a space.
6. Click *OK* to save your changes in each window.
7. Close any Command Prompt instances, open a new instance, and try run the `python` command again.

This should get your Python installation fully working. Windows XP, [has slightly different instructions](#), and so do [Windows 8 installations](#).

2.2.2 Setting Up the PYTHONPATH

With Python now installed, we now need to check that the installation was successful. To do this, we need to check that the `PYTHONPATH` environment variable is setup correctly. `PYTHONPATH` provides the Python interpreter with the location of additional Python packages and modules which add extra functionality to the base Python installation. Without a correctly set `PYTHONPATH`, we'll be unable to install and use Django!

First, let's verify that our `PYTHONPATH` variable exists. Depending on the installation technique that you chose, this may or may not have been done for you. To do this on your UNIX-based operating system, issue the following command in a terminal.

```
$ echo $PYTHONPATH
```

On a Windows-based machine, open a Command Prompt and issue the following.

```
$ echo %PYTHONPATH%
```

If all works, you should then see output that looks something similar to the example below. On a Windows-based machine, you will obviously see a Windows path, most likely originating from the C drive.

```
/opt/local/Library/Frameworks/Python.framework/Versions/2.7/lib/python2.7/site-packages:
```

This is the path to your Python installation's `site-packages` directory, where additional Python packages and modules are stored. If you see a path, you can continue to the next part of this tutorial. If you however do not see anything, you'll need to do a little bit of detective work to find out the path. On a Windows installation, this should be a trivial exercise: `site-packages` is located within the `lib` folder of your Python installation directory. For example, if you installed Python to `C:\Python27`, `site-packages` will be at `C:\Python27\Lib\site-packages\`.

UNIX-based operating systems however require a little bit of detective work to discover the path of your `site-packages` installation. To do this, launch the Python interpreter. The following terminal session demonstrates the commands you should issue.

```
$ python

Python 2.7.5 (v2.7.5:ab05e7dd2788, May 13 2013, 13:18:45)
[GCC 4.2.1 (Apple Inc. build 5666) (dot 3)] on darwin
Type "help", "copyright", "credits" or "license" for more information.

>>> import site
>>> print site.getsitepackages()[0]

'/Library/Frameworks/Python.framework/Versions/2.7/lib/python2.7/site-packages'

>>> quit()
```

Calling `site.getsitepackages()` returns a list of paths that point to additional Python package and module stores. The first typically returns the path to your `site-packages` directory - changing the list index position may be required depending on your installation. If you receive an error stating that `getsitepackages()` is not present within the `site` module, verify you're running the correct version of Python. Version 2.7.5 should include this function. Previous versions of the language do not include this function.

The string which is shown as a result of executing `print site.getsitepackages()[0]` is the path to your installation's `site-packages` directory. Taking the path, we now need to add it to your configuration. On a UNIX-based or UNIX-derived operating system, edit your `.bashrc` file once more, adding the following to the bottom of the file.

```
export PYTHONPATH=$PYTHONPATH:<PATH_TO_SITE-PACKAGES>
```

Replace `<PATH_TO_SITE-PACKAGES>` with the path to your `site-packages` directory. Save the file, and quit and reopen any instances of your terminal.

On a Windows-based computer, you must follow the instructions shown in Section 2.2.1 to bring up the environment variables settings dialog. Add a PYTHONPATH variable with the value being set to your site-packages folder, which is typically C:\Python27\Lib\site-packages\.

2.2.3 Using Setuptools and Pip

Installing and setting up your development environment is a really important part of any project. While it is possible to install Python Packages such as Django separately, this can lead to numerous problems and hassles later on. For example, how would you share your setup with another developer? How would you set up the same environment on your new machine? How would you upgrade to the latest version of the package? Using a package manager removes much of the hassle involved in setting up and configuring your environment. It will also ensure that the package you install is the correct for the version of Python you are using, along with installing any other packages that are dependent upon the one you want to install.

In this book, we will be using *Pip*. Pip is a user-friendly wrapper over the *Setuptools* Python package manager. Because Pip depends on Setuptools, we are required to ensure that both are installed on your computer.

To start, we should download Setuptools from the [official Python package website](#). You can download the package in a compressed .tar.gz file. Using your favourite file extracting program, extract the files. They should all appear in a directory called setuptools-1.1.6 - where 1.1.6 represents the Setuptools version number. From a terminal instance, you can then change into the directory and execute the script ez_setup.py as shown below.

```
$ cd setuptools-1.1.6  
$ sudo python ez_setup.py
```

In the example above, we also use sudo to allow the changes to become system-wide. The second command should install Setuptools for you. To verify that the installation was successful, you should be able to see output similar to that shown below.

```
Finished processing dependencies for setuptools==1.1.6
```

Of course, 1.1.6 is substituted with the version of Setuptools you are installing. If this line can be seen, you can move onto installing Pip. This is a trivial process, and can be completed with one simple command. From your terminal instance, enter the following.

```
$ sudo easy_install pip
```

This command should download and install Pip, again with system-wide access. You should see the following output, verifying Pip has been successfully installed.

```
Finished processing dependencies for pip
```

Upon seeing this output, you should be able to launch Pip from your terminal. To do so, just type pip. Instead of an unrecognised command error, you should be presented with a list of commands and switches that Pip accepts. If you see this, you're ready to move on!

Note: With Windows-based computers, follow the same basic process. You won't need to enter the sudo command, however.

2.2.4 Installing Django

Once the Python package manager Pip is successfully installed on your computer, installing Django is easy. Open a Command Prompt or terminal window, and issue the following command.

```
$ pip install -U django==1.5.4
```

If you are using a UNIX-based operating system and receive complaints about insufficient permissions, you will need to run the command with elevated privileges using the `sudo` command. If this is the case, you must then run the following command instead.

```
$ sudo pip install -U django==1.5.4
```

The package manager will download Django and install it in the correct location for you. Upon completion, Django should be successfully installed. Note, if you didn't include the `==1.5.4`, then a different version of Django may be installed.

2.2.5 Installing the Python Imaging Library

During the course of building Rango, we will be uploading and handling images. This means we will need support from the [Python Imaging Library](#). To install this package issue the following command.

```
$ pip install pil
```

Again, use `sudo` if required.

2.2.6 Installing Other Python Packages

It is worth noting that additional Python packages can be easily downloaded using the same manner. [The Python Package Index](#) provides a listing of all the packages available through Pip.

To get a list of the packages installed, you can run the following command.

```
$ pip list
```

2.2.7 Sharing your Package List

You can also get a list of the packages installed in a format that can be shared with other developers. To do this issue the following command.

```
$ pip freeze > requirements.txt
```

If you examine `requirements.txt` using either the command `more`, `less` or `cat`, you will see the same information but in a slightly different format. The `requirements.txt` can then be used to install the same setup by issuing the following command. This is incredibly useful for setting up your environment on another computer, for example.

```
$ pip install -r requirements.txt
```

2.3 Integrated Development Environment

While not absolutely necessary, a good Python-based integrated development environment (IDE) can be very helpful to you during the development process. Several exist, with perhaps JetBrains' *PyCharm* and *PyDev* (a plugin of the *Eclipse IDE*) standing out as popular choices. The [Python Wiki](#) provides an up-to-date list of Python IDEs.

Research which one is right for you, and be aware that some may require you to purchase a licence. Ideally, you'll want to select an IDE that supports integration with Django. PyCharm and PyDev both support Django integration out of the box - though you will have to point the IDE to the version of Python that you are using.

2.4 Exercises

To get comfortable with your environment, try out the following exercises.

- Install Python 2.7.5 and Pip.
- Play around with your CLI and create a directory called `code`, which we use to create our projects in.
- Install the Django and Pil packages.

2.4.1 Virtual Environments

We're almost all set to go! However, before we continue, it's worth pointing out that while this setup is fine to begin with, there are some drawbacks. What if you had another Python application that requires a different version to run? Or you wanted to switch to the new version of Django, but still wanted to maintain your Django 1.5.4 project?

The solution to this is to use [virtual environments](#). Virtual environments allow multiple installations of Python and their relevant packages to exist in harmony, without disrupting one another. This is the generally accepted approach to configuring a Python setup nowadays. We don't go into much detail about them in this chapter because of their complexity, but in the chapter on [Deploying your Application](#) we will go through setting up a virtual environment. If you are really keen, check out a non-magical introduction to [Pip and Virtualenv for Python Beginners](#) by Jamie Matthews.

2.4.2 Code Repository

We should also point out that when you develop code, you should always house your code within a version-controlled repository such as [SVN](#) or [GIT](#). We won't be going through this right now so that we can get stuck into developing an application in Django. We have however provided a [crash course on GIT](#). We highly recommend that you set up a GIT repository for your own projects. Doing so could save you from disaster.

Django Basics

Let's get started with Django! In this chapter, we'll be giving you an overview of the how to get started with Django. You'll be setting up a new project and a new web application. By the end of this chapter, you will have a simple Django-power webpage up and running!

3.1 Testing your Setup

Let's start by checking that your Python and Django installations are installed correctly, and are at the correct version for this tutorial. To do this, open a new terminal instance and issue the following command.

```
$ python -c "import platform; print platform.python_version()"  
2.7.5
```

This starts your Python interpreter and executes the code within the string provided as part of the `-c` switch. You should see the version of your Python installation printed as the output to the process. If the version displayed is anything but `2.7.5`, you will need to go back to Section 2.2 and verify you have completed all the relevant steps for your operating system.

After verifying your Python installation, check your Django installation by issuing the following command.

```
$ python -c "import django; print(django.get_version())"  
1.5.4
```

The command again executes the code within the string provided as part of the `-c` switch. After importing Django, you should see `1.5.4` printed underneath. If you see a different set of numbers or are greeted with a Python `ImportError`, go back to Section 2.2 or consult the [Django Documentation on Installing Django](#) for more information. If you find that you have got a different version of Django, it is possible that you will come across problems at some point. It's definitely worth making sure you have Django 1.5.4 installed so that you're on the same page.

3.2 Creating your Django Project

To create a new Django Project, go to your `code` directory (i.e. your `<workspace>` directory), and issue the following command:

```
$ django-admin.py startproject tango_with_django_project
```

This command will invoke the `django-admin.py` script, which will set up a new Django project called `tango_with_django_project` for you. Typically, we append `_project` to the end of our Django project directories so we know exactly what they contain - but the naming convention is entirely up to you.

How to Tango with Django, Release 1

You'll now notice within your workspace is a directory set to the name of your new project, `tango_with_django_project`. Within this newly created directory, you should see two items:

- another directory with the same name as your project, `tango_with_django_project`; and
- a Python script called `manage.py`.

For the purposes of this tutorial, we call this nested directory the *project configuration directory*. Within this directory, you will find four Python scripts. We will discuss this scripts in detail later on, but for now you should see:

- `__init__.py`, a blank Python script whose presence indicates to the Python interpreter that the directory is a Python package;
- `settings.py`, the place to store all of your Django project's settings;
- `urls.py`, a Python script to store URL patterns for your project; and
- `wsgi.py`, a Python script used to help run your development server and deploy your project to a production environment.

Note: The project configuration directory has been created with new Django projects since version 1.4. Having two directories with the same name may seem quite a bit odd, but the change was made to separate out project-related components from its individual applications.

In the project directory, you will see there is a file called `manage.py`. We will be calling this script time and time again as we develop our project, as it provides you with a series of commands you can run to maintain your Django project. For example, `manage.py` allows you to run the built-in Django development server to test your work and run database commands. You'll be using this script a lot throughout the development cycle.

Note: See the Django documentation for more details about the [Admin](#) and [Manage](#) scripts.

You can try using the `manage.py` script now, by issuing the following command.

```
$ python manage.py runserver
```

Executing this command will instruct Django to initiate its lightweight development server. You should see the output in your terminal similar to the example shown below:

```
$ python manage.py runserver

Validating models...

0 errors found
September 21, 2013 - 06:12:41
Django version 1.5.4, using settings 'tango_with_django_project.settings'
Development server is running at http://127.0.0.1:8000/
Quit the server with CONTROL-C.
```

Now open up your favourite web browser and enter the URL <http://127.0.0.1:8000/>¹. You should see a webpage similar to the one shown in Figure 3.1.

You can stop the development server at anytime by pushing **CTRL + C** in your terminal window. If you wish to run the development server on a different port, or allow users from other machines to access it, you can do so by supplying optional arguments. Consider the following command:

```
$ python manage.py runserver <your_machines_ip_address>:5555
```

Executing this command will force the development server to respond to incoming requests on TCP port 5555. You will need to replace `<your_machines_ip_address>` with your computer's IP address.

¹ This assumes that you are using the IP address 127.0.0.1 and port 8000 when running your Django development web server. If you do not explicitly provide a port to run the development server on, Django defaults to port 8000 for you.

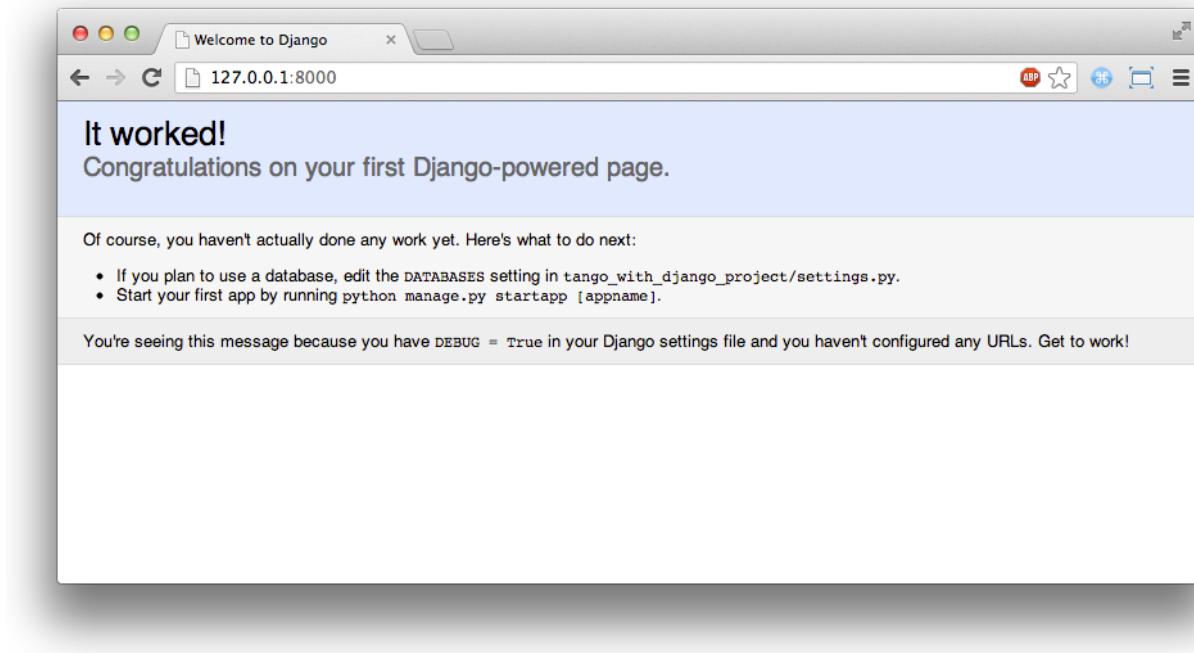


Fig. 3.1: A screenshot of the initial Django page you will see when running the development server for the first time.

When setting ports, it is unlikely that you will be able to use TCP port 80 as this is traditionally reserved for HTTP traffic. Also, any port below 1024 is considered to be [privileged](#) by your operating system.

While you won't be using the lightweight development server to deploy your application, sometimes it is nice to be able to demo your application on a computer of a colleague. Running the server with your machine's IP address will enable others to enter in `http://<your_machines_ip_address>:<port>/` and view your web application. Of course, this will depend on how your network is configured. There may be proxy servers or firewalls in the way which would need to be configured before this would work. Check with the administrator of the network you are using if you can't view the development server remotely.

Note: The `django-admin.py` and `manage.py` scripts provides a lot of useful, time-saving functionality for you. `django-admin.py` allows you to start new projects and apps, along with other commands. Within your project directory, `manage.py` allows you to perform administrative tasks within the scope of your project only. Simply execute the relevant script name without any arguments to see what you can do with each. The [official Django documentation](#) provides a detailed list and explanation of each possible command you can supply for both scripts.

If you are using version control, now may be a good time to commit the changes you have made to your workspace. Refer to the [crash course on GIT](#) if you can't remember the commands and steps involved in doing this.

3.3 Creating a Django Application

A Django project is a collection of *configurations* and *applications* that together make up a given web application or website. One of the intended outcomes of using this approach is to promote good software engineering practices. By developing a small series of applications, the idea is that you can theoretically drop an existing application into a different Django project and have it working with minimal effort. Why reinvent the wheel if it's already there?²

² There are many applications available out there that you can use in your project. Take a look at [PyPI](#) and [Django Packages](#) to search for reusable apps which you can drop into your projects.

A Django application exists to perform a particular task. You need to create specific applications that are responsible for providing your site with particular kinds of functionality. For example, we could imagine that a project might consist of several applications including a polling app, a registration app, and a specific content related app. In another project, we may wish to re-use the polling and registration apps and use them with to dispatch different content. There are many Django applications you can [download](#) and use in your projects. Since we are getting started, we'll kick off by walking through how to create your own application.

To start, create a new application called *Rango*. From within your Django project directory (e.g. <workspace>/tango_with_django_project), run the following command.

```
$ python manage.py startapp rango
```

The `startapp` command creates a new directory within your project's root. Unsurprisingly, this directory is called `rango` - and contained within it are another four Python scripts:

- another `__init__.py`, serving the exact same purpose as discussed previously;
- `models.py`, a place to store your application's data models - where you specify the entities and relationships between data;
- `tests.py`, where you can store a series of functions to test your application's code; and
- `views.py`, where you can store a series of functions that take a client's requests and return responses.

`views.py` and `models.py` are the two files you will use for any given application, and form part of the main architectural design pattern employed by Django, i.e. the *Model-View-Template* pattern. You can check out [the official Django documentation](#) to see how models, views and templates relate to each other in more detail.

Before you can get started with creating your own models and views, you must first tell your Django project about your new application's existence. To do this, you need to modify the `settings.py` file, contained within your project's configuration directory. Open the file and find the `INSTALLED_APPS` tuple. Add the `rango` application to the end of the tuple, which should then look like the following example.

```
INSTALLED_APPS = (
    'django.contrib.auth',
    'django.contrib.contenttypes',
    'django.contrib.sessions',
    'django.contrib.sites',
    'django.contrib.messages',
    'django.contrib.staticfiles',
    # Uncomment the next line to enable the admin:
    # 'django.contrib.admin',
    # Uncomment the next line to enable admin documentation:
    # 'django.contrib.admindocs',
    'rango',
)
```

Verify that Django picked up your new application by running the development server again. If you can start the server without errors, your application was picked up and you will be ready to proceed to the next step.

3.4 Creating a View

With our Rango application created, let's now create a simple view. For our first view, let's just send some simple text back to the client - we won't concern ourselves about using models or templates just yet.

In your favourite IDE, open the file `views.py`, located within your newly created `rango` application directory. Remove the comment `# Create your views here.` so that you now have a blank file.

You can now add in the following code.

```
from django.http import HttpResponse

def index(request):
    return HttpResponse("Rango says hello world!")
```

Breaking down the three lines of code, we observe the following points about creating this simple view.

- We first import the `HttpResponse` object from the `django.http` module.
- Each view exists within the `views.py` file as a series of individual functions. In this instance, we only created one view - called `index`.
- Each view takes in at least one argument - a `HttpRequest` object, which also lives in the `django.http` module. Convention dictates that this is named `request`, but you can rename this to whatever you want if you so desire.
- Each view must return a `HttpResponse` object. A simple `HttpResponse` object takes a string parameter representing the content of the page we wish to send to the client requesting the view.

With the view created, you're only part of the way to allowing a user to access it. For a user to see your view, you must map a [Uniform Resources Locator \(URL\)](#) to the view.

3.5 Mapping URLs

Within the `rango` application directory, we now need to create a new file called `urls.py`. The contents of the file will allow you to map URLs for your application (e.g. `http://www.tangowithdjango.com/rango/`) to specific views. Check out the simple `urls.py` file below.

```
from django.conf.urls import patterns, url
from rango import views

urlpatterns = patterns('',
    url(r'^$', views.index, name='index'))
```

This code imports the relevant Django machinery that we use to create URL mappings. Importing the `views` module from `rango` also provides us with access to our simple view implemented previously, allowing us to reference the view in the URL mapping we will create.

To create our mappings, we use a `tuple`. For Django to pick your mappings up, this tuple *must* be called `urlpatterns`. The `urlpatterns` tuple contains a series of calls to the `django.conf.urls.url()` function, with each call handling a unique mapping. In the code example above, we only use `url()` once, so we have therefore defined only one URL mapping. The first parameter we provide to the `django.conf.urls.url()` function is the regular expression `^$`, which matches to an empty string. Any URL supplied by the user that matches this pattern means that the view `views.index()` would be invoked by Django. The view would be passed a `HttpRequest` object as a parameter, containing information about the user's request to the server. We also make use of the optional parameter to the `url()` function, `name`, using the string '`index`' as the associated value.

Note: You might be thinking that matching a blank URL is pretty pointless - what use would it serve? When the URL pattern matching takes place, only a portion of the original URL string is considered. This is because our Django project will first process the original URL string (i.e. `http://www.tangowithdjango.com/rango/`). Once this has been processed, it is removed, with the remained being passed for pattern matching. In this instance, there would be nothing left - so an empty string would match!

Note: The `name` parameter is optional to the `django.conf.urls.url()` function. This is provided by Django to allow you to distinguish one mapping from another. It is entirely plausible that two separate URL mappings expressions could end calling the same view. `name` allows you to differentiate between them - something which is useful for *reverse URL matching*. Check out the [Official Django documentation on this topic](#) for more information.

You may have seen that within your project configuration directory a `urls.py` file already exists. Why make another? Technically, you can put *all* the URLs for your project's applications within this file. However, this is considered bad practice as it increases coupling on your individual applications. A separate `urls.py` file for each application allows you to set URLs for individual applications. With minimal coupling, you can then join them up to your project's master `urls.py` file later.

This means we need to configure the `urls.py` of our project `tango_with_django_project` and connect up our main project with our Rango application.

How do we do this? It's quite simple. Open the project's `urls.py` file which is located inside your project configuration directory. As a relative path from your workspace directory, this would be the file `<workspace>/tango_with_django_project/tango_with_django_project/urls.py`. Update the `urlpatterns` tuple as shown in the example below.

```
urlpatterns = patterns('',
    # Examples:
    # url(r'^$', 'tango_with_django_project.views.home', name='home'),
    # url(r'^tango_with_django_project/', include('tango_with_django_project.foo.urls')),

    # Uncomment the admin/doc line below to enable admin documentation:
    # url(r'^admin/doc/', include('django.contrib.admindocs.urls')),

    # Uncomment the next line to enable the admin:
    # url(r'^admin/', include(admin.site.urls)),

    url(r'^rango/', include('rango.urls')), # ADD THIS NEW TUPLE!
)
```

The added mapping looks for url strings that match the patterns `^rango/`. When a match is made the remainder of the url string is then passed onto and handled by `rango.urls` (which we have already configured). This is done with the help of the `include()` function from within `django.conf.urls`. Think of this as a chain that processes the URL string - as illustrated in Figure 3.2. In this chain, the domain is stripped out and the remainder of the url string (`rango/`) is passed on to `tango_with_django` project, where it finds a match and strips away `rango/` leaving an empty string to be passed on to the application `rango`. `Rango` now tries to match the empty string, which it does, and this then dispatches the `index()` view that we created.

Kick off the development server and visit `http://127.0.0.1:8000/rango`. If all went well you should see the text `Rango says hello world!`, just like the screenshot shown in Figure 3.3.

Fig. 3.2: An illustration of a URL, showing how the different parts of the URL are the responsibility of different `url.py` files.

Within each application, you will create a number of URL to view mappings. This initial mapping is quite simple. As we progress, we will create more sophisticated mappings that allow the URLs to be parameterised.

It's important to have a good understanding of how URLs are handled in Django. If you are still bit confused or would like to know more check out the [official Django documentation on URLs](#) for further details and further examples.

Note: The URL patterns use [regular expressions](#) to perform the matching. It is worthwhile familiarising yourself on how to use regular expressions in Python. The official Python documentation contains a [useful guide on regular expressions](#), while [regexcheatsheet.com](#) provides a [neat summary of regular expressions](#).

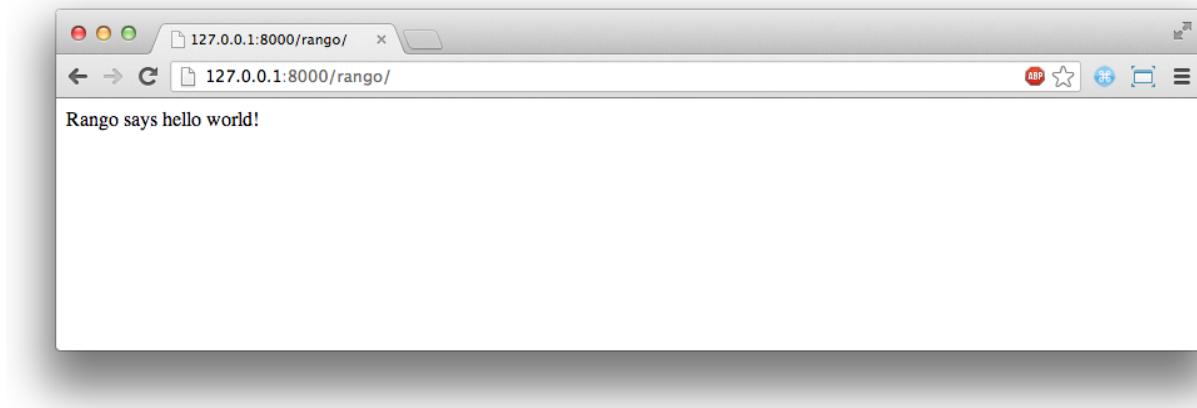


Fig. 3.3: A screenshot of Google Chrome displaying our first Django-powered webpage. Hello, Rango!

3.6 Basic Workflows

What you've just learnt in this chapter can be succinctly summarised into a list of actions. Here, we provide these lists for the two distinct tasks you have performed. You can use this section for a quick reference if you need to remind yourself about particular actions.

3.6.1 Creating a new Django Project

1. To create the project run, `django-admin.py startproject <name>`, where `<name>` is the name of the project you wish to create.

3.6.2 Creating a new Django application

1. To create a new application run, `$ python manage.py startapp <appname>`, where `<appname>` is the name of the application you wish to create.
2. Tell your Django project about the new application by adding it to the `INSTALLED_APPS` tuple in your project's `settings.py` file.
3. In your project `urls.py` file, add a mapping to the application.
4. In your application's directory, create a `urls.py` file to direct incoming URL strings to views.
5. In your application's `view.py`, create the required views ensuring that they return a `HttpResponse` object.

3.7 Exercises

Congratulations! You have got Rango up and running. This is a significant landmark in working with Django. Creating views and mapping URLs to views is the first step towards developing more complex and usable web applications. Now try the following exercises to reinforce what you've learnt.

- Revise the procedure and make sure you follow how the URLs are mapped to views.
- Next create a new view called about page - which returns the following: "Rango Says: Here is the about page."

- Now map the this view to /rango/about/. Note you will only have to edit the `urls.py` of the rango application.
- Revise the `HttpResponse` in the index view to include a link to the about page.
- In the `HttpResponse` in the about view include a link back to the main page.
- If you haven't done so already, it is a good point to go off an complete part one of the official Django Tutorial.

3.7.1 Hints

If you're struggling to get the exercises done, the following hints will hopefully provide you with some inspiration on how to progress.

- The `HttpResponse` in the index view should be something like: "Rango says: Hello world! About"
- The regular expression to match `about/` is `r'^about/'`
- The HTML to link back to the index page is `Index`

Templates and Static Media

In this chapter, we'll be extending your knowledge of Django by introducing you to the template engine as well as how to serve *static media* within your web pages.

4.1 Using Templates

Up to this point, you have plugged a few things together to create a Django-powered webpage. This is coupled a view, which is in turn coupled with a series of URL mappings. Here we will delve into how to combine templates into the mix.

Well-designed websites use a lot of repetition in their structure or layout. Whether you see a common header or footer on a website's pages, the [repetition of page layouts](#) aids users with navigation, promotes organisation of the website and reinforces a sense of continuity. Django provides [templates](#) to make it easier for developers to achieve this design goal, as well as separating application logic from presentational concerns. In this chapter, you'll create a basic template which will be used to create a HTML page. This template will then be dispatched via a Django view. In Chapter 6, we will take this a step further by using templates in conjunction with models to dispatch dynamically generated data.

4.1.1 Configuring the Templates Directory

To get templates up and running, you will need to setup a directory in which template files are stored.

In your Django project's directory (e.g. `<workspace>/tango_with_django_project/`), create a new directory called `templates`. Within the new templates directory, create another directory called `rango`. So the directory `<workspace>/tango_with_django_project/templates/rango/` will be the location in which we will be storing templates associated with our `rango` application.

To tell your Django project where the templates will be housed, open your project's `settings.py` file. Find the tuple `TEMPLATE_DIRS` and add in the path to your newly created `templates` directory, so it looks like the following example.

```
TEMPLATE_DIRS = (
    # Put strings here, like "/home/html/django_templates" or "C:/www/django/templates".
    # Always use forward slashes, even on Windows.
    # Don't forget to use absolute paths, not relative paths.
    '<workspace>/tango_with_django_project/templates/',
)
```

Note that you are *required to use absolute paths* to locate the `templates` directory. If you are part of a team or working on different computers, this may become a problem in the future. You'll have different usernames, meaning different paths to your `<workspace>` directory. The *hard-coded* path you entered above would not be the same on

different computers. Of course, you could add in the template directory for each different setup, but that would be a pretty nasty way to tackle the problem. So, what can we do?

Warning: The road to hell is paved with hard-coded paths. Hard-coding paths is considered to be a [software engineering anti-pattern](#), and will make your project less [portable](#).

4.1.2 Dynamic Paths

The solution to the problem of hard-coding paths is to make use of built-in Python functions to work out the path of our `templates` directory automatically for us. This way, an absolute path can be obtained regardless of where you place your Django project's code on your filesystem. This in turn means that your project's code becomes more *portable*. At the expense of a little bit more complexity in setting up your Django project now, you can make your life much easier later on. No pain, no gain!

To start, modify the `settings.py` file to include a variable called `SETTINGS_DIR`. This will store the path to the directory in which your project's `settings.py` module will be contained. This is obtained by using the special Python `__file__` attribute, which is [set to the absolute path of your settings module](#). We then take the absolute path to the `settings` file, extracting the path to the directory in which the `settings` module is contained. For example, calling `os.path.dirname()` on the absolute path of `<workspace>/tango_with_django_project/tango_with_django_project/settings.py` would yield a directory of `<workspace>/tango_with_django_project/tango_with_django_project/`.

To implement this functionality, add the following code to the top of `settings.py`. Note the importing of the `os` module.

```
import os
SETTINGS_DIR = os.path.dirname(__file__)
```

Now that we have the absolute path to our project's configuration directory, we need to get the root directory of our project. This is exactly one level up from the project configuration directory. We can therefore make use of the `pardir` string within the Python `os` module to help us obtain the path to the [parent directory](#). A project configuration directory path of `<workspace>/tango_with_django_project/tango_with_django_project/` would for example return the path `<workspace>/tango_with_django_project/`.

Add the following two lines to `settings.py` directly underneath what you entered above.

```
PROJECT_PATH = os.path.join(SETTINGS_DIR, os.pardir)
PROJECT_PATH = os.path.abspath(PROJECT_PATH)
```

First, we use the `os.path.join()` function to join together our `SETTINGS_DIR` variable and the `os.pardir` string. Typically, this will return `<workspace>/tango_with_django_project/tango_with_django_project/...`. Calling `os.path.abspath()` on that path will then yield the absolute path to the next directory up the hierarchy, which would be `<workspace>/tango_with_django_project/` in our example. In other words, `PROJECT_PATH` now contains the absolute path to our Django project's root. We can then use this path to point to special directories in our project, such as the directory which stores our templates!

Warning: When joining or concatenating system paths together, using `os.path.join()` is the preferred approach. Using this function ensures that the correct slashes are used depending on your operating system. On a POSIX-compatible operating system, forward slashes would be used to separate directories, whereas a Windows operating system would use backward slashes. If you manually append slashes to paths, you may end up with path errors when attempting to run your code on a different operating system.

Let's make use of it now. Create a new variable in `settings.py` called `TEMPLATE_PATH` and store the path to the `templates` directory you created earlier. Using the `os.path.join()` function, your code should look like the following example.

```
TEMPLATE_PATH = os.path.join(PROJECT_PATH, 'templates')
```

We again make use of `os.path.join()` to mash together the `PROJECT_PATH` variable and '`templates`', which would for example yield `<workspace>/tango_with_django_project/templates/`. We can then replace the hard-coded path we put in the `TEMPLATE_DIRS` tuple earlier with `TEMPLATE_PATH`, just like in the example below.

```
TEMPLATE_DIRS = (
    # Put strings here, like "/home/html/django_templates" or "C:/www/django/templates".
    # Always use forward slashes, even on Windows.
    # Don't forget to use absolute paths, not relative paths.
    TEMPLATE_PATH,
)
```

We can keep the `TEMPLATE_PATH` variable at the top of our `settings.py` module to make it easy to access should it ever need to be changed. This is why we created an additional variable to store the template path.

4.1.3 Adding a Template

With your template directory and path set up, create a file called `index.html` and place it in the `templates/rango/` directory. Within this new file, add the following HTML code:

```
<!DOCTYPE html>
<html>

    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Rango says...</h1>
        hello world! <strong>{{ boldmessage }}</strong><br />
        <a href="/rango/about/">About</a><br />
    </body>

</html>
```

From this HTML code, it should be clear that a simple HTML page is going to be generated that greets a user with a *hello world* message. You might also notice some non-HTML in the form of `{{ boldmessage }}`. This is a *Django template variable*, and we will be able to set a value for this variable to be displayed within the rendered output. We'll get to that in a moment.

To use this template, we need to re-configure the `index()` view that we created earlier. Instead of dispatching a simple message, we will change it to dispatch our template.

In `rango/views.py`, add the following import statements at the top of the file.

```
from django.template import RequestContext
from django.shortcuts import render_to_response
```

You can then update the `index()` view function as follows. Check out the inline commentary to see what each line does.

```
def index(request):
    # Request the context of the request.
    # The context contains information such as the client's machine details, for example.
    context = RequestContext(request)
```

```
# Construct a dictionary to pass to the template engine as its context.
# Note the key boldmessage is the same as {{ boldmessage }} in the template!
context_dict = {'boldmessage': "I am bold font from the context"}

# Return a rendered response to send to the client.
# We make use of the shortcut function to make our lives easier.
# Note that the first parameter is the template we wish to use.
return render_to_response('rango/index.html', context_dict, context)
```

In our updated view, we use the `RequestContext` class to gain access to settings related to the user's request. We then create a dictionary to store any data we want to send through to the template, then finally call the `render_to_response()` helper function. We pass as parameters the template we wish to use, the dictionary with our template variables, and the context we obtained from the user's request. The `render_to_response()` function will take this data and mash it together with the template to produce a complete HTML page. This is then returned and dispatched to the user's web browser.

When a template file is loaded with the Django templating system, a *template context* is created. In simple terms, a template context is essentially a Python dictionary that maps template variable names with Python variables. In the template we created earlier, we included a template variable name called `boldmessage`. In our `index(request)` view example, the string `I am bold font from the context` is mapped to template variable `boldmessage`. The string `I am bold font from the context` therefore replaces any instance of `{{ boldmessage }}` within the template.

Now that you have updated the view to employ the use of your template, run the Django development server and visit <http://127.0.0.1:8000/rango/>. You should see your template rendered in all its glory, just like the example shown in Figure 4.1.

If you don't, read the error message presented to see what the problem is, and then double check all the changes that you have made. Ensure that all the changes required have been made. One of the most common issues people have with templates is that the path is set incorrectly in `settings.py`. Sometimes it's worth adding a `print` statement to `settings.py` to report the `PROJECT_PATH` and `TEMPLATE_PATH`.

This example demonstrates how to use templates within your views. However, we have only touched upon some of the functionality provided by Django regarding templates. We will use templates in more sophisticated ways as we progress through this tutorial. In the meantime, you can find out more about [templates](#) from the official Django documentation.

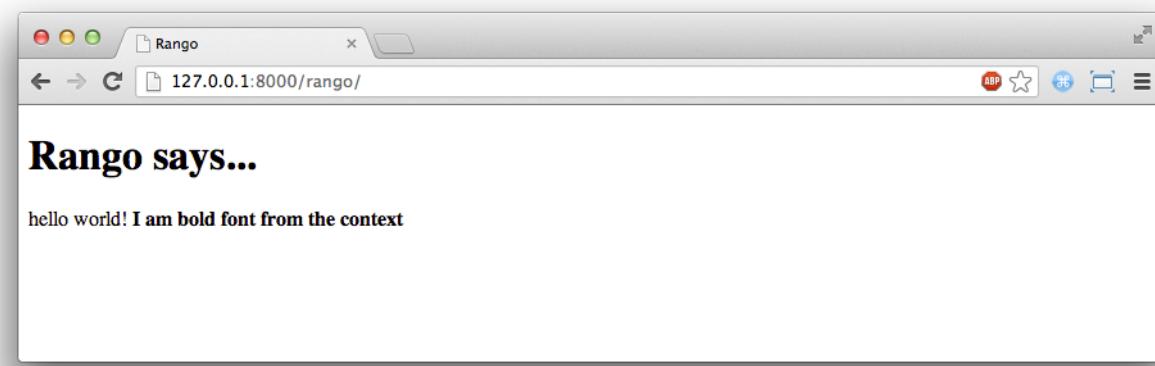


Fig. 4.1: A screenshot of Google Chrome rendering the template used with this tutorial.

4.2 Serving Static Media

Admittedly, the *Rango* website is pretty plain as we have not included any styling or imagery. [Cascading Style Sheets \(CSS\)](#), [JavaScript](#) and images are essentially *static media* files which we can include in our webpages to add style and introduce dynamic behaviour. These files are served in a slightly different way from webpages. This is because they aren't generated on the fly like our HTML pages. This section shows you how to setup your Django project to serve static media to the client. We'll also modify our template to include some example static media.

4.2.1 Configuring the Static Media Directory

To get static media up and running, you will need to set up a directory in which static media files are stored. In your project directory (e.g. <workspace>/tango_with_django_project/), create a new directory called `static`.

Now place an image within the `static` directory. As shown in Figure 4.2, we chose a picture of the chameleon, *Rango* - a fitting mascot, if ever there was one.

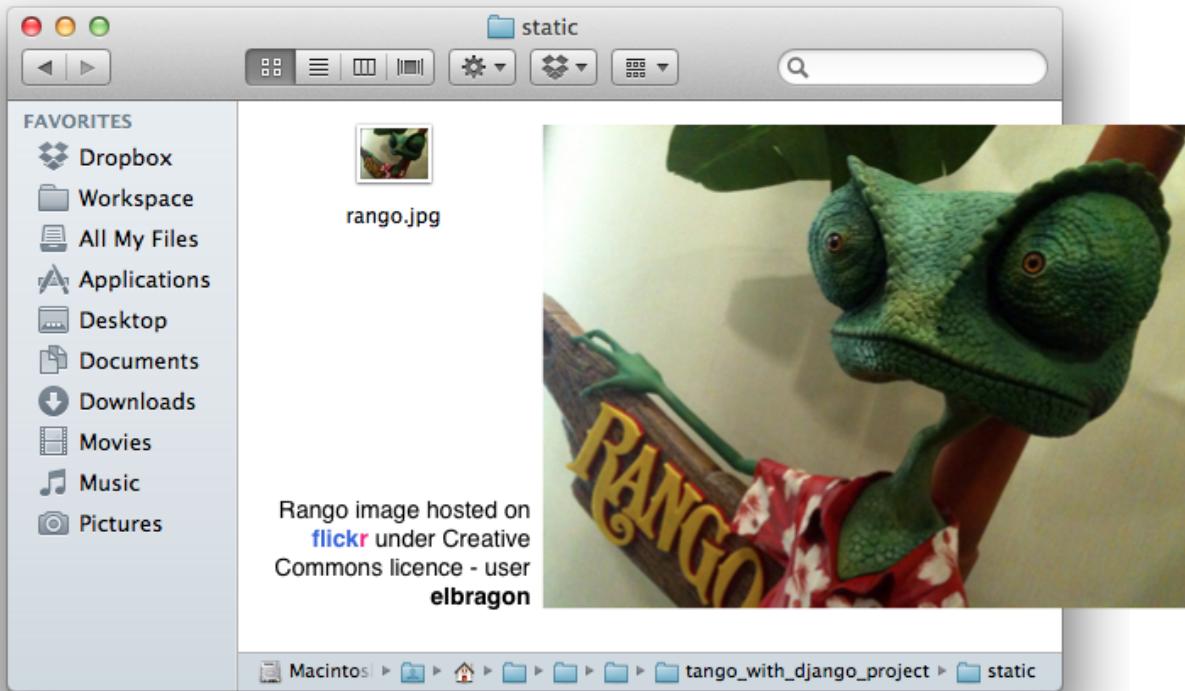


Fig. 4.2: Rango the chameleon within our static media directory.

With our `static` directory created, we need to tell Django about it, just like we did with our `templates` directory earlier. In `settings.py` file, we need to update two variables: `STATIC_URL` and the `STATICFILES_DIRS` tuple. First, create a variable to store the path to the static directory (`STATIC_PATH`) as follows.

```
STATIC_PATH = os.path.join(PROJECT_PATH, 'static')

STATIC_URL = '/static/' # You may find this is already defined as such.
```

```
STATICFILES_DIRS = (
    STATIC_PATH,
)
```

You've typed in some code, but what does it represent? The first variable STATIC_URL defines the base URL with which your Django applications will find static media files when the server is running. For example, when running the Django development server with STATIC_URL set to /static/ like in the code example above, static media will be available at `http://127.0.0.1:8000/static/`. The [official documentation on serving up static media](#) warns that it is vitally important to make sure that those slashes are there. Not configuring this problem can lead to a world of pain.

While STATIC_URL defines the URL to access media via the web server, STATICFILES_DIRS allows you to specify the location of the newly created `static` directory on your local disk. Just like the TEMPLATE_DIRS tuple, STATICFILES_DIRS requires an absolute path to the `static` directory. Here, we re-used the PROJECT_PATH defined in Section 4.1 to create the STATIC_PATH.

With those two settings updated, run your Django project's development server once more. If we want to view our image of Rango, visit the URL `http://127.0.0.1:8000/static/rango.jpg`. If it doesn't appear, you will want to check to see if everything has been correctly spelt and that you saved your `settings.py` file, and restart the development server. If it does appear, try putting in additional file types into the `static` directory and request them via your browser.

Caution: While using the Django development server to serve your static media files is fine for a development environment, it's highly unsuitable for a production - or *live* - environment. The [official Django documentation on Deployment](#) provides further information about deploying static files in a production environment.

4.3 Static Media Files and Templates

Now that you have your Django project set up to handle static media, you can now access such media within your templates.

To demonstrate how to include static media, open up `index.html` located in the <workspace>/templates/rango/ directory. Modify the HTML source code as follows. The two lines that we add are shown with a HTML comment next to them for easy identification.

```
<!DOCTYPE html>

{# load static #} <!-- New line -->

<html>

    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Rango says...</h1>
        hello world! <strong>{{ boldmessage }}</strong><br />
        <a href="/rango/about/">About</a><br />
         <!-- New line -->
    </body>

</html>
```

First, we need to inform Django's template system that we will be using static media with the `{% load static %}` tag. This allows us to call the `static` template tag as done in `{% static "rango.jpg" %}`. As you can see, Django template tags are denoted by curly brackets `{ }`. In this example, the `static` tag will combine the `STATIC_URL` with `"rango.jpg"` so that the rendered HTML looks like the following.

```
 <!-- New line -->
```

If for some reason the image cannot be loaded, it is always nice to specify an alternative text tagline. This is what the `alt` attribute provides - the text here is used in the event the image fails to load.

With these minor changes in place, kick off the Django development server once more and visit `http://127.0.0.1:8000/rango`. Hopefully, you will see web page something like the one shown in Figure 4.3.

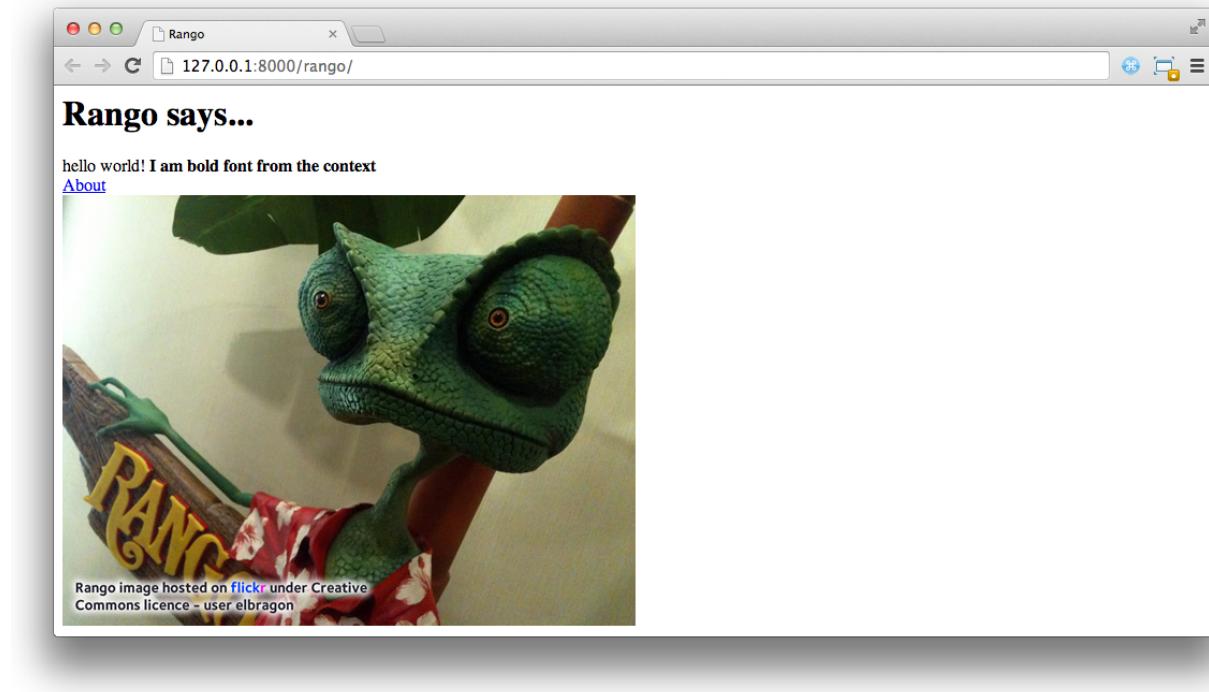


Fig. 4.3: Our first Rango template, complete with a picture of Rango the chameleon.

The `{% static %}` function call should be used whenever you wish to reference static media within a template. The code example below demonstrates how you could include JavaScript, CSS and images into your templates - all with the correct HTML markup.

```
<!DOCTYPE html>

{% load static %}

<html>

    <head>
        <title>Rango</title>
        <link rel="stylesheet" href="{% static "css/base.css" %}" /> <!-- CSS -->
        <script src="{% static "js/jquery.js" %}"></script> <!-- JavaScript -->
    </head>

    <body>
```

```
<h1>Including Static Media</h1>
 <!-- Images -->
</body>

</html>
```

Static files you reference will obviously need to be present within your `static` directory. If the file is not there or you have referenced it incorrectly, the console output provided by Django's lightweight development server will flag up any errors. Try referencing a non-existent file and see what happens.

For further information about including static media you can read through the official [Django documentation on working with static files in templates](#).

Caution: Care should be taken in your templates to ensure that any [document type declaration](#) (e.g. `<!DOCTYPE html>`) you use in your webpages appears in the rendered output on the *first line*. This is why we put the Django template command `{% load static %}` on a line underneath the document type declaration, rather than at the very top. It is a requirement of HTML/XHTML variations that the document type declaration be declared on the very first line. Django commands placed before will obviously be removed in the final rendered output, but they may leave behind residual whitespace which means your output [will fail validation on the W3C markup validation service](#).

4.4 The Static Media Server

Now that you can dispatch static files, let's look at uploading media. Many websites provide their users with the ability to do this - for example, to upload a profile image. This section shows you how to add a simple development media server to your Django project. The development media server can be used in conjunction with file uploading forms which we will touch upon in Chapter 8.

So, how do we go about setting up a development media server? The first step is to create another new directory called `media` within our Django project's root (e.g. `<workspace>/tango_with_django_project/`). The new `media` directory should now be sitting alongside your `templates` and `static` directories. After you create the directory, you must then modify your Django project's `urls.py` file, located in the project configuration directory (e.g. `<workspace>/tango_with_django_project/tango_with_django_project/`). Add the following code to the `urls.py` file.

```
# At the top of your urls.py file, add the following line:
from django.conf import settings

# UNDERNEATH your urlpatterns definition, add the following two lines:
if settings.DEBUG:
    urlpatterns += patterns(
        'django.views.static',
        (r'media/(?P<path>.*)',
        'serve',
        {'document_root': settings.MEDIA_ROOT}), )
```

The `settings` module from `django.conf` allows us access to the variables defined within our project's `settings.py` file. The conditional statement then checks if the Django project is being run in `DEBUG` mode. If the project's `DEBUG` setting is set to `True`, then an additional URL matching pattern is appended to the `urlpatterns` tuple. The pattern states that for any file requested with a URL starting with `media/`, the request will be passed to the `django.views.static` view. This view handles the dispatching of uploaded media files for you.

With your `urls.py` file updated, we now need to modify our project's `settings.py` file. We now need to set the values of two variables. In your file, find `MEDIA_URL` and `MEDIA_ROOT`, setting them to the values as shown below.

```
MEDIA_URL = '/media/'
MEDIA_ROOT = os.path.join(PROJECT_PATH, 'media') # Absolute path to the media directory
```

The first variable `MEDIA_URL` defines the base URL from which all media files will be accessible on your development server. Setting the `MEDIA_URL` for example to `/media/` will mean that user uploaded files will be available from the URL `http://127.0.0.1:8000/media/`. `MEDIA_ROOT` is used to tell Django where uploaded files should be stored on your local disk. In the example above, we set this variable to the result of joining our `PROJECT_PATH` variable defined in Section 4.1 with `/media/`. This gives an absolute path of `<workspace>/tango_with_django_project/media/`.

Caution: As previously mentioned, the development media server supplied with Django is very useful for debugging purposes. However, it should **not** be used in a production environment. The official [Django documentation on static files](#) warns that such an approach is “*grossly inefficient and insecure*”. If you do come to deploying your Django project, read the documentation to see an alternative solution for file uploading that can handle a high volume of requests in a much more secure manner.

You can test this setup works by placing an image file in your newly created `media` directory. Drop the file in, start the Django development server, and request the image in your browser. For example, if you added the file `rango.jpg` to `media`, the URL you should enter would look like `http://127.0.0.1:8000/media/rango.jpg`. The image should show in your browser. If it doesn’t, you’ll need to go back and check your setup.

4.5 Basic Workflow

With the chapter complete, you should now know how to setup and create templates, use templates within your views, setup and use Django to send static media files, include images within your templates *and* setup Django’s static media server to allow for file uploads. We’ve actually covered quite a lot!

Creating a template and integrating it within a Django view is a key concept for you to understand. It takes several steps, but becomes second nature to you after a few attempts.

1. First, create the template you wish to use and save it within the `templates` directory you specified in your project’s `settings.py` file. You may wish to use Django template variables (e.g. `{ { variable_name } }`) within your template. You’ll be able to replace these with whatever you like within the corresponding view.
2. Find or create a new view within an application’s `views.py` file.
3. Add your view-specific logic (if you have any) to the view. For example, this may involve extracting data from a database.
4. Within the view, construct a dictionary object which you can pass to the template engine as part of the template’s `context`.
5. Make use of the `RequestContext()` class and `render_to_response()` helper function to generate the rendered response. Ensure you reference the correct template file for the first `render_to_response()` parameter!
6. If you haven’t already done so, map the view to a URL by modifying your project’s `urls.py` file - and the application-specific `urls.py` file if you have one.

The steps involved for getting a static media file onto one of your pages is another important process you should be familiar with. Check out the steps below on how to do this.

1. Take the static media file you wish to use and place it within your project’s `static` directory. This is the directory you specify in your project’s `STATICFILES_DIRS` tuple within `settings.py`.

2. Add a reference to the static media file to a template. For example, an image would be inserted into an HTML page through the use of the `` tag. Remember to use the `{% load static %}` and `{% static "filename" %}` commands within the template to make your life easier!
3. Load the view that utilises the template you modified in your browser. Your static media should appear.

The next chapter will look at databases. We'll see how to make use of Django's excellent database layer to make your life easier and SQL free!

4.6 Exercises

Give the following exercises a go to reinforce what you've learnt from this chapter.

- Convert the `about` page to use a template too from a template called `about.html`.
- Within the `about.html` template, add a picture stored within your project's static media.

Models and Databases

Working with databases often requires you to get your hands dirty messing about with SQL. In Django, a lot of this hassle is taken care of for you by Django's *object relational mapping (ORM)* functions, and how Django encapsulates databases tables through models. Essentially, a model is a Python object that describes your data model/table. Instead of directly working on the database table via SQL, all you have to do is manipulate the corresponding Python object. In this chapter, we'll walkthrough how to setup a database and the models required for Rango.

5.1 Rango's Requirements

First, let's go over the data requirements for Rango. The following list provides the key details of Rango's data requirements.

- Rango is essentially a *web page directory* - a site containing links to other websites.
- There are a number of different *webpage categories*, and each category houses a number of links. We assumed in Chapter 1 that this is a one-to-many relationship. See the Entity Relationship Diagram below.
- A category has a name, number of visits, and number of likes.
- A page refers to a category, has a title, URL and a number of views.

Fig. 5.1: The Entity Relationship Diagram of Rango's two main entities.

5.2 Telling Django About Your Database

Before we can create any models, we must first setup our database configuration. To do this, open the `settings.py` and locate the dictionary called `DATABASES`. Now, modify the `default` key/value pair so it looks something like the following example.

```
DATABASES = {  
    'default': {  
        'ENGINE': 'django.db.backends.sqlite3',  
        'NAME': DATABASE_PATH,  
    }  
}
```

Also create a new variable called `DATABASE_PATH` and add that to the top of your `settings.py` along with your other previously defined paths.

```
DATABASE_PATH = os.path.join(PROJECT_PATH, 'rango.db')
```

Here, we have defined the *default* database to use the SQLite Django backend. This provides us with access to the lightweight python database, [SQLite](#), which is great for development purposes. The only other value we need to set is the NAME key/value pair, which we have set to DATABASE_PATH. For SQLite databases, the remaining keys of USER, PASSWORD, HOST and PORT are not required and can thus be safely removed.

Note: While using an SQLite engine for this tutorial is fine, it may not perhaps be the best option when it comes to deploying your application. Instead, it may be better to use a more robust and scalable database engine. Django comes with out of the box support for several other popular database engines, such as [PostgreSQL](#) and [MySQL](#). See the [official Django documentation on Database Engines](#) for more details. You can also check out [this excellent article](#) on the SQLite website which explains situation where you should and you shouldn't consider using the lightweight SQLite engine.

5.3 Creating Models

With your database configured in `settings.py`, let's now create the two initial data models for the Rango application.

In `rango/models.py`, we will define two classes - both of which must inherit from `django.db.models.Model`. The two Python classes will be the definitions for models representing *categories* and *pages*. Define the Category and Page models as follows.

```
class Category(models.Model):
    name = models.CharField(max_length=128, unique=True)

    def __unicode__(self):
        return self.name

class Page(models.Model):
    category = models.ForeignKey(Category)
    title = models.CharField(max_length=128)
    url = models.URLField()
    views = models.IntegerField(default=0)

    def __unicode__(self):
        return self.title
```

When you define a model, you need to specify the list of attributes and their associated types along with any optional parameters. Django provides a number of built-in fields. Some of the most commonly used are listed below.

- `CharField`, a field for storing character data (e.g. strings). Specify `max_length` to provide a maximum number of characters the field can store.
- `URLField`, much like a `CharField`, but designed for storing resource URLs. You may also specify a `max_length` parameter.
- `IntegerField`, which stores integers.
- `DateField`, which stores a Python `datetime.date`.

Check out the [Django documentation on model fields](#) for a full listing.

For each field, you can specify the `unique` attribute. If set to `True`, only one instance of a particular value in that field may exist throughout the entire database model. For example, take a look at our `Category` model defined above. The field `name` has been set to `unique` - thus every category name must be unique.

This is useful if you wish to use a particular field as an additional database key. You can also specify additional attributes for each field such as specifying a default value (`default='value'`), and whether the value for a field can be `NULL` (`null=True`) or not.

Django also provides simple mechanisms that allows us to relate models/database tables together. These mechanisms are encapsulated in three further field types, and are listed below.

- `ForeignKey`, a field type that allows us to create a one-to-many relationship.
- `OneToOneField`, a field type that allows us to define a strict one-to-one relationship.
- `ManyToManyField`, a field type which allows us to define a many-to-many relationship.

From our model examples above, the field `category` in model `Page` is of type `ForeignKey`. This allows us to create a one-to-many relationship with model/table `Category`, which is specified as an argument to the field's constructor. *You should be aware that Django creates an ID field for you automatically in each table relating to a model. You therefore do not need to explicitly define a primary key for each model.* Django makes your life so much easier!

Note: When creating a Django model, it's good practice to make sure you include the `__unicode__()` method - a method almost identical to the `__str__()` method. If you're unfamiliar with both of these, think of them as methods analogous to the `toString()` method in a Java class. The method is used to provide a unicode representation of a model instance. Our `Category` model for example returns the name of the category in the `__unicode__()` method - something which will be incredibly handy to you when you begin to use the Django admin interface later on in this chapter.

5.4 Creating and Synchronising the Database

With our models defined, we can now let Django work its magic and create the table representations in our database. To do this, we need to synchronise the database using the `manage.py` script by invoking the following command.

```
$ python manage.py syncdb
```

You can then follow the on-screen instructions. You'll notice that Django creates a number of tables within your selected database. Of particular interest to you are the two tables `rango_category` and `rango_page`. These correspond to your two Rango models `Category` and `Page` respectively. Check out the example output the script produces below.

```
$ python manage.py syncdb

Creating tables ...
Creating table auth_permission
Creating table auth_group_permissions
Creating table auth_group
Creating table auth_user_groups
Creating table auth_user_user_permissions
Creating table auth_user
Creating table django_content_type
Creating table django_session
Creating table django_site
Creating table rango_category <-- Rango Category model
Creating table rango_page      <-- Rango Page model

You just installed Django's auth system, which means you don't have any superusers defined.
Would you like to create one now? (yes/no): yes
Username (leave blank to use '<YOUR USERNAME>'): <ENTER A USERNAME>
Email address: <ENTER YOUR E-MAIL ADDRESS>
```

```
Password: <ENTER PASSWORD>
Password (again): <ENTER PASSWORD AGAIN>
Superuser created successfully.
Installing custom SQL ...
Installing indexes ...
Installed 0 object(s) from 0 fixture(s)
```

You may be asked if you wish to create a superuser account, just like in the example above. The superuser account will be used to access the Django admin interface later on in this tutorial, so we recommend that you bite the bullet and set up an account now. Enter yes, then enter a username for the account, your e-mail address and provide a password when prompted. Once completed, the script should finish successfully. Make sure you take a note of the username and password for your superuser account.

After this process has been completed you will be able to see a new file called `rango.db` in your project directory.

Note: Whenever you change your database models, you will have to delete the database file and then re-sync the database by running `python manage.py syncdb` again. If you add a new model however, you can `syncdb` your database without deleting it. This is a known drawback of Django, and can be quite frustrating. *New models will be synchronised, but changes to existing models will not be.* You may have also noticed that we forgot to include a couple of fields in our `Category` model. We will add these in later to remind you of this process. One solution to this problem is using a third party application like [South](#) that handles schema migrations (changes to models). South is currently in active development and is considered a standard solution for schema migrations until something like this becomes part of the standard Django codebase. However, we won't be covering South in this book.

5.5 Generated SQL

The database synchronisation process essentially converts the Django models into SQL tables. To see the SQL that is produced, you can run the command `python manage.py sql rango`. This will output the SQL code that was issued to the database backend on your behalf. Notice that while we did not specify an `id` attribute in our model, Django automatically assigns an `id` to every model. You should also note that in the `rango_page` table, the reference to `Category` is created as a foreign key.

```
BEGIN;
CREATE TABLE "rango_category" (
    "id" integer NOT NULL PRIMARY KEY,
    "name" varchar(128) NOT NULL UNIQUE
)
;
CREATE TABLE "rango_page" (
    "id" integer NOT NULL PRIMARY KEY,
    "category_id" integer NOT NULL REFERENCES "rango_category" ("id"),
    "title" varchar(128) NOT NULL,
    "url" varchar(200) NOT NULL,
    "views" integer NOT NULL
)
;
COMMIT;
```

Since Django provides a wrapper over the database engines all we need to worry about is interfacing with the models (and the Object Relation Mapping). However, you still have the ability to issue SQL commands directly to the database if you wish. This is detailed in the [Official Django Documentation on running custom SQL](#).

5.6 Django Models and the Django Shell

Before we turn our attention to demonstrating the Django admin interface, it's worth noting that you can interact with Django models from the Django shell - a very useful aid for debugging purposes. We'll demonstrate how to create a Category instance using this method.

To access the shell, we need to call `manage.py` from within your Django project's root directory once more. Run the following command.

```
$ python manage.py shell
```

This will start an instance of the Python interpreter and load in your project's settings for you. You can then interact with the models. The following terminal session demonstrates this functionality. Check out the inline commentary to see what each command does.

```
# Import the Category model from the Rango application
>>> from rango.models import Category

# Show all the current categories
>>> print Category.objects.all()
[] # Returns an empty list (no categories have been defined!)

# Create a new category object, and save it to the database.
>>> c = Category(name="Test")
>>> c.save()

# Now list all the category objects stored once more.
>>> print Category.objects.all()
[<Category: test>] # We now have a category called 'test' saved in the database!

# Quit the Django shell.
>>> quit()
```

In the example, we first import the model that we want to manipulate. We then print out all the existing categories, of which there are none because our table is empty. Then we create and save a Category, before printing out all the categories again. This second `print` should then show the Category just added.

Note: The example we provide above is only a very basic taster on database related activities you can perform in the Django shell. If you have not done so already, it is good time to complete part one of the [official Django Tutorial](#) to learn more about interacting with the models. Also check out the [official Django documentation](#) on the list of available commands for working with models.

5.7 Configuring the Admin Interface

One of the stand-out features of Django is that it provides a built in, web-based administrative interface that allows us to browse and edit data stored within our models and corresponding database tables. Like everything else, we need to configure the admin interface for use before we can access it.

First, we must open our Django project's `settings.py` file. This is located within the project configuration directory. Within the file, locate the `INSTALLED_APPS` tuple, and uncomment the line which adds `django.contrib.admin`. The tuple should now look something like the following example. Check out the inline comment to easily spot see the uncommented line.

```
INSTALLED_APPS = (
    'django.contrib.auth',
```

```
'django.contrib.contenttypes',
'django.contrib.sessions',
'django.contrib.sites',
'django.contrib.messages',
'django.contrib.staticfiles',
# Uncomment the next line to enable the admin:
'django.contrib.admin', # THIS LINE SHOULD NOW BE UNCOMMENTED
# Uncomment the next line to enable admin documentation:
# 'django.contrib.admindocs',
    'rango',
)
```

When this has been done, save the file and synchronise the database. We need to do this as the `django.contrib.admin` application needs to add some additional database tables in order to function correctly. To synchronise the database, run the `syncdb` command once more.

```
$ python manage.py syncdb
```

You should see that the table `django_admin_log` is created for you. Once done, open your project's `urls.py` file. This was created in the project configuration directory. Within the file, first uncomment the two lines after the import statements. We also need to ensure that the URL pattern for `/admin/` is present, and points to the `admin.site.urls` module as shown below.

```
from django.conf.urls import patterns, include, url
from django.conf import settings

# Uncomment the next two lines to enable the admin:
from django.contrib import admin # UNCOMMENT THIS LINE
admin.autodiscover() # UNCOMMENT THIS LINE, TOO!

urlpatterns = patterns('',
    url(r'^rango/', include('rango.urls')),
    url(r'^admin/', include(admin.site.urls)), # ADD THIS LINE
)

if settings.DEBUG:
    urlpatterns += patterns(
        'django.views.static',
        (r'media/(?P<path>.*)',
        'serve',
        {'document_root': settings.MEDIA_ROOT}), )
```

We're almost there! Next, we need to now tell the Django admin application what models we wish to make available to the admin interface. To do this, you will need to create a new python file in `rango` application directory called `admin.py`. Add the following code to the new file.

```
from django.contrib import admin
from rango.models import Category, Page

admin.site.register(Category)
admin.site.register(Page)
```

This will *register* the models with the admin interface. If we were to have another model, it would be a trivial case of calling the `admin.site.register()` function, passing the model in as a parameter.

With all of these changes made, start or restart the Django development server and visit: `http://127.0.0.1:8000/admin/`. You should then see a login box, prompting you for a username and password. Enter the superuser username and password you created when setting up your database, and you should then see a webpage similar to that shown in Figure 5.2.

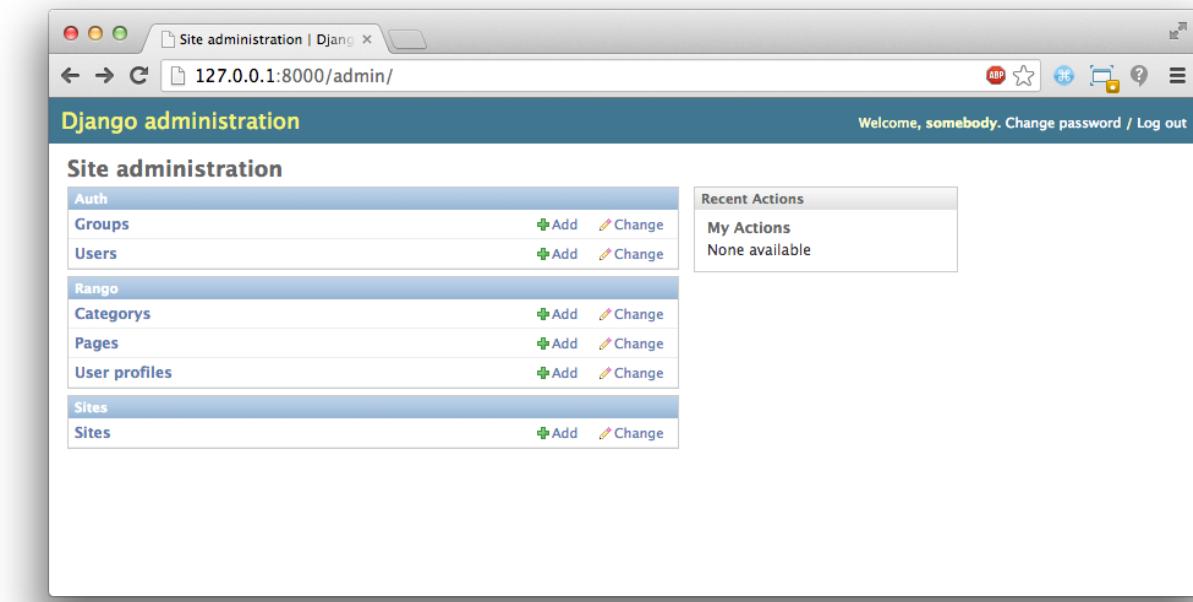


Fig. 5.2: The Django admin interface. Note the Rango category, and the two models contained within.

Try clicking the `Categorys` link within the Rango section. From here, you should see the `test` category that we created via the Django shell. Try deleting the category as we'll be populating the database with a population script next. The interface is easy to use. Spend a few minutes creating, modifying and deleting both categories and pages. You can also add new users who can login to the Django admin interface for your project by adding a user to the `User` in the `Auth` application.

Note: Note the typo within the admin interface (`categorys`, not `categories`). This problem can be fixed by adding a nested `Meta` class into your model definitions with the `verbose_name_plural` attribute. Check out [Django's official documentation on models](#) for more information.

Note: The example `admin.py` file for our Rango application is the most simple, functional example available. There are many different features which you can use in the `admin.py` to perform all sorts of cool customisations, such as changing the way models appear in the admin interface. For this tutorial, we'll stick with the bare-bones admin interface, but you can check out the [official Django documentation on the admin interface](#) for more information if you're interested.

5.8 Creating a Population Script

It's highly likely that during the course of development, you'll come to a point where you will need to modify a Django model. When you do this, the easiest option - without external software - is to re-create your entire database and run `python manage.py syncdb ...again!` Since this slow and repetitive task can be such a pain, it's good practice to create what we call a *population script* for your database. This script is designed to automatically populate your database with test data for you, which can potentially save you lots of time.

To create a population script for Rango's database, we start by creating a new Python module within our Django project's root directory (e.g. `<workspace>/tango_with_django_project/`). Create the

How to Tango with Django, Release 1

populate_rango.py file and add the following code.

```
import os

def populate():
    python_cat = add_cat('Python')

    add_page(cat=python_cat,
        title="Official Python Tutorial",
        url="http://docs.python.org/2/tutorial/")

    add_page(cat=python_cat,
        title="How to Think like a Computer Scientist",
        url="http://www.greenteapress.com/thinkpython/")

    add_page(cat=python_cat,
        title="Learn Python in 10 Minutes",
        url="http://www.korokithakis.net/tutorials/python/")

    django_cat = add_cat("Django")

    add_page(cat=django_cat,
        title="Official Django Tutorial",
        url="https://docs.djangoproject.com/en/1.5/intro/tutorial01/")

    add_page(cat=django_cat,
        title="Django Rocks",
        url="http://www.djangorocks.com/")

    add_page(cat=django_cat,
        title="How to Tango with Django",
        url="http://www.tangowithdjango.com/")

    frame_cat = add_cat("Other Frameworks")

    add_page(cat=frame_cat,
        title="Bottle",
        url="http://bottlepy.org/docs/dev/")

    add_page(cat=frame_cat,
        title="Flask",
        url="http://flask.pocoo.org")

    # Print out what we have added to the user.
    for c in Category.objects.all():
        for p in Page.objects.filter(category=c):
            print "- {0} - {1}".format(str(c), str(p))

    def add_page(cat, title, url, views=0):
        p = Page.objects.get_or_create(category=cat, title=title, url=url, views=views)[0]
        return p

    def add_cat(name):
        c = Category.objects.get_or_create(name=name)[0]
        return c

    # Start execution here!
    if __name__ == '__main__':
        print "Starting Rango population script..."
```

```
os.environ.setdefault('DJANGO_SETTINGS_MODULE', 'tango_with_django_project.settings')
from rango.models import Category, Page
populate()
```

While this looks like a lot of code, what it does is relatively simple. As we define a series of functions at the top of the file, code execution begins towards the bottom - look for the line `if __name__ == '__main__'`. We then import Rango's settings and our models `Category` and `Page` and call the `populate()` function.

Warning: When importing Django models, make sure you have imported your project's settings by that stage. If you don't, an exception will be raised. This is why we import `Category` and `Page` towards the end of the population script, rather than at the top.

The `populate()` function is responsible for calling the `add_cat()` and `add_page()` functions, who are in turn responsible for the creation of new categories and pages respectively. `populate()` keeps tabs on category references for us as we create each individual `Page` model instance and store them within our database. Finally, we loop through our `Category` and `Page` models to print to the user all the `Page` instances and their corresponding categories.

Note: We make use of the convenience `get_or_create()` function for creating model instances. Check out the [official Django documentation](#) for more information on this function. The documentation will explain why we pass index `[0]` of the function's returned value!

When saved, we can run the script by changing the current working directory in a terminal to our Django project's root and executing the module with the command `$ python populate_rango.py`. You should then see output similar to that shown below.

```
$ python populate_rango.py
Starting Rango population script...
- Python - Official Python Tutorial
- Python - How to Think like a Computer Scientist
- Python - Learn Python in 10 Minutes
- Django - Official Django Tutorial
- Django - Django Rocks
- Django - How to Tango with Django
- Other Frameworks - Bottle
- Other Frameworks - Flask
```

Now let's verify that the population script populated the database. Restart the Django development server, navigate to the admin interface, and check that you have some new categories and pages. Do you see all the pages if you click Pages, like in Figure 5.3?

A population script takes a bit of time to write but if you ever change your models, you'll be glad you have it. Remember to update your population script when you update your models too - you don't want the two to get out of sync with each other.

5.9 Basic Workflows

Now that we've covered the core principles of dealing with Django's models functionality, now is a good time to summarise the processes involved in setting everything up. We've split the core tasks into separate sections for you.

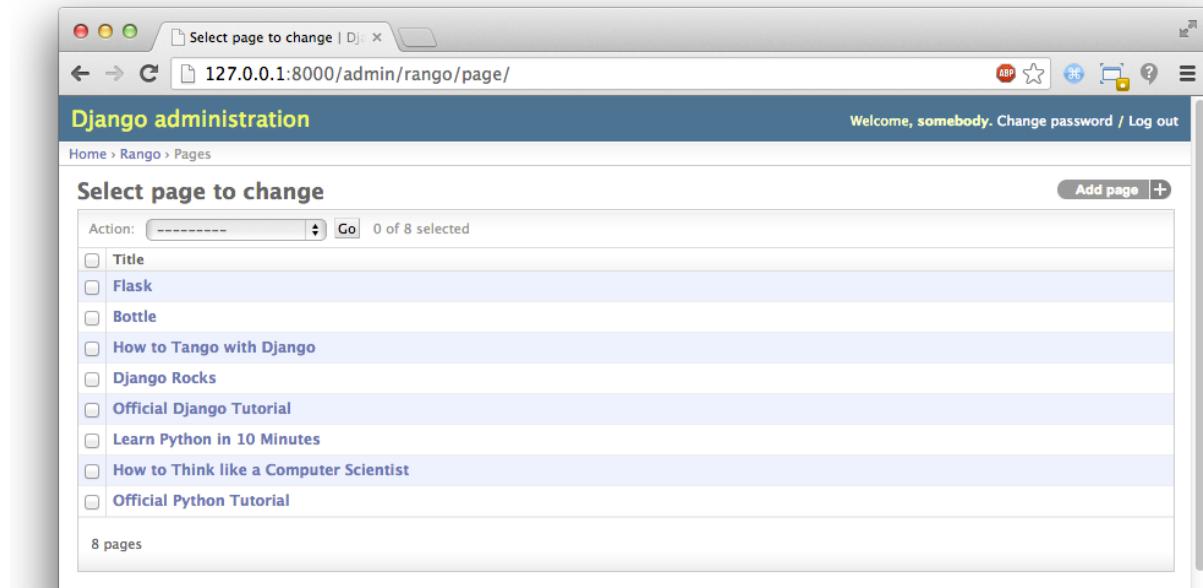


Fig. 5.3: The Django admin interface, showing the Page table populated with sample data from our population script.

5.9.1 Setting up your Database

With a new Django project, you should first tell Django about the database you intend to use (i.e. configure DATABASES in settings.py). Without doing so, Django won't have anywhere to store your data. You can also enable the admin interface to make your life a little bit easier - and remember, you can always disable it later if you want to.

5.9.2 Adding a Model

The workflow for adding models can be broken down into five steps.

1. First, create your new model(s) in your Django application's `models.py` file.
2. With the model created, reconfigure the admin interface to include your new model(s), if you are using it.
3. You should then synchronise or resynchronise your database with the `$ python manage.py syncdb` command. This will create the necessary infrastructure within the database for your new model(s).
4. Create/Edit and then run your population script for your new model(s).

You should also remember about the nuances of the `syncdb` command. Note that the command can be used only for adding new models to your database - if you wish to amend an existing model, you must recreate the database.

5.10 Exercises

Now that you've completed the chapter, try out these exercises to reinforce and practice what you have learnt.

- Update the Category model to include the additional attributes, `views` and `likes` where the default value is zero.

- Re-sync your database, and update your population script so that the Python category has 128 views and 64 likes, the Django category has 64 views and 32 likes, and the Other Frameworks category has 32 views and 16 likes.
- Undertake the [part two of official Django tutorial](#) if you have not done so. This will help to reinforce further what you have learnt here, and to learn more about customising the admin interface.
- Customise the Admin Interface - so that when you view the Page model it displays in a list the category, the name of the page and the url.

5.10.1 Hints

If you require some help or inspiration to get these exercises done, these hints will hopefully help you out.

- To customise the admin interface, you will need to edit `rango/admin.py` and create a `PageAdmin` class that inherits from `admin.ModelAdmin`.
- Within your new `PageAdmin` class, add `list_display = ('title', 'category', 'url')`.
- Finally, register the `PageAdmin` class with Django's admin interface. You should modify the line `admin.site.register(Page)`. Change it to `admin.site.register(Page, PageAdmin)` in Rango's `admin.py` file.

Action:	Title	Category	Url
<input type="checkbox"/>	Flask	Other Frameworks	http://flask.pocoo.org
<input type="checkbox"/>	Bottle	Other Frameworks	http://bottlepy.org/docs/dev/
<input type="checkbox"/>	How to Tango with Django	Django	http://www.tangowithdjango.com/
<input type="checkbox"/>	Django Rocks	Django	http://www.djangoproject.com/
<input type="checkbox"/>	Official Django Tutorial	Django	https://docs.djangoproject.com/en/1.5/intro/tutorial01/
<input type="checkbox"/>	Learn Python in 10 Minutes	Python	http://www.korokithakis.net/tutorials/python/
<input type="checkbox"/>	How to Think like a Computer Scientist	Python	http://www.greenteapress.com/thinkpython/
<input type="checkbox"/>	Official Python Tutorial	Python	http://docs.python.org/2/tutorial/

8 pages

Fig. 5.4: The updated admin interface page view, complete with columns for category and URL.

Models, Templates and Views

Now that we have our models set up and populated with some data, we can now start putting things together. We'll be figuring out how to source data from our models in our views, and how to present this data within our templates.

6.1 Basic Workflow: Data Driven Pages

There are five main steps that you must undertake to create a data driven webpage in Django.

1. First, import the models you wish to use into your application's `views.py` file.
2. Within the view you wish to use, query the model to get the data you want to present.
3. Pass the results from your model into the template's context.
4. Setup your template to present the data to the user in whatever way you wish.
5. If you have not done so already, map a URL to your view.

These steps highlight how Django's framework separates the concerns between models, views and templates.

6.2 Showing Categories on Rango's Homepage

One of the requirements regarding the main pages was to show the top five rango'ed categories.

6.2.1 Importing Required Models

To fulfil this requirement, we will go through each of the above steps. First, open `rango/views.py` and import the `Category` model from Rango's `models.py` file.

```
# Import the Category model
from rango.models import Category
```

6.2.2 Modifying the Index View

With the first step out of the way, we then want to modify our `index()` function. If we cast our minds back, we should remember the `index()` function is responsible for the main page view. Modify the function to look like the example below.

```
def index(request):
    # Obtain the context from the HTTP request.
    context = RequestContext(request)

    # Query the database for a list of ALL categories currently stored.
    # Order the categories by no. likes in descending order.
    # Retrieve the top 5 only - or all if less than 5.
    # Place the list in our context_dict dictionary which will be passed to the template engine.
    category_list = Category.objects.order_by('-likes')[:5]
    context_dict = {'categories': category_list}

    # Render the response and send it back!
    return render_to_response('rango/index.html', context_dict, context)
```

Here we have performed steps two and three in one go. First, we queried the `Category` model to retrieve the top five categories. Here we used the `order_by()` method to sort by the number of likes in descending order - hence the inclusion of the `-`. We then restricted this list to the first 5 `Category` objects in the list.

With the query complete, we passed a reference to the list (stored as variable `category_list`) to a new dictionary, `context_dict`. This dictionary is then passed as part of the context for the template engine in the `render_to_response()` call.

6.2.3 Modifying the Index Template

With the view updated, all that is left for us to do is update the template `rango/index.html`, located within your project's `templates` directory. Change the HTML code of the file so that it looks like the example shown below.

```
<!DOCTYPE html>
<html>
    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Rango says...hello world!</h1>

        {% if categories %}
            <ul>
                {% for category in categories %}
                    <li>{{ category.name }}</li>
                {% endfor %}
            </ul>
        {% else %}
            <strong>There are no categories present.</strong>
        {% endif %}

        <a href="/rango/about/">About</a>
    </body>
</html>
```

Here, we make use of Django's template language to present the data using `if` and `for` control statements. Within the `<body>` of the page, we test to see if `categories` - the name of the context variable containing our list - actually contains any categories (i.e. `{% if categories %}`).

If so, we proceed to construct an unordered HTML list (within the `` tags). The for loop (`{% for category in categories %}`) then iterates through the list of results, printing out each category's name (`{{ category.name }}`) within a pair of `` tags to indicate a list element.

If no categories exist, a message is displayed instead indicating so.

As the example shows in Django's template language, all commands are enclosed within the tags { % and % }, while variables are referenced within { { and } } brackets.

If you now visit Rango's homepage at <http://127.0.0.1:8000/rango/>, you should see a list of three categories underneath the page title just like in Figure 6.1.

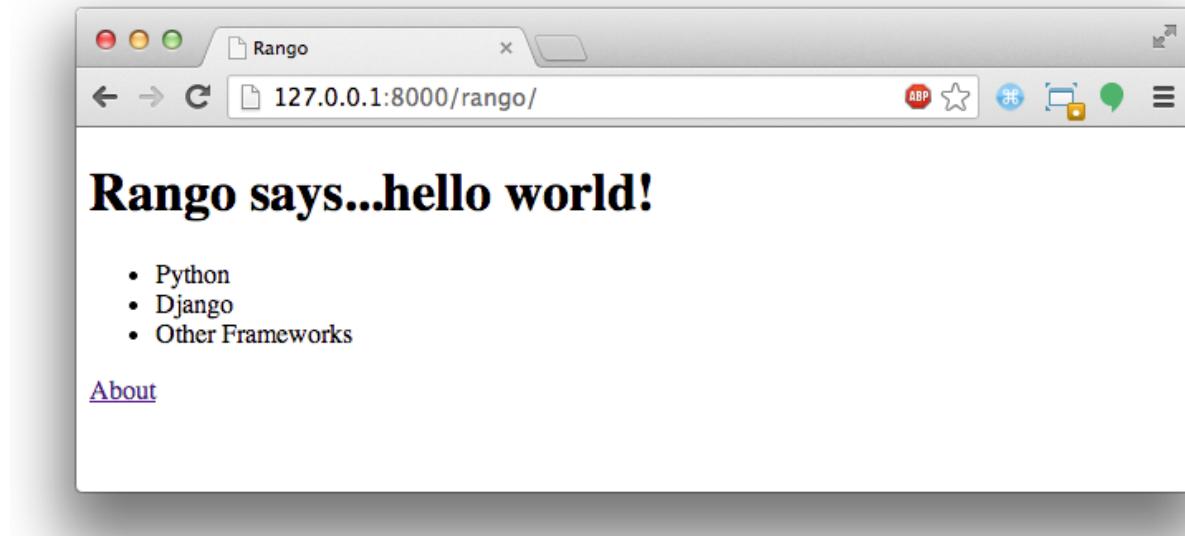


Fig. 6.1: The Rango homepage - now dynamically generated - showing a list of categories. How exciting!

6.3 Creating a Details Page

According to Rango's specification, we also need to show a list of pages that are associated with each category. We have a number of challenges here to overcome. A new view must be created, which should be parameterised. We also need to create URL patterns and URL strings that encode category names.

6.3.1 URL Design and Mapping

Let's start by considering the URL problem. One way we could handle this problem is to use the unique ID for each category within the URL. For example, we could create URLs like `/rango/category/1/` or `/rango/category/2/`, where the numbers correspond to the categories with unique IDs 1 and 2 respectively. However, these URLs are not easily understood by humans. Although we could probably infer that the number relates to a category, how would a user know what category relates to unique IDs 1 or 2? The user wouldn't know without trying.

Instead, we could just use the category name as part of the URL. `/rango/category/Python/` should give us a list of pages related to the Python category. This is a simple, readable and meaningful URL. If we go with this approach, we'll have to handle categories which have multiple words, like 'Other Frameworks', etc.

Note: Designing clean URLs is an important aspect of web design. See [Wikipedia's article on Clean URLs](#) for more details.

6.3.2 Category Page Workflow

With our URLs design chosen, let's get started. We'll undertake the following steps.

1. Import the Page model into `rango/views.py`.
2. **Create a new view in `rango/views.py` - called `category` - The `category` view will take an additional parameter, `category_name_url`.**
 - We will need some help functions to encode and decode the `category_name_url`.
3. Create a new template, `templates/rango/category.html`.
4. Update Rango's `urlpatterns` to map the new `category` view to a URL pattern in `rango/urls.py`.

We'll also need to update the `index()` view and `index.html` template to provide links to the category page view.

6.3.3 Category View

In `rango/views.py`, we first need to import the `Page` model. This means we must add the following import statement at the top of the file.

```
from rango.models import Page
```

Next, we can add our new view, `category()`.

```
def category(request, category_name_url):  
    # Request our context from the request passed to us.  
    context = RequestContext(request)  
  
    # Change underscores in the category name to spaces.  
    # URLs don't handle spaces well, so we encode them as underscores.  
    # We can then simply replace the underscores with spaces again to get the name.  
    category_name = category_name_url.replace('_', ' ')  
  
    # Create a context dictionary which we can pass to the template rendering engine.  
    # We start by containing the name of the category passed by the user.  
    context_dict = {'category_name': category_name}  
  
    try:  
        # Can we find a category with the given name?  
        # If we can't, the .get() method raises a DoesNotExist exception.  
        # So the .get() method returns one model instance or raises an exception.  
        category = Category.objects.get(name=category_name)  
  
        # Retrieve all of the associated pages.  
        # Note that filter returns >= 1 model instance.  
        pages = Page.objects.filter(category=category)  
  
        # Adds our results list to the template context under name pages.  
        context_dict['pages'] = pages  
        # We also add the category object from the database to the context dictionary.  
        # We'll use this in the template to verify that the category exists.  
        context_dict['category'] = category  
    except Category.DoesNotExist:  
        # We get here if we didn't find the specified category.  
        # Don't do anything - the template displays the "no category" message for us.  
        pass
```

```
# Go render the response and return it to the client.
return render_to_response('rango/category.html', context_dict, context)
```

Our new view follows the same basic steps as our `index()` view. We first obtain the context of the request, then build a context dictionary, render the template, and send the result back. In this case, the difference is that the context dictionary building is a little more complex. To build our context dictionary, we need to determine which category to look at by using the value passed as parameter `category_name_url` to the `category()` view function. Once we have determined which category to look for, we can pull the relevant information from the database and append the results to our context dictionary, `context_dict`. We'll figure out how to get the value for `category_name_url` from the URL shortly.

You will have also seen in the `category()` view function we assume that the `category_name_url` is the category name where spaces are converted to underscores. We therefore replace all the underscores with spaces. This is unfortunately a pretty crude way to handle the decoding and encoding of the category name within the URL. As an exercise later, it'll be your job to create two functions to encode and decode category name.

Warning: While you can use spaces in URLs, it is considered to be unsafe to use them. Check out [IETF Memo on URLs](#) to read more.

6.3.4 Category Template

Now let's create our template for the new view. In `<workspace>/tango_with_django_project/templates/rango/` directory, create `category.html`. In the new file, add the following code.

```
<!DOCTYPE html>
<html>
    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>{{ category_name }}</h1>
        {% if category %}
            {% if pages %}
                <ul>
                    {% for page in pages %}
                        <li><a href="{{ page.url }}">{{ page.title }}</a></li>
                    {% endfor %}
                </ul>
            {% else %}
                <strong>No pages currently in category.</strong>
            {% endif %}
        {% else %}
            The specified category {{ category_name }} does not exist!
        {% endif %}
    </body>
</html>
```

The HTML code example again demonstrates how we utilise the data passed to the template via its context. We make use of the `category_name` variable and our `category` and `pages` objects. If `category` is not defined within our template context, the category was not found within the database, and a friendly error message is displayed stating this fact. If the opposite is true, we then proceed to check for `pages`. If `pages` is undefined or contains no elements, we display a message stating there are no pages present. Otherwise, the pages within the category are presented in a HTML list. For each page in the `pages` list, we present their `title` and `url` attributes.

Note: This a Django template `{% if %}` statement with an object is a really neat way of determining the existence of the object within the template's context. Try getting into the habit of performing these checks to reduce the scope for potential exceptions that could be raised within your code.

6.3.5 Parameterised URL Mapping

Now let's have a look at how we actually pass the value of the `category_name_url` parameter to the `category()` function. To do so, we need to modify Rango's `urls.py` file and update the `urlpatterns` tuple as follows.

```
urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^about/$', views.about, name='about'),
    url(r'^category/(?P<category_name_url>\w+)/$', views.category, name='category'),) # New!
```

As you can see, we have added in a rather complex entry that will invoke `view.category()` when the regular expression `r'^ (?P<category_name_url>\w+)/$'` is matched. We set up our regular expression to look for any sequence of word characters (e.g. a-z, A-Z, _, or 0-9) before the trailing URL slash. This value is then passed to the view `views.category()` as parameter `category_name_url`, the only argument after the mandatory `request` argument. Essentially, the name you hard-code into the regular expression is the name of the argument that Django looks for in your view's function definition.

Note: Regular expressions may seem horrible and confusing at first, but there are tons of resources online to help you. [This cheat sheet](#) provides you with an excellent resource for fixing pesky regular expression problems.

6.3.6 Modifying the Index View and Template

Our new view is set up and ready to go - but we need to do one more thing. Our index page view needs to be updated to provide users with a means to view the category pages that are listed. Update in the `index()` in `rango/views.py` as follows.

```
def index(request):
    # Obtain the context from the HTTP request.
    context = RequestContext(request)

    # Query for categories - add the list to our context dictionary.
    category_list = Category.objects.order_by('-likes')[:5]
    context_dict = {'categories': category_list}

    # The following two lines are new.
    # We loop through each category returned, and create a URL attribute.
    # This attribute stores an encoded URL (e.g. spaces replaced with underscores).
    for category in category_list:
        category.url = category.name.replace(' ', '_')

    # Render the response and return to the client.
    return render_to_response('rango/index.html', context_dict, context)
```

As explained in the inline commentary, we take each category that the database returns, then iterate through the list of categories encoding the name to make it URL friendly. This URL friendly value is then placed as an attribute inside the `Category` object (i.e. we take advantage of Python's dynamic typing to add this attribute on the fly).

We then pass the list of categories - `category_list` - to the context of the template so it can be rendered. With a `url` attribute now available for each category, we can update our `index.html` template to look like the example

below.

```
<!DOCTYPE html>
<html>
    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Rango says..hello world!</h1>

        {%
            if categories
        %}
            <ul>
                {%
                    for category in categories
                %}
                    <!-- Following line changed to add an HTML hyperlink -->
                    <li><a href="/rango/category/{{ category.url }}"/>{{ category.name }}</a></li>
                {%
                    endfor
                %}
            </ul>
        {%
            else
        %}
            <strong>There are no categories present.</strong>
        {%
            endif
        }

    </body>
</html>
```

Here we have updated each list element (``) adding a HTML hyperlink (`<a>`). The hyperlink has an `href` attribute, which we use to specify the target URL defined by `{{ category.url }}`.

6.3.7 Demo

Let's try everything out now by visiting the Rango's homepage. You should see your homepage listing all the categories. The categories should now be clickable links. Clicking on Python should then take you to the Python detailed category view, as demonstrated in Figure 6.2. If you see a list of links like Official Python Tutorial, then you've successfully set up the new view. Try navigating a category which doesn't exist, like `/rango/category/computers`. You should see a message telling you that no pages exist in the category.

6.4 Exercises

Reinforce what you've learnt in this chapter by trying out the following exercises.

- Modify the index page to also include the top 5 most viewed pages.
- The encoding and decoding of the Category name to a URL is not particularly robust as every time we need to access the page we need to decode the URL. One way we can make sure the encoding and decoding is to use some functions to perform the encoding and decoding.
- While the encode and decode function help it is only part of the solution. How would you redesign the URL encoding so that it is more seamless and doesn't require encoding and decoding every time?
- Undertake the part three of official Django tutorial if you have not done so already to further what you've learnt here.

6.4.1 Hints

To help you with the exercises above, the following hints may be of some use to you. Good luck!

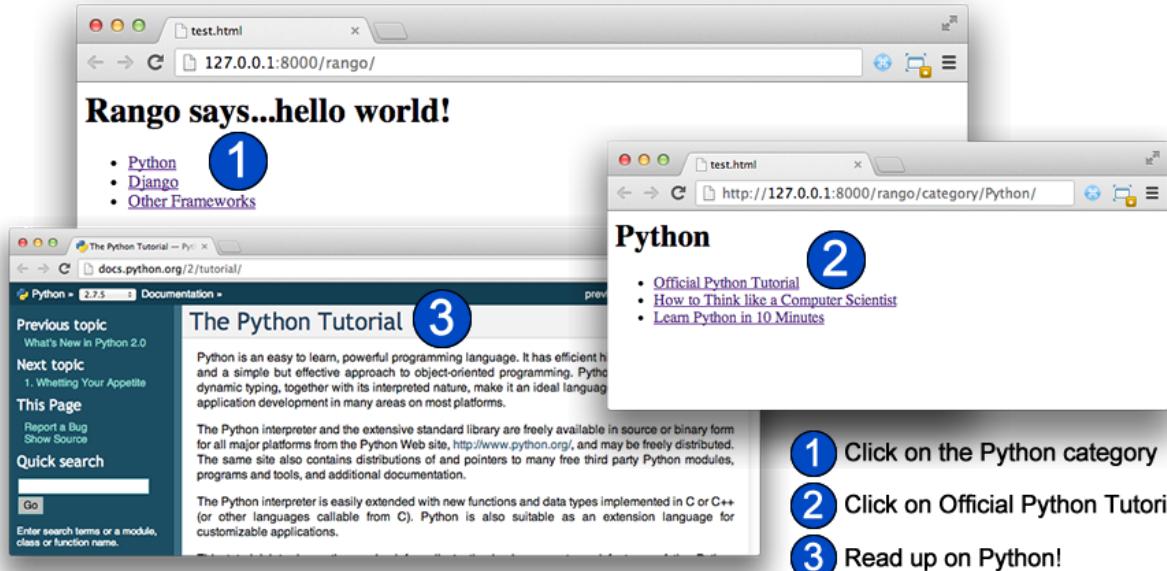


Fig. 6.2: What your link structure should now look like. Starting with the Rango homepage, you are then presented with the category detail page. Clicking on a page link takes you to the linked website.

- Update the population script to add some value to the views count for each page.
- Create an encode and decode function to convert `category_name_url` to `category_name` and vice versa.

Fun with Forms

So far we have only presented data through the views and templates that we have created. In this chapter, we will run through how to capture data through web forms. Django comes with some neat form handling functionality, making it a pretty straightforward process to gather information from users and send it back to your web application. According to [Django's documentation on forms](#), the form handling functionality allows you to:

1. display an HTML form with automatically generated *form widgets* (like a text field or date picker);
2. check submitted data against a set of validation rules;
3. redisplay a form in case of validation errors; and
4. convert submitted form data to the relevant Python data types.

One of the major advantages of using Django's forms functionality is that it can save you a lot of time and HTML hassle. This part of the tutorial will look at how to implement the necessary infrastructure that will allow users of Rango to add categories and pages to the database via forms.

7.1 Basic Workflow

The basic steps involved in creating a form and allowing users to enter data via the form is as follows.

1. If you haven't already got one, create a `forms.py` file within your Django application's directory to store form-related classes.
2. Create a `ModelForm` class for each model that you wish to represent as a form.
3. Customise the forms as you desire.
4. Create or update a view to handle the form - including *displaying* the form, *saving* the form data, and *flagging up errors* which may occur when the user enters incorrect data (or no data at all) in the form.
5. Create or update a template to display the form.
6. Add a `urlpattern` to map to the new view (if you created a new one).

This workflow is a bit more complicated than previous workflows, and the views that we have to construct have a lot more complexity as well. However, once you undertake the process a few times it will be pretty clear how everything pieces together.

7.2 Page and Category Forms

First, create a file called `forms.py` within the `rango` application directory. While this step is not absolutely necessary, as you could put the forms in the `models.py`, this makes the codebase a lot cleaner and clearer to understand.

7.2.1 Creating ModelForm Classes

Within Rango's `forms.py` module, we will be creating a number of classes that inherit from Django's `ModelForm`. In essence, a `ModelForm` is a *helper class* that allows you to create a Django Form from a pre-existing model. As we've already got two models defined for Rango (`Category` and `Page`), we'll create `ModelForms` for both.

In `rango/forms.py` add the following code.

```
from django import forms
from rango.models import Page, Category

class CategoryForm(forms.ModelForm):
    name = forms.CharField(max_length=128, help_text="Please enter the category name.")
    views = forms.IntegerField(widget=forms.HiddenInput(), initial=0)
    likes = forms.IntegerField(widget=forms.HiddenInput(), initial=0)

    # An inline class to provide additional information on the form.
    class Meta:
        # Provide an association between the ModelForm and a model
        model = Category

class PageForm(forms.ModelForm):
    title = forms.CharField(max_length=128, help_text="Please enter the title of the page.")
    url = forms.URLField(max_length=200, help_text="Please enter the URL of the page.")
    views = forms.IntegerField(widget=forms.HiddenInput(), initial=0)

    # What fields do we want to include in our form?
    # This way we don't need every field in the model present.
    # Some fields may allow NULL values, so we may not want to include them...
    # Here, we are hiding the foreign key.
    fields = ('title', 'url', 'views')
```

Django provides us with a number of ways to customise the forms that are created on our behalf. In the code sample above, we've specified the widgets that we wish to use for each field to be displayed. For example, in our `PageForm` class, we've defined `forms.CharField` for the `title` field, and `forms.URLField` for `url` field. Both fields provide text entry for users. Note the `max_length` parameters we supply to our fields - the lengths that we specify are identical to the maximum length of each field we specified in the underlying data models. Go back to Chapter 5 to check for yourself, or have a look at Rango's `models.py` file.

You will also notice that we have included several `IntegerField` entries for the `views` and `likes` fields in each form. Note that we have set the `widget` to be hidden with the parameter setting `widget=forms.HiddenInput()`, and then set the value to zero with `initial=0`. This is one way to set the field to zero without giving the control to the user as the field will be hidden, yet the form will provide the value to the model. However, as you can see in the `PageForm`, despite the fact that we have a hidden field, we still need to include the field in the form. If in `fields` we excluded `views`, then the form would not contain that field (despite it being specified) and so the form would not return the value zero for that field. This may raise an error depending on how the model has been set up. If in the models we specified that the `default=0` for these fields then we can rely on the model to automatically populate

field with the default value - and thus avoid a `not null` error. In this case, it would not be necessary to have these hidden fields. Essentially, you need to be careful when you define your models and forms to make sure that form is going to contain and pass on all the data that is required to populate your model correctly.

Besides the `CharField` and `IntegerField` widget, many more are available for use. As an example, Django provides `IntegerField` (integer entry), `ChoiceField` (radio input buttons), and `DateField` (for date/time entry). There are many other field types you can use, which perform error checking for you (e.g. *is the value provided a valid integer?*). We highly recommend you have a look at the [official Django documentation on widgets](#) to see what components exist and the arguments you can provide to customise them.

Perhaps the most important aspect of a class inheriting from `ModelForm` is the need to define *which model we're wanting to provide a form for*. We take care of this through our nested `Meta` class. Set the `model` attribute of the nested `Meta` class to the model you wish to use. For example, our `CategoryForm` class has a reference to the `Category` model. This is a crucial step enabling Django to take care of creating a form in the image of the specified model. It will also help in handling flagging up any errors along with saving and displaying the data in the form.

We also use the `Meta` class to specify which fields that we wish to include in our form through the `fields` tuple. Use a tuple of field names to specify the fields you wish to include.

Note: We highly recommend you check out the [official Django documentation on forms](#) for further information about how to customise them.

7.2.2 Creating an Add Category View

With our `CategoryForm` class now defined, we're now ready to create a new view to display the form and handle the posting of form data. To do this, add the following code to `rango/views.py`.

```
from rango.forms import CategoryForm

def add_category(request):
    # Get the context from the request.
    context = RequestContext(request)

    # A HTTP POST?
    if request.method == 'POST':
        form = CategoryForm(request.POST)

        # Have we been provided with a valid form?
        if form.is_valid():
            # Save the new category to the database.
            form.save(commit=True)

            # Now call the index() view.
            # The user will be shown the homepage.
            return index(request)
        else:
            # The supplied form contained errors - just print them to the terminal.
            print form.errors
    else:
        # If the request was not a POST, display the form to enter details.
        form = CategoryForm()

        # Bad form (or form details), no form supplied...
        # Render the form with error messages (if any).
    return render_to_response('rango/add_category.html', {'form': form}, context)
```

The new `add_category()` view adds several key pieces of functionality for handling forms. First, we access the context surrounding the HTTP request. This then allows us to determine the type of request being made - whether it be a HTTP GET or POST. This allows us to handle different requests appropriately - whether we want to show a form (i.e. on GET), or process form data (i.e. on POST) - all from the same URL. The `add_category()` view function can handle three different scenarios:

- showing a new, blank form for adding a category;
- saving form data provided by the user to the associated model, and rendering the Rango homepage; and
- if there are errors, redisplay the form with error messages.

Note: What do we mean by GET and POST? They are two different types of *HTTP requests*.

- A HTTP GET is used to *request a representation of the specified resource*. In other words, we use a HTTP GET to retrieve a particular resource, whether it be a webpage, image or other file.
- In contrast, a HTTP POST *submits data from the client's web browser to be processed*. This type of request is used for example when submitting the contents of a HTML form.
- Ultimately, a HTTP POST may end up being programmed to create a new resource (e.g. a new database entry) on the server. This can later be accessed through a HTTP GET request.

Django's form handling machinery has also been utilised to process the data returned from a user's browser via a HTTP POST request. It not only handles the saving of form data into the chosen model, but will also automatically generate any error messages for each form field (if any are required). This means that Django will not store any submitted forms with missing information which could potentially cause problems for your database's referential integrity. For example, supplying no value in the category name field will return an error, as the field cannot be blank.

You'll notice from the line in which we call `render_to_response()` that we refer to a new template called `add_category.html` which will contain the relevant Django template code and HTML for the form and page.

7.2.3 Creating the Add Category Template

Create the file `templates/rango/add_category.html`. Within the file, add the following HTML markup and Django template code.

```
<!DOCTYPE html>
<html>
    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Add a Category</h1>

        <form id="category_form" method="post" action="/rango/add_category/">

            {% csrf_token %}
            {% for hidden in form.hidden_fields %}
                {{ hidden }}
            {% endfor %}

            {% for field in form.visible_fields %}
                {{ field.errors }}
                {{ field.help_text}}
                {{ field }}
            {% endfor %}

        </form>
    </body>
</html>
```

```

<input type="submit" name="submit" value="Create Category" />
</form>

</body>

</html>

```

Now, what does this code do? You can see that within the `<body>` of the HTML page that we place a `<form>` element. Looking at the attributes for the `<form>` element, you can see that all data captured within this form is sent to the URL `/rango/add_category/` as a HTTP POST request (the `method` attribute is case insensitive, so you can do `POST` or `post` - both provide the same functionality). Within the form, we have two for loops - one controlling *hidden* form fields, the other *visible* form fields - with visible fields controlled by the `fields` attribute of your `ModelForm Meta` class. These loops produce HTML markup for each form element. For visible form fields, we also add in any errors that may be present with a particular field and help text which can be used to explain to the user what he or she needs to enter.

Note: The need for hidden as well as visible form fields is necessitated by the fact that HTTP is a stateless protocol. You can't persist state between different HTTP requests which can make certain parts of web applications difficult to implement. To overcome this limitation, hidden HTML form fields were created which allow web applications to pass important information to a client (which cannot be seen on the rendered page) in a HTML form, only to be sent back to the originating server when the user submits the form.

You should also take note of the code snippet `{% csrf_token %}`. This is a *Cross-Site Request Forgery (CSRF) token*, which helps to protect and secure the HTTP POST action that is initiated on the subsequent submission of a form. *The CSRF token is required by the Django framework. If you forget to include a CSRF token in your forms, a user may encounter errors when he or she submits the form.* Check out the official Django documentation on `CSRF tokens` for more information about this.

7.2.4 Mapping the Add Category View

Now we need to map the `add_category()` view to a URL. In the template we have used the URL `/rango/add_category/` in the form's submit attribute. So we will need to follow suit in `rango/urls.py` and modify the `urlpatterns` as follows.

```

urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^about/$', views.about, name='about'),
    url(r'^add_category/$', views.add_category, name='add_category'), # NEW MAPPING!
    url(r'^category/(?P<category_name_url>\w+)$', views.category, name='category'),
)

```

Ordering doesn't necessarily matter in this instance. However, take a look at the [official Django documentation](#) on how Django process a request for more information. Our new URL for adding a category is `/rango/add_category/`.

7.2.5 Modifying the Index Page View

As a final step let's put a link on the index page so that we can easily add categories. Edit the template `rango/index.html` and add the following HTML hyperlink just before the `</body>` closing tag.

```

<a href="/rango/add_category/">Add a New Category</a>

```

7.2.6 Demo

Now let's try it out! Run your Django development server, and navigate to `http://127.0.0.1:8000/rango/`. Use your new link to jump to the add category page, and try adding a category. Figure 7.1 shows screenshots of the Add Category and Index Pages.

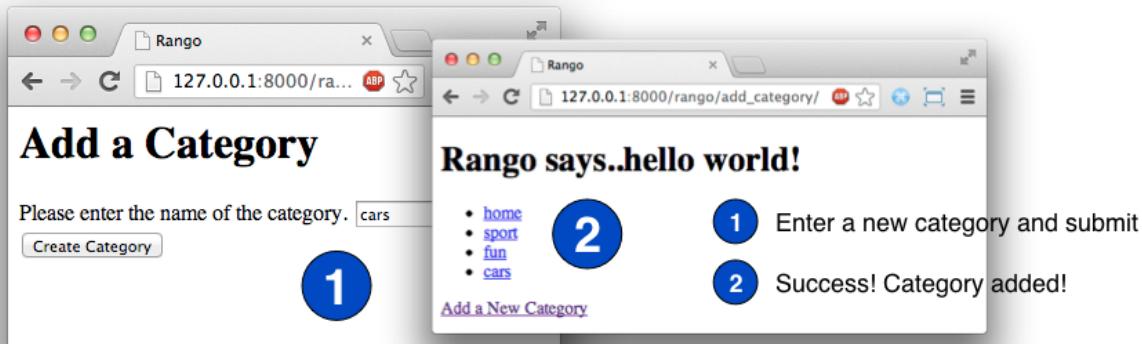


Fig. 7.1: Adding a new category to Rango with our new form. The diagram illustrates the steps involved.

7.2.7 Cleaner Forms

Recall that our `Page` model has a `url` attribute set to an instance of the `URLField` type. In a corresponding HTML form, Django would reasonably expect any text entered into a `url` field to be a well-formed, complete URL. However, users can find entering something like `http://www.url.com` to be cumbersome - indeed, users *may* not even know what forms a correct URL!

In scenarios where user input may not be entirely correct, we can *override* the `clean()` method implemented in `ModelForm`. This method is called upon before saving form data to a new model instance, and thus provides us with a logical place to insert code which can verify - and even fix - any form data the user inputs. In our example above, we can check if the value of `url` field entered by the user starts with `http://` - and if it doesn't, we can prepend `http://` to the user's input.

```
class PageForm(forms.ModelForm):  
    ...  
  
    def clean(self):  
        cleaned_data = self.cleaned_data  
        url = cleaned_data.get('url')  
  
        # If url is not empty and doesn't start with 'http://', prepend 'http://'.  
        if url and not url.startswith('http://'):   
            url = 'http://' + url  
            cleaned_data['url'] = url  
  
        return cleaned_data
```

Within the `clean()` method, a simple pattern is observed which you can replicate in your own Django form handling code.

1. Form data is obtained from the `ModelForm` dictionary attribute `cleaned_data`.
2. Form fields that you wish to check can then be taken from the `cleaned_data` dictionary. Use the `.get()` method provided by the dictionary object to obtain the form's values. If a user does not enter a value into a form field, its entry will not exist in the `cleaned_data` dictionary. In this instance, `.get()` would return `None` rather than raise a `KeyError` exception. This helps your code look that little bit cleaner!
3. For each form field that you wish to process, check that a value was retrieved. If something was entered, check what the value was. If it isn't what you expect, you can then add some logic to fix this issue before *reassigning* the value in the `cleaned_data` dictionary for the given field.
4. You *must* always end the `clean()` method by returning the reference to the `cleaned_data` dictionary. If you don't, you'll get some really frustrating errors!

This trivial example shows how we can clean the data being passed through the form before being stored. This is pretty handy, especially when particular fields need to have default values - or data within the form is missing, and we need to handle such data entry problems.

Note: Overriding methods that Django provides can provide you with an elegant way to add that extra bit of functionality for your application. There are many methods which you can safely override for your benefit, just like the `clean()` method in `ModelForm` as shown above. Check out the [Official Django Documentation on Models](#) for more examples on how you can override default functionality to slot your own in.

7.3 Exercises

Now that you've worked through the chapter, try these exercises to solidify your knowledge on Django's form functionality.

- What happens when you don't enter in a category name on the add category form?
- What happens when you try to add a category that already exists?
- What happens when you visit a category that does not exist?
- How could you gracefully handle when a user visits a category that does not exist?
- Undertake the [part four of the official Django Tutorial](#) if you have not done so already to reinforce what you have learnt here.

7.3.1 Creating an Add Pages View, Template and URL Mapping

A next logical step would be to allow users to add pages to a given category. To do this, repeat the same workflow above for Pages - create a new view (and URL mapping), a new template, the URL mapping and then a link from the category page. To get you started, here's the view logic for you.

```
from rango.forms import PageForm

def add_page(request, category_name_url):
    context = RequestContext(request)

    category_name = decode_url(category_name_url)
    if request.method == 'POST':
        form = PageForm(request.POST)

        if form.is_valid():
            # This time we cannot commit straight away.
```

```
# Not all fields are automatically populated!
page = form.save(commit=False)

# Retrieve the associated Category object so we can add it.
cat = Category.objects.get(name=category_name)
page.category = cat

# Also, create a default value for the number of views.
page.views = 0

# With this, we can then save our new model instance.
page.save()

# Now that the page is saved, display the category instead.
return category(request, category_name_url)
else:
    print form.errors
else:
    form = PageForm()

return render_to_response('rango/add_page.html',
    {'category_name_url': category_name_url,
     'category_name': category_name, 'form': form},
    context)
```

7.3.2 Hints

To help you with the exercises above, the following hints may be of some use to you.

- Update the `category()` view to pass `category_name_url` by inserting it to the view's `context_dict` dictionary.
- Update the `category.html` with a link to `/rango/category/<category_name_url>/add_page/`.
- Update `rango/urls.py` with a URL mapping to handle the above link.
- In the sample code above, we make use of the `decode_url()` function created in Chapter 6. If you haven't created this function, you'll need to do so now.

User Authentication

The aim of this next part of the tutorial is to get you familiar with the user authentication mechanisms provided by Django. We'll be using the `auth` application provided as part of a standard Django installation in package `django.contrib.auth`. According to [Django's official documentation on Authentication](#), the application consists of the following aspects.

- *Users*.
- *Permissions*: a series of binary flags (e.g. yes/no) determining what a user may or may not do.
- *Groups*: a method of applying permissions to more than one user.
- A configurable *password hashing system*: a must for ensuring data security.
- *Forms and view tools for logging in users*, or restricting content.
- A pluggable *backend system*, allowing you to provide your own authentication-related functionality.

There's lots that Django can do for you in the area of user authentication. We'll be covering the basics to get you started. This'll help you build your confidence with the available tools and their underlying concepts.

8.1 Setting up Authentication

Before you can begin to play around with Django's authentication offering, you'll need to make sure that the relevant settings are present in your Rango project's `settings.py` file.

Within the `settings.py` file find the `INSTALLED_APPS` tuple and check that `django.contrib.auth` and `django.contrib.contenttypes` are listed, so that it looks like the code below:

```
INSTALLED_APPS = (
    'django.contrib.auth', # THIS LINE SHOULD BE PRESENT AND UNCOMMENTED
    'django.contrib.contenttypes', # THIS LINE SHOULD BE PRESENT AND UNCOMMENTED
    'django.contrib.sessions',
    'django.contrib.sites',
    'django.contrib.messages',
    'django.contrib.staticfiles',
    # Uncomment the next line to enable the admin:
    'django.contrib.admin',
    # Uncomment the next line to enable admin documentation:
    # 'django.contrib.admindocs',
    'rango',
)
```

While `django.contrib.auth` provides Django with access to the authentication system, `django.contrib.contenttypes` is used by the authentication application to track models installed in your database. Check out the [Official Django documentation](#) for more details on what `django.contrib.contenttypes` is and does to make your life easier.

Note: Remember, if you had to add either one of the `auth` or `contenttypes` applications to your `INSTALLED_APPS` tuple, you will need to resynchronise your database with the `$ python manage.py syncdb` command.

Passwords are stored by default in Django using the [PBKDF2 algorithm](#), providing a good level of security for your user's data. You can read more about this as part of the [official Django documentation](#) on how Django stores passwords. The documentation also provides an explanation of how to use different password hashers if you require a greater level of security.

8.2 The User Model

The core of Django's authentication system is the `User` object, located at `django.contrib.auth.models.User`. A `User` object represents each of the people interacting with a Django application. The [Django documentation on User objects](#) states that they are used to allow aspects of the authentication system like access restriction, registration of new user profiles and the association of creators with site content.

The `User` model comes complete with five primary attributes. They are:

- the username for the user account;
- the account's password;
- the user's email address;
- the user's first name; and
- the user's surname.

The model also comes with other attributes such as `is_active` (which determines whether a particular account is active or not). Check the [official Django documentation](#) on the `User` model for a full list of attributes provided by the base `User` model.

8.3 Additional User Attributes

However, what if all the provided attributes that the `User` model provides isn't enough? For our Rango application, we want to include two more additional attributes for each user account. Specifically, we wish to include:

- a `URLField`, allowing a user of Rango to specify their own website; and
- a `ImageField`, which allows users to specify a picture for their user profile.

Fortunately, this is a relatively easy task to accomplish. This is achieved through the creation of an additional model in Rango's `models.py` file. Let's add the new model - add the following code.

```
class UserProfile(models.Model):  
    # This line is required. Links UserProfile to a User model instance.  
    user = models.OneToOneField(User)  
  
    # The additional attributes we wish to include.  
    website = models.URLField(blank=True)
```

```

picture = models.ImageField(upload_to='profile_images', blank=True)

# Override the __unicode__() method to return out something meaningful!
def __unicode__(self):
    return self.user.username

```

As we also reference the `User` model, we'll need to include the model into the `models.py` namespace. Add it with the following import statement at the top of the file.

```
from django.contrib.auth.models import User
```

So, how do we accomplish our goal of adding additional user profile fields? This isn't achieved through inheritance, instead the `UserProfile` model inherits from Django's `Model` class and is linked to the base `User` class through a one-to-one relationship via attribute `user`. This is because various applications may all want to use the `User` model and extend upon it in different ways.

For Rango, we've added two fields to complete our user profile, and provided a `__unicode__()` method to return a meaningful value when a unicode representation of a `UserProfile` model instance is requested.

For the two fields `website` and `picture`, we have set `blank=True` for both. This allows each of the fields to be blank if necessary, meaning that users need not supply values for the attributes if they do not wish to.

Note that the `ImageField` field has an `upload_to` attribute. The value of this attribute is conjoined with the project's `MEDIA_ROOT` setting to provide a path with which uploaded profile images will be stored. For example, a `MEDIA_ROOT` of `<workspace>/tango_with_django_project/media/` and `upload_to` attribute of `profile_images` will result in all profile images being stored in the directory `<workspace>/tango_with_django_project/media/profile_images/`.

Warning: The Django `ImageField` field makes use of the *Python Imaging Library (PIL)*. Back in Chapter 2, we discussed installing PIL along with Django to your setup. If you haven't got PIL installed, you'll need to install it now. If you don't, you'll be greeted with exceptions stating that the module `pil` cannot be found!

With our `UserProfile` model defined, we now edit Rango's `admin.py` file to include the new `UserProfile` model in the Django administration web interface. In the `admin.py` file, add the following line.

```
admin.site.register(UserProfile)
```

You also need to import the `UserProfile` model by adding one of the following lines at the top of the `admin.py` file. Choose which one you like - the first imports `UserProfile` with a separate import statement, while the second combines the import of `UserProfile` with Rango models that we have used previously in `admin.py`.

```

# Import the UserProfile model individually.
from rango.models import UserProfile

# Import the UserProfile model with Category and Page.
# If you choose this option, you'll want to modify the import statement you've already got to include
from rango.models import Category, Page, UserProfile

```

Note: Remember that your database must be synchronised with the creation of a new model. Run `$ python manage.py syncdb` from your terminal to synchronise the new `UserProfile` model. This process involves Django creating one or more underlying database tables for the given model. Forgetting to synchronise your changes will result in errors explaining that the required database tables cannot be found.

8.4 Creating a *User Registration View and Template*

With our authentication infrastructure laid out, we can now begin to build onto it by providing users of our application with the opportunity to create new user accounts. We will achieve this via the creation of a new view and template combination.

Note: We feel it's important to note that there are several off-the-shelf user registration packages available for you to download and use in your Django projects. Examples include the [Django Registration application](#), and you can also check out the table on [this webpage](#) which lists other registration packages. While these exist, we'll be showing you how to set up everything from scratch. While this is at odds with the DRY principle, it is also important to get a feeling for the user authentication package and feature. It will also re-enforce your understanding of working with forms, how to extend upon the user model, and how to upload media.

To set everything the user registration functionality will we go through the following steps:

1. Create a `UserForm` and `UserProfileForm`.
2. Add a view to handle the creation of a new user.
3. Create a template that displays the `UserForm` and `UserProfileForm`.
4. Map a URL to the view created.
5. Link the index page to the register page

8.4.1 Creating the `UserForm` and `UserProfileForm`

In `rango/forms.py`, we now need to create two classes inheriting from `forms.ModelForm`. We'll be creating one for the base `User` class, as well as one for the new `UserProfile` model that we just created. The two `ModelForm` inheriting classes allow us to display a HTML form displaying the necessary form fields for a particular model, taking away a significant amount of work for us. Neat!

In `rango/forms.py`, let's create our two classes which inherit from `forms.ModelForm`. Add the following code to the module.

```
class UserForm(forms.ModelForm):  
    password = forms.CharField(widget=forms.PasswordInput())  
  
    class Meta:  
        model = User  
        fields = ('username', 'email', 'password')  
  
class UserProfileForm(forms.ModelForm):  
    class Meta:  
        model = UserProfile  
        fields = ('website', 'picture')
```

You'll notice that within both classes we create, we add a `nested` `Meta` class. As [the name of the nested class may suggest](#), anything within a nested `Meta` class describes additional properties about the particular `ModelForm` class it belongs to. Each `Meta` class must at a bare minimum supply a `model` field, which references back to the model the `ModelForm` inheriting class should relate to. Our `UserForm` class is therefore associated with the `User` model, for example. By default, Django then renders a HTML form for *all* fields within the associated model.

However, there may be scenarios where we would not want a user to provide information for *all* fields within the associated model. For example, certain form fields may need to be filled in automatically by your code - such as in the `UserProfileForm`. Recall that the `UserProfile` model contains a `user` attribute, providing a one-to-one relationship to the `User` model. We don't want users to see this abstraction - we want Rango to handle it for them!

With the `fields` attribute, we can fine tune what fields the user sees in a rendered form. `UserProfileForm` will therefore display entries for the `website` and `picture` fields, but will not provide anything for the `user` field.

You'll also notice that `UserForm` includes a definition of the `password` attribute. While a `User` model instance contains a `password` attribute by default, the rendered HTML form element is of the incorrect type. If a user types a password, the password will be visible. By updating the `password` attribute definition, we can then specify that the `CharField` instance should hide a user's input from prying eyes through use of the `PasswordInput()` widget.

You shouldn't forget to include the required classes at the top of the `forms.py` module!

```
from rango.models import UserProfile
from django.contrib.auth.models import User
```

8.4.2 Creating the `register()` View

Next we need to handle both the rendering of the form, and the processing of form input data. Within `Rango`'s `views.py` file, add the following view function:

```
from rango.forms import UserForm, UserProfileForm

def register(request):
    # Like before, get the request's context.
    context = RequestContext(request)

    # A boolean value for telling the template whether the registration was successful.
    # Set to False initially. Code changes value to True when registration succeeds.
    registered = False

    # If it's a HTTP POST, we're interested in processing form data.
    if request.method == 'POST':
        # Attempt to grab information from the raw form information.
        # Note that we make use of both UserForm and UserProfileForm.
        user_form = UserForm(data=request.POST)
        profile_form = UserProfileForm(data=request.POST)

        # If the two forms are valid...
        if user_form.is_valid() and profile_form.is_valid():
            # Save the user's form data to the database.
            user = user_form.save()

            # Now we hash the password with the set_password method.
            # Once hashed, we can update the user object.
            user.set_password(user.password)
            user.save()

            # Now sort out the UserProfile instance.
            # Since we need to set the user attribute ourselves, we set commit=False.
            # This delays saving the model until we're ready to avoid integrity problems.
            profile = profile_form.save(commit=False)
            profile.user = user

            # Did the user provide a profile picture?
            # If so, we need to get it from the input form and put it in the UserProfile model.
            if 'picture' in request.FILES:
                profile.picture = request.FILES['picture']

            # Now we save the UserProfile model instance.
            profile.save()

            registered = True

    # This needs to be processed after the above code.
    # It can't be done before because the profile_form depends on the user_form
    # which has not been processed yet.
    if registered:
        return HttpResponseRedirect('/rango/register/complete')
    else:
        return render(request, 'rango/register.html', {'form': user_form, 'profile_form': profile_form, 'registered': registered})
```

```
# Update our variable to tell the template registration was successful.
registered = True

# Invalid form or forms - mistakes or something else?
# Print problems to the terminal.
# They'll also be shown to the user.
else:
    print user_form.errors, profile_form.errors

# Not a HTTP POST, so we render our form using two ModelForm instances.
# These forms will be blank, ready for user input.
else:
    user_form = UserForm()
    profile_form = UserProfileForm()

# Render the template depending on the context.
return render_to_response(
    'rango/register.html',
    {'user_form': user_form, 'profile_form': profile_form, 'registered': registered},
    context)
```

Is the view a lot more complex? It might look so at first, but it isn't really. The only added complexity from our previous `add_category()` view is the need to handle two distinct `ModelForm` instances - one for the `User` model, and one for the `UserProfile` model. We also need to handle a user's profile image, if he or she chooses to upload one.

We also establish a link between the two model instances that we create. After creating a new `User` model instance, we reference it in the `UserProfile` instance with the line `profile.user = user`. This is where we populate the `user` attribute of the `UserProfileForm` form, which we hid from users in Section 8.4.1.

8.4.3 Creating the *Registration* Template

Now create a new template file, `rango/register.html` and add the following code:

```
<!DOCTYPE html>
<html>
    <head>
        <title>Rango</title>
    </head>

    <body>
        <h1>Register with Rango</h1>

        {%
            if registered %
        Rango says: <strong>thank you for registering!</strong>
        <a href="/rango/">Return to the homepage.</a><br />
        {%
            else %
        Rango says: <strong>register here!</strong><br />

        <form id="user_form" method="post" action="/rango/register/">
            enctype="multipart/form-data">

            {% csrf_token %}

            <!-- Display each form. The as_p method wraps each element in a paragraph
                (<p>) element. This ensures each element appears on a new line,
                making everything look neater. -->
```

```

{{ user_form.as_p }}
{{ profile_form.as_p }}

<!-- Provide a button to click to submit the form. -->
<input type="submit" name="submit" value="Register" />
</form>
{%
  % endif %}
</body>
</html>

```

This HTML template makes use of the `register` variable we used in our view indicating whether registration was successful or not. Note that `registered` must be `False` in order for the template to display the registration form - otherwise, apart from the title, only a success message is displayed.

Warning: You should be aware of the `enctype` attribute for the `<form>` element. When you want users to upload files from a form, it's an absolute *must* to set `enctype` to `multipart/form-data`. This attribute and value combination instructs your browser to send form data in a special way back to the server. Essentially, the data representing your file is split into a series of chunks and sent. For more information, check out [this great Stack Overflow answer](#). You should also remember to include the CSRF token, too. Ensure that you include `{% csrf_token %}` within your `<form>` element.

8.4.4 The `register()` View URL Mapping

Now we can add a URL mapping to our new view. In `rango/urls.py` modify the `urlpatterns` tuple as shown below:

```

urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^about/$', views.about, name='about'),
    url(r'^category/(?P<category_name_url>\w+)$', views.category, name='category'),
    url(r'^add_category/$', views.add_category, name='add_category'),
    url(r'^category/(?P<category_name_url>\w+)/add_page/$', views.add_page, name='add_page'),
    url(r'^register/$', views.register, name='register'), # ADD NEW PATTERN!
)

```

The newly added pattern points the URL `/rango/register/` to the `register()` view.

8.4.5 Linking Together

Finally, we can add a link pointing to that URL in our homepage `index.html` template. Underneath the link to the category addition page, add the following hyperlink.

```
<a href="/rango/register/">Register Here</a>
```

8.4.6 Demo

Easy! Now you'll have a new hyperlink with the text `Register Here` that'll take you to the registration page. Try it out now! Start your Django development server and try to register a new user account. Upload a profile image if you wish. Your registration form should look like the one illustrated in Figure 8.1.

Upon seeing the message indicating your details were successfully registered, the database should have two new entries in its tables corresponding to the `User` and `UserProfile` models.

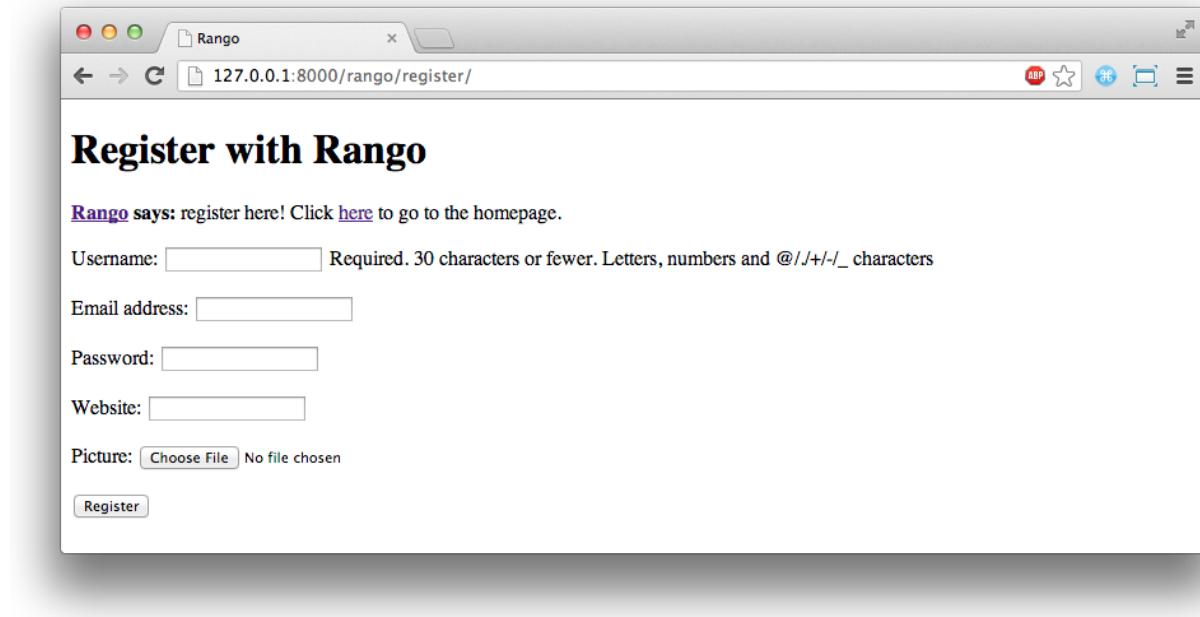


Fig. 8.1: A screenshot illustrating the basic registration form you create as part of this tutorial.

8.5 Adding Login Functionality

With the ability to register accounts completed, we now need to add login functionality. To achieve this we will need to undertake the workflow below:

- Create a login view to handle user credentials
- Create a login template to display the login form
- Map the login view to a url
- Provide a link to login from the index page

8.5.1 Creating the `login()` View

In `rango/views.py` create a new function called `user_login()` and add the following code:

```
def user_login(request):  
    # Like before, obtain the context for the user's request.  
    context = RequestContext(request)  
  
    # If the request is a HTTP POST, try to pull out the relevant information.  
    if request.method == 'POST':  
        # Gather the username and password provided by the user.  
        # This information is obtained from the login form.  
        username = request.POST['username']  
        password = request.POST['password']  
  
        # Use Django's machinery to attempt to see if the username/password  
        # combination is valid - a User object is returned if it is.  
        user = authenticate(username=username, password=password)
```

```

# If we have a User object, the details are correct.
# If None (Python's way of representing the absence of a value), no user
# with matching credentials was found.
if user is not None:
    # Is the account active? It could have been disabled.
    if user.is_active:
        # If the account is valid and active, we can log the user in.
        # We'll send the user back to the homepage.
        login(request, user)
        return HttpResponseRedirect('/rango/')
    else:
        # An inactive account was used - no logging in!
        return HttpResponse("Your Rango account is disabled.")

else:
    # Bad login details were provided. So we can't log the user in.
    print "Invalid login details: {0}, {1}".format(username, password)
    return HttpResponse("Invalid login details supplied.")

# The request is not a HTTP POST, so display the login form.
# This scenario would most likely be a HTTP GET.
else:
    # No context variables to pass to the template system, hence the
    # blank dictionary object...
    return render_to_response('rango/login.html', {}, context)

```

This view may seem rather complicated as it has to handle a variety of situations. Like in previous examples, the `user_login()` view handles form rendering and processing.

First, if the view is accessed via the HTTP GET method, then the login form is displayed. However, if the form has been posted via the HTTP POST method, then we can handle processing the form.

If a valid form is sent, the username and password are extracted from the form. These details are then used to attempt to authenticate the user (with Django's `authenticate()` function). `authenticate()` then returns a `User` object if the username/password combination exists within the database - or `None` if no match was found.

If we retrieve a `User` object, we can then check if the account is active or inactive - and return the appropriate response to the client's browser.

However, if an invalid form is sent, because the user did not add both a username and password the login form is presented back to the user will form error messages (i.e. `username/password` is missing).

Of particular interest in the code sample above is the use of the built-in Django machinery to help with the authentication process. Note the use of the `authenticate()` function to check whether the username and password provided match to a valid user account, and the `login()` function to signify to Django that the user is to be logged in.

You'll also notice that we make use of a new class, `HttpResponseRedirect`. As the name may suggest to you, the response generated by an instance of the `HttpResponseRedirect` class tells the client's browser to redirect to the URL you provide as the argument. Note that this will return a HTTP status code of 302, which denotes a redirect, as opposed to an status code of 200 i.e. OK. See the [official Django documentation on Redirection](#), to learn more.

All of these functions and classes are provided by Django, and as such you'll need to import them, so add the following imports to `rango/views.py`:

```

from django.contrib.auth import authenticate, login
from django.http import HttpResponseRedirect, HttpResponse

```

8.5.2 Creating a *Login* Template

With our new view created, we'll need to create a new template allowing users to login. While we know that the template will live in the `templates/rango/` directory, we'll leave you to figure out the name of the file. Look at the code example above to work out the name. In your new template file, add the following code:

```
<!DOCTYPE html>
<html>
    <head>
        <!-- Is anyone getting tired of repeatedly entering the header over and over?? -->
        <title>Rango</title>
    </head>

    <body>
        <h1>Login to Rango</h1>

        <form id="login_form" method="post" action="/rango/login/">
            {% csrf_token %}
            Username: <input type="text" name="username" value="" size="50" />
            <br />
            Password: <input type="password" name="password" value="" size="50" />
            <br />

            <input type="submit" value="submit" />
        </form>

    </body>
</html>
```

Ensure that you match up the input name attributes to those that you specified in the `user_login()` view - i.e. `username` for the username, and `password` for password. Don't forget the `{% csrf_token %}`, either!

8.5.3 Mapping the Login View to a URL

With your login template created, we can now match up the `user_login()` view to a URL. Modify Rango's `urls.py` file so that its `urlpatterns` tuple now looks like the code below:

```
urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^about/$', views.about, name='about'),
    url(r'^category/(?P<category_name_url>\w+)$', views.category, name='category'),
    url(r'^add_category/$', views.add_category, name='add_category'),
    url(r'^category/(?P<category_name_url>\w+)/add_page/$', views.add_page, name='add_page'),
    url(r'^register/$', views.register, name='register'),
    url(r'^login/$', views.user_login, name='login'),
)
```

8.5.4 Linking Together

Our final step is to provide users of Rango with a handy link to access the login page. To do this, we'll edit the `index.html` template inside of the `templates/rango/` directory. Find the previously created category addition and registration links, and add the following hyperlink underneath. You may wish to include a line break (`
`) before the link.

```
<a href="/rango/login/">Login</a>
```

If you like, you can also modify the header of the index page to provide a personalised message if a user is logged in, and a more generic message if the user isn't. Within the `index.html` template, find the header, as shown in the code snippet below.

```
<h1>Rango says..hello world!</h1>
```

Replace this header with the following markup and Django template code. Note that we make use of the `user` object, which is available to Django's template system via the context. We can tell from this object if the user is logged in (authenticated). If he or she is logged in, we can also obtain details about him or her.

```
{% if user.is_authenticated %}
<h1>Rango says... hello {{ user.username }}!</h1>
{% else %}
<h1>Rango says... hello world!</h1>
{% endif %}
```

As you can see we have used Django's Template Language to check if the user is authenticated with `{% if user.is_authenticated %}`. The context variable which we pass through to the template will include a `user` variable if the user is logged in - so we can check whether they are authenticated or not. If so they will receive a personalised greeting in the header, i.e. Rango says... hello leifos!. Otherwise, the generic Rango says... hello world! header is displayed.

8.5.5 Demo

Check out Figure 8.2 for screenshots of what everything should look like.

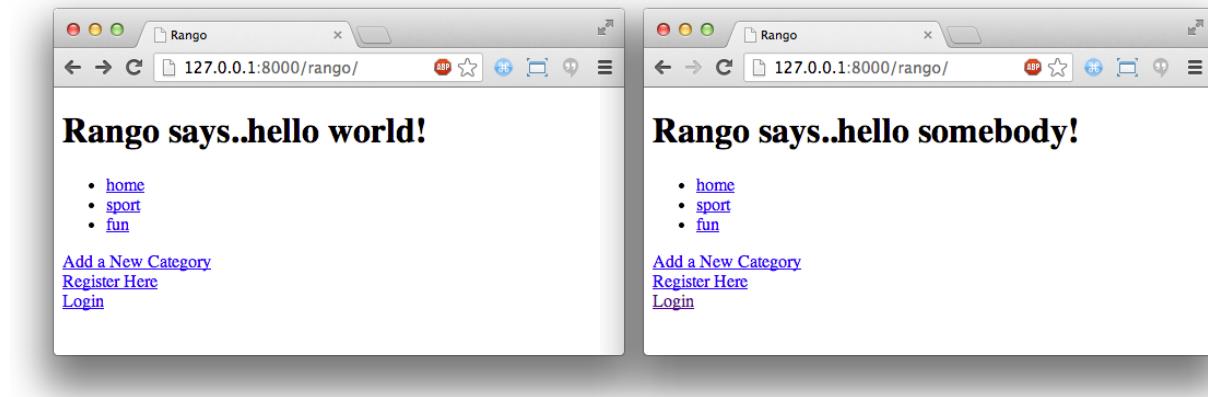


Fig. 8.2: Screenshots illustrating the header users receive when not logged in, and logged in with username somebody.

With this completed, user logins should now be completed! To test everything out, try starting Django's development server and attempt to register a new account. After successful registration, you should then be able to login with the details you just provided.

8.6 Restricting Access

Now that users can login to Rango, we can now go about restricting access to particular parts of the application as per the specification, i.e. that only registered users can add categories and pages. With Django, there are two ways in which we can achieve this goal:

- directly, by examining the `request` object and check if the user is authenticated, or,
- using a convenience *decorator* function that check if the user is authenticated.

The direct approach checks to see whether a user is logged in, via the `user.is_authenticated()` method. The `user` object is available via the `request` object passed into a view. The following example demonstrates this approach.

```
def some_view(request):
    if not request.user.is_authenticated():
        return HttpResponseRedirect("You are logged in.")
    else:
        return HttpResponseRedirect("You are not logged in.")
```

The second approach uses [Python decorators](#). Decorators are named after a [software design pattern by the same name](#). They can dynamically alter the functionality of a function, method or class without having to directly edit the source code of the given function, method or class.

Django provides decorator called, `login_required()` which we can attach to any view, where we require the user to be logged in. If a user is not logged in and they try to access a page which calls that view, then the user is redirected to another page which you can set, typically the login page.

8.6.1 Restricting Access with a Decorator

To try this out, create a view in Rango's `views.py` file, called `restricted()` and add the following code:

```
@login_required
def restricted(request):
    return HttpResponseRedirect("Since you're logged in, you can see this text!")
```

Note that to use a decorator, you place it *directly above* the function signature, and put a @ before naming the decorator. Python will execute the decorator before executing the code of your function/method. To use the decorator you will have to import it, so also add the following import:

```
from django.contrib.auth.decorators import login_required
```

We'll also add in another pattern to Rango's `urlpatterns` tuple in the `urls.py` file. Our tuple should then look something like the following example. Note the inclusion of mapping of the `views.restricted` view - this is the mapping you need to add.

```
urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^add_category/$', views.add_category, name='add_category'),
    url(r'^register/$', views.register, name='register'),
    url(r'^login/$', views.user_login, name='login'),
    url(r'^^(?P<category_name_url>\w+)$', views.category, name='category'),
    url(r'^restricted/$', views.restricted, name='restricted'),
)
```

We'll also need to handle the scenario where a user attempts to access the `restricted()` view, but is not logged in. What do we do with the user? The simplest approach is to redirect his or her browser. Django allows us to specify this in our project's `settings.py` file, located in the project configuration directory. In `settings.py`, define the variable `LOGIN_URL` with the URL you'd like to redirect users to that aren't logged in, i.e. the login page located at `/rango/login/`:

```
LOGIN_URL = '/rango/login/'
```

This ensures that the `login_required()` decorator will redirect any user not logged in to the URL `/rango/login/`.

8.7 Logging Out

To enable users to log out gracefully it would be nice to provide a logout option to users. Django comes with a handy `logout()` function that takes care of ensuring that the user is logged out, that their session is ended, and that if they subsequently try to access a view, that it will deny them access.

To provide log out functionality in `rango/views.py` add the a view called `user_logout()` with the following code:

```
from django.contrib.auth import logout

# Use the login_required() decorator to ensure only those logged in can access the view.
@login_required
def user_logout(request):
    # Since we know the user is logged in, we can now just log them out.
    logout(request)

    # Take the user back to the homepage.
    return HttpResponseRedirect('/rango/')
```

Note: Where's `RequestContext()`? In this simple user logout view, there's no need to obtain the request's context from Django's backend. If we don't need it, why ask for it?

With the view created, map the URL `/rango/logout/` to the `user_logout()` view by modifying the `urlpatterns` tuple in Rango's `urls.py` module:

```
urlpatterns = patterns('',
    url(r'^$', views.index, name='index'),
    url(r'^about/$', views.about, name='about'),
    url(r'^category/(?P<category_name_url>\w+)/$', views.category, name='category'),
    url(r'^add_category/$', views.add_category, name='add_category'),
    url(r'^category/(?P<category_name_url>\w+)/add_page/$', views.add_page, name='add_page'),
    url(r'^register/$', views.register, name='register'),
    url(r'^login/$', views.user_login, name='login'),
    url(r'^restricted/$', views.restricted, name='restricted'),
    url(r'^logout/$', views.user_logout, name='logout'),
)
```

Now that all the machinery for logging a user out has been completed, it'd be handy to provide a link from the homepage to allow users to simply click a link to logout. However, let's be smart about this: is there any point providing the logout link to a user who isn't logged in? Perhaps not - it may be more beneficial for a user who isn't logged in to be given the chance to register, for example.

Like in the previous section, we'll be modifying Rango's `index.html` template, and making use of the `user` object in the template's context to determine what links we want to show. Find your growing list of links at the bottom of the page and replace it with the following HTML markup and Django template code. Note we also add a link to our restricted page at `/rango/restricted/`.

```
{% if user.is_authenticated %}
<a href="/rango/restricted/">Restricted Page</a><br />
<a href="/rango/logout/">Logout</a><br />
{% else %}
<a href="/rango/register/">Register Here</a><br />
<a href="/rango/login/">Login</a><br />
{% endif %}

<a href="/rango/about/">About</a><br/>
<a href="/rango/add_category/">Add a New Category</a><br />
```

Simple - when a user is authenticated and logged in, he or she can see the Restricted Page and Logout links. If he or she isn't logged in, Register Here and Login are presented. As About and Add a New Category are not within the template conditional blocks, these links are available to both anonymous and logged in users.

8.8 Exercises

This chapter has covered several important aspects of managing user authentication within Django. We've covered the basics of installing Django's `django.contrib.auth` application into our project. Additionally, we have also shown how to implement a user profile model that can provide additional fields to the base `django.contrib.auth.models.User` model. We have also detailed how to setup the functionality to allow user registrations, login, logout, and to control access. For more information about user authentication and registration consult [Django's official documentation on Authentication](#).

- Customise the application so that only registered users can add/edit, while non-registered can only view/use the categories/pages. You'll also have ensure links to add/edit pages appear only if the user browsing the website is logged in.
- Provide informative error messages when users incorrectly enter their username or password.

In most applications you are going to require different levels of security when registering and managing users - for example, making sure the user enters an email address that they have access to, or sending users passwords that they have forgotten. While we could extend the current approach and build all the necessary infrastructure to support such functionality a `django-registration` application has been developed which greatly simplifies the process - visit <https://django-registration.readthedocs.org/en/latest/> to find out more about using this package.

Working with Templates

So far we've created several Django HTML templates for different pages in the application. You've probably already noticed that there is a lot of repeated HTML code in these templates.

While most sites will have lots of repeated structure (i.e. headers, sidebars, footers, etc) repeating the HTML in each template is a not good way to handle this. So instead of doing the same cut and paste hack job, we can minimize the amount of repetition in our code base by employing *template inheritance* provided by Django's Template Language.

The basic approach to using inheritance in templates is as follows.

1. Identify the re-occurring parts of each page that are repeated across your application (i.e. header bar, sidebar, footer, content pane)
2. In a *base template*, provide the skeleton structure of a standard page along with any common content (i.e. the copyright notice that goes in the footer, the logo and title that appears in the section), and then define a number of *blocks* which are subject to change depending on which page the user is viewing.
3. Create specific templates - all of which inherit from the base template - and specify the contents of each block.

9.1 Reoccurring HTML and The Base Template

Given the templates that we have created so far it should be pretty obvious that we have been repeating a fair bit of HTML code. Below we have abstracted away any page specific details to show the skeleton structure that we have been repeating within each template.

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango</title>
  </head>

  <body>
    <!-- Page specific content goes here -->
  </body>
</html>
```

Let's make this our base template, for the time being, and save it as `base.html` in Rango's `templates` directory (e.g. `templates/rango/base.html`).

Note: You should always aim to extract as much reoccurring content for your base templates. While it may be a bit more of a challenge for you to do initially, the time you will save in maintenance of your templates in the future

far outweighs the initial overhead. Think about it: would you rather maintain one copy of your markup or multiple copies?

Warning: Remember that your page `<!DOCTYPE html>` declaration absolutely must be placed on the first line for your page! Not doing so will mean your markup will not comply with the W3C HTML5 guidelines.

9.2 Template Blocks

Now that we've identified our base template, we can prepare it for our inheriting templates. To do this, we need to include a Template Tag to indicate what can be overridden in the base template - this is done through the use of *blocks*.

Add a `body_block` to the base template as follows:

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango</title>
  </head>

  <body>
    {%
      block body_block %} {%
      endblock %}
  </body>
</html>
```

Recall that standard Django template commands are denoted by `{%` and `%}` tags. To start a block, the command is `block <NAME>`, where `<NAME>` is the name of the block you wish to create. You must also ensure that you close the block with the `endblock` command, again enclosed within Django template tags.

You can also specify ‘default content’ for your blocks, if you so desire. Our `body_block` defined above presently has no default content associated with it. This means that if no inheriting template were to employ the use of `body_block`, nothing would be rendered - as shown in the code snippet below.

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango</title>
  </head>

  <body>

  </body>
</html>
```

However, we can overcome this by placing default content within the block definition, like so:

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango</title>
  </head>

  <body>
    {%
      block body_block %}This is body_block's default content.{%
      endblock %}
```

```
</body>
</html>
```

If a template were to inherit from the base template without employing the use of `body_block`, the rendered outcome would now look something like the markup shown below.

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango</title>
  </head>

  <body>
    This is body_block's default content.
  </body>
</html>
```

Hopefully this all makes sense - and for now, we'll be leaving `body_block` blank by default. All of our inheriting templates will be making use of `body_block`. You can place as many blocks in your templates as you so desire. For example, you could create a block for the page title, meaning you can alter the title of each page while still inheriting from the same base template.

Blocks are a really powerful feature of Django's template system to learn more about them check out the [official Django documentation on templates](#).

9.2.1 Abstracting Further

Now that you have an understanding of Django blocks, let's take the opportunity to abstract our base template a little bit further. Reopen the `base.html` template and modify it to look like the following.

```
<!DOCTYPE html>

<html>
  <head>
    <title>Rango - {% block title %}How to Tango with Django!{% endblock %}</title>
  </head>

  <body>
    <div>
      {% block body_block %}{% endblock %}
    </div>

    <hr />

    <div>
      <ul>
        {% if user.is_authenticated %}
          <li><a href="/rango/restricted/">Restricted Page</a></li>
          <li><a href="/rango/logout/">Logout</a></li>
          <li><a href="/rango/add_category/">Add a New Category</a></li>
        {% else %}
          <li><a href="/rango/register/">Register Here</a></li>
          <li><a href="/rango/login/">Login</a></li>
        {% endif %}
      </ul>
    </div>

    <li><a href="/rango/about/">About</a></li>
```

```
</ul>
</div>
</body>
</html>
```

We introduce two new features into the template.

- The first is a new Django template block, `title`. This will allow us to specify a custom page title for each page inheriting from our base template. If an inheriting page does not make use of this feature, the title is defaulted to Rango – How to Tango with Django!
- We also bring across the list of links from our current `index.html` template and place them into a HTML `<div>` tag underneath our `body_block` block. This will ensure the links are present across all pages inheriting from the base template. The links are preceded by a *horizontal rule* (`<hr />`) which provides a visual separation between the `body_block` content and the links.

Also note that we enclose the `body_block` within a HTML `<div>` tag - we'll be explaining the meaning of the `<div>` tag in Chapter 19. Our links are also converted to an unordered HTML list through use of the `` and `` tags.

9.3 Template Inheritance

Now that we've created a base template with a block, we can now update the templates we have created to inherit from the base template. For example, let's refactor the template `rango/category.html`.

To do this, first remove all the repeated HTML code leaving only the HTML and Template Tags/Commands specific to the page. Then at the beginning of the template add the following line of code:

```
{% extends 'rango/base.html' %}
```

The `extends` command takes one parameter, the template which is to be extended/inherited from (i.e. `rango/base.html`). We can then modify the `category.html` template so it looks like the following complete example.

Note: The parameter you supply to the `extends` command should be relative from your project's templates directory. For example, all templates we use for Rango should extend from `rango/base.html`, not `base.html`.

```
{% extends 'rango/base.html' %}

{% block title %}{{ category_name }}{% endblock %}

{% block body_block %}
<h1>{{ category_name }}</h1>

{% if pages %}
<ul>
    {% for page in pages %}
        <li><a href="{{ page.url }}">{{ page.title }}</a></li>
    {% endfor %}
</ul>
{% else %}
    <strong>No pages currently in category.</strong>
{% endif %}

{% if user.is_authenticated %}
    <a href="/rango/category/{{ category_name_url }}/add_page/">Add a Page</a>
```

```
{% endif %}

{% endblock %}
```

Now that we inherit from `base.html`, all that exists within the `category.html` template is the `extends` command, the `title` block and the `body_block` block. You don't need a well-formatted HTML document because `base.html` provides all the groundwork for you. All you're doing is plugging in additional content to the base template to create the complete HTML document which is sent to the client's browser.

Note: Templates are very powerful and you can even create your own template tags. Here we have shown how we can minimise the repetition of structure HTML in our templates.

However, templates can also be used to minimise code within your application's views. For example, if you wanted to include the same database-driven content on each page of your application, you could construct a template that calls a specific view to handle the repeating portion of your webpages. This then saves you from having to call the Django ORM functions which gather the required data for the template in every view that renders it.

To learn more about the extensive functionality offered by Django's template language, check out the official [Django documentation on templates](#).

9.4 Exercises

Now that you've worked through this chapter, we've got several exercises for you to work through. After completing them, you'll be a Django templating pro.

- Update all other existing templates within Rango's repertoire to extend from the `rango/base.html` template. Follow the same process as we demonstrated above. Once completed, your templates should all inherit from `base.html`, as demonstrated in Figure 9.1. While you're at it, make sure you remove the links from our `index.html` template. We don't need them anymore! You can also remove the link to Rango's homepage within the `about.html` template.
- Convert the restricted page to use a template. Call the template `restricted.html`, and ensure that it too extends from our `base.html` template.
- Add another link to our growing link collection that allows users to navigate back to Rango's homepage from anywhere on the website.

Warning: Remember to add `{% load static %}` to the top of each template that makes use of static media. If you don't, you'll get an error! Django template modules must be imported individually for each template that requires them - *you can't make use of modules included in templates you extend from!*

Fig. 9.1: A class diagram demonstrating how your templates should inherit from `base.html`.

Cookies and Sessions

In this chapter, we will be going through *sessions* and *cookies*, both of which go hand in hand, and are of paramount importance in modern day web applications. In the previous chapter, the Django framework used sessions and cookies to handle the login and logout functionality (all behind the scenes). Here we will explore how to explicitly use cookies for other purposes.

10.1 Cookies, Cookies Everywhere!

If you are already comfortable with the ideas and concepts behind cookies, then skip straight to Section 10.2, if not read on.

Whenever a request to a website is made, the webserver returns the content of the requested page. In addition, one or more cookies may also be sent to the client, which are in turn stored in a persistent browser cache. When a user requests a new page from the same web server, any cookies that are matched to that server are sent with the request. The server can then interpret the cookies as part of the request's context and generate a response to suit.

The term *cookie* wasn't actually derived from the food that you eat, but from the term *magic cookie*, a packet of data a program receives and sends again unchanged. In 1994, MCI sent a request to *Netscape Communications* to implement a way of implementing persistence across HTTP requests. This was in response to their need to reliably store the contents of a user's virtual shopping basket for an e-commerce solution they were developing. Netscape programmer Lou Montulli took the concept of a magic cookie and applied it to web communications. You can find out more about [cookies and their history on Wikipedia](#). Of course, with such a great idea came a software patent - and you can read [US patent 5774670](#) that was submitted by Montulli himself.

As an example, you may login to a site with a particular username and password. When you have been authenticated, a cookie may be returned to your browser containing your username, indicating that you are now logged into the site. At every request, this information is passed back to the server where your login information is used to render the appropriate page - perhaps including your username in particular places on the page. Your session cannot last forever, however - cookies *have* to expire at some point in time - they cannot be of infinite length. A web application containing sensitive information may expire after only a few minutes of inactivity. A different web application with trivial information may expire half an hour after the last interaction - or even weeks into the future.

The passing of information in the form of cookies can open up potential security holes in your web application's design. This is why developers of web applications need to be extremely careful when using cookies - does the information you want to store as a cookie *really* need to be sent? In many cases, there are alternate - and more secure - solutions to the problem. Passing a user's credit card number on an e-commerce site as a cookie for example would be highly unwise. What if the user's computer is compromised? The cookie could be taken by a malicious program. From there, hackers would have his or her credit card number - all because your web application's design is fundamentally flawed. You'll however be glad to know that a majority of websites use cookies for application specific functionality.

Note: Because of the potentially sensitive nature of cookies, lawmakers have taken a particularly keen interest in

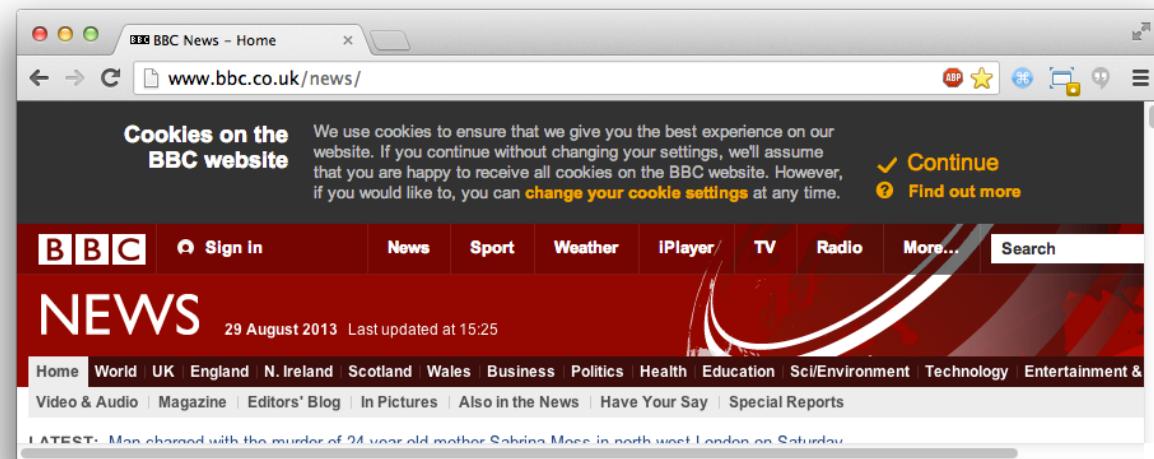


Fig. 10.1: A screenshot of the BBC News website (hosted in the United Kingdom) with the cookie warning message presented at the top of the page.

them. In particular, EU lawmakers in 2011 introduced an EU-wide ‘cookie law’, where all hosted sites within the EU should present a cookie warning message when a user visits the site for the first time. Check out Figure 10.1, demonstrating such a warning on the BBC News website. You can read about the law [here](#).

10.2 Sessions and the Stateless Protocol

All correspondence between web browsers (clients) and servers is achieved through the [HTTP protocol](#). As we *very briefly* touched upon in Chapter 7, HTTP is a [stateless protocol](#). This therefore means that a client computer running a web browser must establish a new network connection (a TCP connection) to the server each time a resource is requested (HTTP GET) or sent (HTTP POST)¹.

Without a persistent connection between client and server, the software on both ends cannot simply rely on connections alone to *hold session state*. For example, the client would need to tell the server each time who is logged on to the web application on a particular computer. This is known as a form of *dialogue* between the client and server, and is the basis of a *session* - a [semi-permanent exchange of information](#). Being a stateless protocol, HTTP makes holding session state pretty challenging (and frustrating) - but there are luckily several techniques we can use to circumnavigate this problem.

The most commonly used way of holding state is through the use of a *session ID* stored as a cookie on a client’s computer. A session ID can be considered as a token (a sequence of characters) to identify a unique session within a particular web application. Instead of storing all kinds of information as cookies on the client (such as usernames, names, passwords...), only the session ID is stored, which can then be mapped to a data structure on the web server. Within that data structure, you can store all of the information you require. This approach is a **much more secure** way to store information about users. This way, the information cannot be compromised by a insecure client or a connection which is being snooped.

If your browser supports cookies, pretty much all websites create a new session for you when you visit. You can see this for yourself now - check out Figure 10.2. In Google Chrome’s developer tools, you can view cookies which are

¹ The latest version of the HTTP standard HTTP 1.1 actually supports the ability for multiple requests to be sent in one TCP network connection. This provides huge improvements in performance, especially over high-latency network connections (such as via a traditional dial-up modem and satellite). This is referred to as *HTTP pipelining*, and you can read more about this technique on [Wikipedia](#).

sent by the web server you've accessed. In Figure 10.2, you can observe the selected cookie `sessionid`. The cookie contains a series of letters and numbers which Django uses to uniquely identify your session. From there, all your session details can be accessed - but only on the server side.

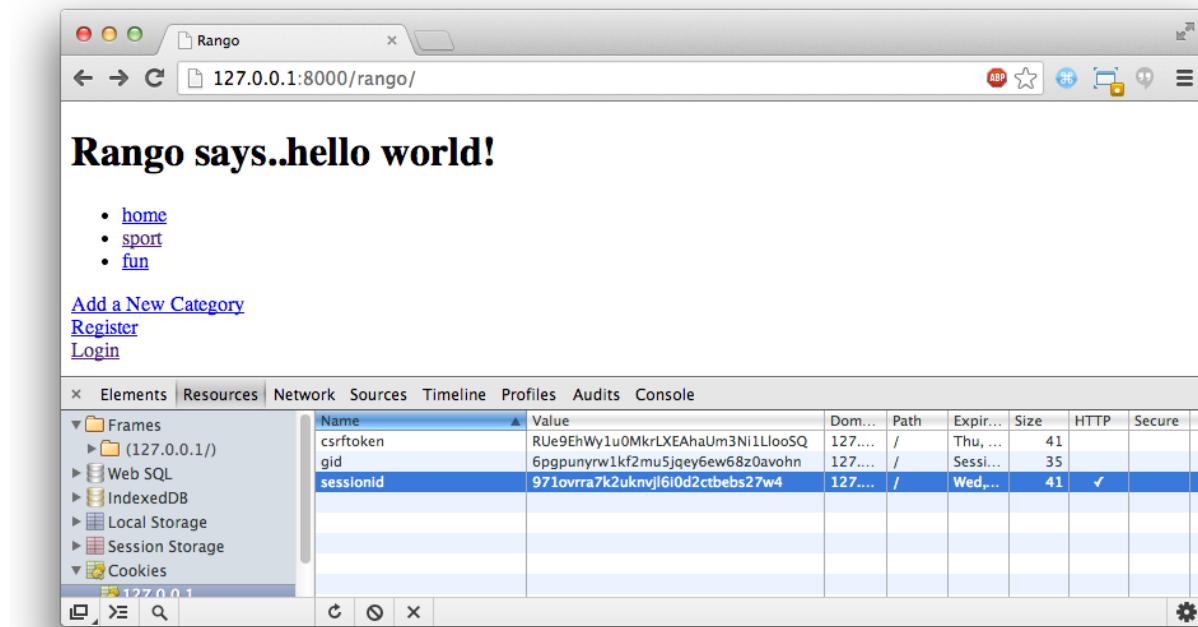


Fig. 10.2: A screenshot of Google Chrome with the Developer Tools opened - check out the cookie `sessionid`...

Session IDs don't have to be stored with cookies, either. Legacy PHP applications typically include them as a *querystring*, or part of the URL to a given resource. If you've ever come across a URL like `http://www.site.com/index.php?sessid=omgPhPwtfIsThisIdDoingHere332i942394`, that's probably uniquely identifying you to the server. Interesting stuff!

Note: Have a closer look at Figure 10.2. Do you notice the token `csrfmiddlewaretoken`? This cookie is to help prevent any cross-site forgery.

10.3 Setting up Sessions in Django

Although this should already be setup and working correctly, it's nevertheless good practice to learn which Django modules provide which functionality. In the case of sessions, Django provides `middleware` that implements session functionality.

To check that everything is in order, open your Django project's `settings.py` file. Within the file, locate the `MIDDLEWARE_CLASSES` tuple. You should find the `django.contrib.sessions.middleware.SessionMiddleware` module listed as a string in the tuple - if you don't, add it to the tuple now. It is the `SessionMiddleware` middleware which enables the creation of unique `sessionid` cookies.

The `SessionMiddleware` is designed to work flexibly with different ways to store session information. There are many approaches that can be taken - you could store everything in a file, in a database, or even in a cache. The most straightforward approach is to use the `django.contrib.sessions` application to store session information in a

Django model/database (specifically, the model `django.contrib.sessions.models.Session`). To use this approach, you'll also need to make sure that `django.contrib.sessions` is in the `INSTALLED_APPS` tuple of your Django project's `settings.py` file. If you add the application now, you'll need to synchronise your database using the `python manage.py syncdb` command to add the new tables to your database.

Note: If you are looking for lightning fast performance, you may want to consider a cached approach for storing session information. You can check out the [official Django documentation for advice on cached sessions](#).

10.4 A Cookie Tasting Session

We can now test out whether your browser supports cookies. While all modern web browsers do support cookies it is worthwhile checking your browser's settings regarding cookies. If you have your browser's security level set to a high level, certain cookies may get blocked. Look up your browser's documentation for more information, and enable cookies.

10.4.1 Testing Cookie Functionality

To test out cookies, you can make use of some convenience methods provided by Django's `request` object. The three of particular interest to us are `set_test_cookie()`, `test_cookie_worked()` and `delete_test_cookie()`. In one view, you will need to set a cookie. In another, you'll need to test that the cookie exists. Two different views are required for testing cookies because you need to wait to see if the client has accepted the cookie from the server.

We'll use two pre-existing views for this simple exercise, `index()` and `register()`. You'll need to make sure that you are logged out of Rango if you've implemented the user authentication functionality. Instead of displaying anything on the pages themselves, we'll be making use of the terminal output from the Django development server to verify whether cookies are working correctly. After we successfully determine that cookies are indeed working, we can remove the code we add to restore the two views to their previous state.

In Rango's `views.py` file, locate your `index()` view. Add the following line to the view. To ensure the line is executed, make sure you put it as the first line of the view, outside any conditional blocks.

```
request.session.set_test_cookie()
```

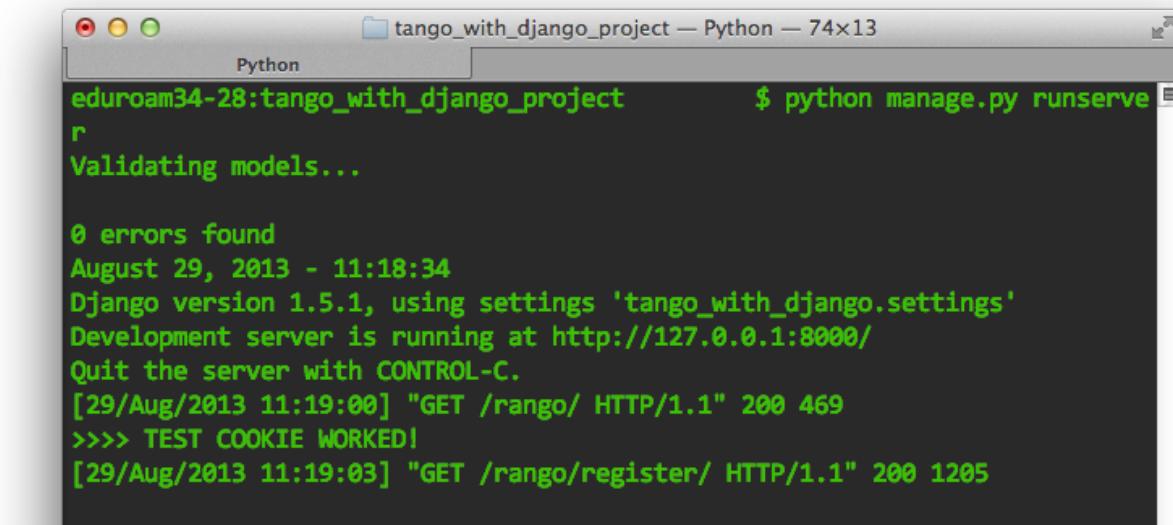
In the `register()` view, add the following three lines to the top of the function - again, to ensure that they are executed.

```
if request.session.test_cookie_worked():
    print ">>> TEST COOKIE WORKED!"
    request.session.delete_test_cookie()
```

With these small changes saved, run the Django development server and navigate to Rango's homepage, `http://127.0.0.1:8000/rango/`. Once the page is loaded, navigate to the registration page. When the registration page is loaded, you should see `>>> TEST COOKIE WORKED!` appear in your Django development server's console, like in Figure 10.3. If you do, everything works as intended!

If the message isn't displayed, you'll want to check your browser's security settings. The settings may be preventing the browser from accepting the cookie.

Note: You can delete the code you added in this section - we only used it to demonstrate cookies in action.



```

eduroam34-28:tango_with_django_project      $ python manage.py runserver
r
Validating models...

0 errors found
August 29, 2013 - 11:18:34
Django version 1.5.1, using settings 'tango_with_django.settings'
Development server is running at http://127.0.0.1:8000/
Quit the server with CONTROL-C.
[29/Aug/2013 11:19:00] "GET /rango/ HTTP/1.1" 200 469
>>> TEST COOKIE WORKED!
[29/Aug/2013 11:19:03] "GET /rango/register/ HTTP/1.1" 200 1205

```

Fig. 10.3: A screenshot of the Django development server's console output with the >>> TEST COOKIE WORKED! message.

10.5 Client Side Cookies: A Site Counter Example

Now we know cookies work, let's implement a very simple site visit counter. To achieve this, we're going to be creating two cookies: one to track the number of times the user has visited the Rango website, and the other to track the last time he or she accessed the site. Keeping track of the date and time of the last access will allow us to only increment the site counter once per day, for example.

The sensible place to assume a user enters the Rango site is at the index page. Open `rango/index.py` and edit the `index()` view as follows:

```

def index(request):
    context = RequestContext(request)

    category_list = Category.objects.all()
    context_dict = {'categories': category_list}

    for category in category_list:
        category.url = encode_url(category.name)

    page_list = Page.objects.order_by('-views')[:5]
    context_dict['pages'] = page_list

    ##### NEW CODE #####
    # Obtain our Response object early so we can add cookie information.
    response = render_to_response('rango/index.html', context_dict, context)

    # Get the number of visits to the site.
    # We use the COOKIES.get() function to obtain the visits cookie.
    # If the cookie exists, the value returned is casted to an integer.
    # If the cookie doesn't exist, we default to zero and cast that.
    visits = int(request.COOKIES.get('visits', '0'))

```

```
# Does the cookie last_visit exist?
if request.COOKIES.has_key('last_visit'):
    # Yes it does! Get the cookie's value.
    last_visit = request.COOKIES['last_visit']
    # Cast the value to a Python date/time object.
    last_visit_time = datetime.strptime(last_visit[:-7], "%Y-%m-%d %H:%M:%S")

    # If it's been more than a day since the last visit...
    if (datetime.now() - last_visit_time).days > 0:
        # ...reassign the value of the cookie to +1 of what it was before...
        response.set_cookie('visits', visits+1)
        # ...and update the last visit cookie, too.
        response.set_cookie('last_visit', datetime.now())
else:
    # Cookie last_visit doesn't exist, so create it to the current date/time.
    response.set_cookie('last_visit', datetime.now())

# Return response back to the user, updating any cookies that need changed.
return response
#### END NEW CODE ####
```

For reading through the code, you will see that a majority of the code deals with checking the current date and time. For this, you'll need to include Python's `datetime` module by adding the following import statement at the top of the `views.py` file.

```
from datetime import datetime
```

There's a `datetime` object within the `datetime` module, that's not a typo. Make sure you import the module correctly, otherwise you'll get frustrating import errors.

In the added code we check to see if the cookie `last_visit` exists. If it does, we can take the value from the cookie using the syntax `request.COOKIES['cookie_name']`, where `request` is the name of the `request` object, and '`cookie_name`' is the name of the cookie you wish to retrieve. **Note that all cookie values are returned as strings;** do not assume that a cookie storing whole numbers will return you an integer. You have to manually cast this to the correct type yourself. If a cookie does not exist, you can create a cookie with the `set_cookie()` method of the `response` object you create. The method takes in two values, the name of the cookie you wish to create (as a string), and the value of the cookie. In this case, it doesn't matter what type you pass as the value - it will be automatically cast as a string.

Now if you visit the Rango homepage, and inspect the developer tools provided by your browser, you should be able to see the cookies `visits` and `last_visit`. Figure 10.4 demonstrates the cookies in action.

10.6 Session Data

In the previous example, we used client side cookies. However, a more secure way to save session information is to store any such data on the server side. We can then use the session ID cookie which is stored on the client side (but is effectively anonymous) as the key to unlock the data.

To use session based cookies you need to perform the following steps.

1. Make sure that `MIDDLEWARE_CLASSES` in `settings.py` contains `django.contrib.sessions.middleware.SessionMiddleware`.
2. Configure your session backend. By default a database backend is assumed - so you will have make sure that your database is set up and synchronised. See the [official Django Documentation on Sessions](#) for other backend configurations.

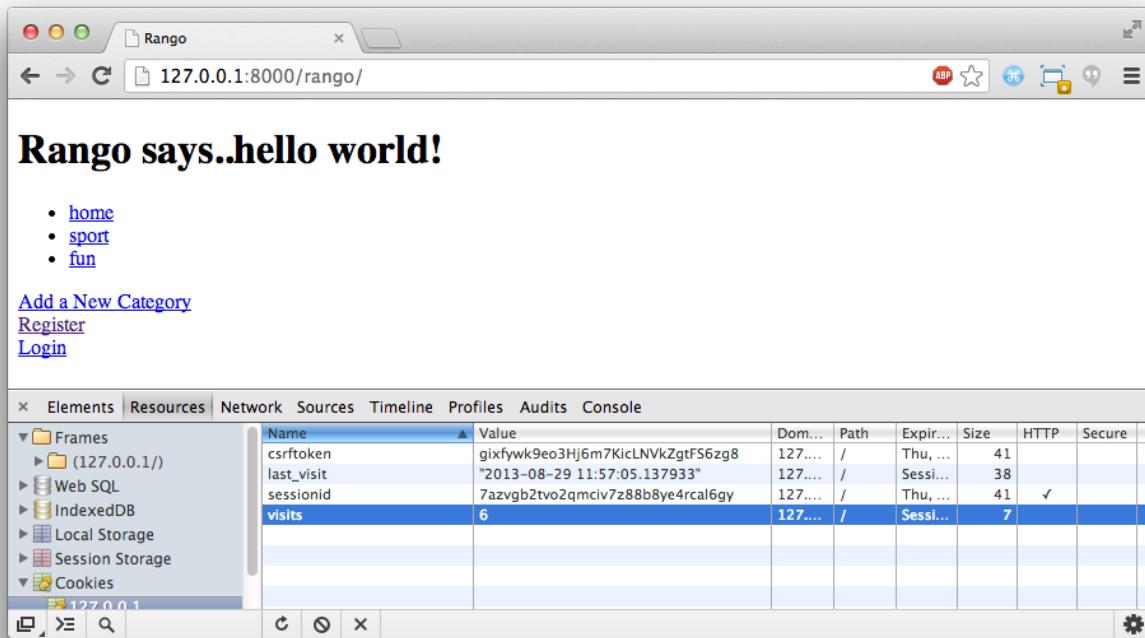


Fig. 10.4: A screenshot of Google Chrome with the Developer Tools open showing the cookies for Rango. Note the visits cookie - the user has visited a total of six times, with each visit at least one day apart.

Now if you want to check if the cookie has been stored you can do so by accessing the `request.session` object, where `request` is the name of your view's required parameter. Check out the modified `index()` function below to see how to do this.

```
def index(request):
    context = RequestContext(request)

    category_list = Category.objects.all()
    context_dict = {'categories': category_list}

    for category in category_list:
        category.url = encode_url(category.name)

    page_list = Page.objects.order_by('-views')[:5]
    context_dict['pages'] = page_list

    ##### NEW CODE #####
    if request.session.get('last_visit'):
        # The session has a value for the last visit
        last_visit_time = request.session.get('last_visit')
        visits = request.session.get('visits', 0)

        if (datetime.now() - datetime.strptime(last_visit_time[:-7], "%Y-%m-%d %H:%M:%S")).days > 0:
            request.session['visits'] = visits + 1
            request.session['last_visit'] = str(datetime.now())
    else:
        # The get returns None, and the session does not have a value for the last visit.
        request.session['last_visit'] = str(datetime.now())
        request.session['visits'] = 1
```

```
#### END NEW CODE ####

# Render and return the rendered response back to the user.
return render_to_response('rango/index.html', context_dict, context)
```

Note: A nice extra advantage to storing session data server-side is that you don't need to always cast data from strings to the desired type. Be careful though: this only seems to hold for simple data types such as strings, integers, floats and booleans.

10.7 Browser-Length and Persistent Sessions

When using cookies you can use Django's session framework to set cookies as either *browser-length sessions* or *persistent sessions*. As the names of the two types suggest:

- browser-length sessions expire when the user closes his or her browser; and
- persistent sessions can last over several browser instances - expiring at a time of your choice. This could be half an hour, or even as far as a month in the future.

By default, browser-length sessions are disabled. You can enable them by modifying your Django project's `settings.py` file. Add the variable `SESSION_EXPIRE_AT_BROWSER_CLOSE`, setting it to `True`.

Alternatively, persistent sessions are enabled by default, with `SESSION_EXPIRE_AT_BROWSER_CLOSE` either set to `False`, or not being present in your project's `settings.py` file. Persistent sessions have an additional setting, `SESSION_COOKIE_AGE`, which allows you to specify the age of which a cookie can live to. This value should be an integer, representing the number of seconds the cookie can live for. For example, specifying a value of `1209600` will mean your website's cookies expire after a two week period.

Check out the available settings you can use on the [official Django documentation on cookies](#) for more details. You can also check out [Eli Bendersky's blog](#) for an excellent tutorial on cookies and Django.

10.8 Basic Considerations and Workflow

When using cookies within your Django application, there's a few things you should consider:

- First, consider what type of cookies your web application requires. Does the information you wish to store need to persist over a series of user browser sessions, or can it be safely disregarded upon the end of one session?
- Think carefully about the information you wish to store using cookies. Remember, storing information in cookies by their definition means that the information will be stored on client's computers, too. This is a potentially huge security risk: you simply don't know how compromised a user's computer will be. Consider server-side alternatives if potentially sensitive information is involved.
- As a follow-up to the previous bullet point, remember that users may set their browser's security settings to a high level which could potentially block your cookies. As your cookies could be blocked, your site may function incorrectly. You *must* cater for this scenario - *you have no control over the client browser's setup*.

If client-side cookies are the right approach for you then work through the following steps:

1. You must first perform a check to see if the cookie you want exists. This can be done by checking the `request` parameter. The `request.COOKIEs.has_key('<cookie_name>')` function returns a boolean value indicating whether a cookie `<cookie_name>` exists on the client's computer or not.
2. If the cookie exists, you can then retrieve its value - again via the `request` parameter - with `request.COOKIEs[]`. The `COOKIES` attribute is exposed as a dictionary, so pass the name of the cookie

you wish to retrieve as a string between the square brackets. Remember, cookies are all returned as strings, regardless of what they contain. You must therefore be prepared to cast to the correct type.

3. If the cookie doesn't exist, or you wish to update the cookie, pass the value you wish to save to the response you generate. `response.set_cookie('<cookie_name>', value)` is the function you call, where two parameters are supplied: the name of the cookie, and the value you wish to set it to.

If you need more secure cookies, then use session based cookies:

1. Make sure that `MIDDLEWARE_CLASSES` in `settings.py` contains `'django.contrib.sessions.middleware.SessionMiddleware'`.
2. Configure your session backend `SESSION_ENGINE`. See the [official Django Documentation on Sessions](#) for the various backend configurations.
3. Check to see if the cookie exists via `requests.sessions.get()`
4. Update or set the cookie via the session dictionary, `requests.session['<cookie_name>']`

10.9 Exercises

Now you've read through this chapter and tried out the code, give these exercises a go.

- Change your cookies from client side to server side to make your application more secure. Clear the browser's cache and cookies, then check to make sure can't see the `last_visit` and `visits` variables in the browser. Note you will still see the `sessionid` cookie.
- Update the *About* page view and template telling the visitors how many times they have visited the site.

10.9.1 Hint

To aid you in your quest to complete the above exercises, the following hint may help you.

You'll have to pass the value from the cookie to the template context for it to be rendered as part of the page, as shown in the example below.

```
# If the visits session variable exists, take it and use it.
# If it doesn't, we haven't visited the site so set the count to zero.
if request.session.get('visits'):
    count = request.session.get('visits')
else:
    count = 0

# remember to include the visit data
return render_to_response('rango/about.html', {'visits': count}, context)
```

Bootstrapping Rango

In this chapter, we will be styling Rango using the *Twitter Bootstrap 2.3.2* toolkit. We won't go into the details about how Bootstrap works, and we will be assuming you have some familiarity with CSS. If you don't, check out the CSS chapter so that you understand the basics and then check out some Bootstrap tutorials. However, you should be able to go through this section and piece things together.

To get started take a look at the [Bootstrap 2.3.2 website](#) - it provides you with sample code and examples of the different components and how to style them by added in the appropriate style tags, etc.

On the Bootstrap website they provide a number of [example layouts](#) which we can base our design on, see <http://getbootstrap.com/2.3.2/getting-started.html#examples> for more information.

To style Rango we have identified that the [fluid style](#) more or less meets our needs in terms of the layout of Rango, i.e. it has a menu bar at the top, a side bar (which we will use to show categories) and a main content pane; see <http://getbootstrap.com/2.3.2/examples/fluid.html> for more information.

11.1 Setting up The Base Template

Before we can set up the base template to use this style we need to download Bootstrap, *JQuery* and the *Fluid Template*.

11.1.1 Download Bootstrap

Go to the [Bootstrap 2.3.2 website](#) and download the toolkit. When you unzip the files you will see that in the directory you have a `imgs`, `js` and `css` directory plus associated files. Copy the three subdirectories in Rango's `static` folder so that you will be able to reference them via the URL `/static/imgs`, `/static/js/` and `/static/css/`.

11.1.2 Download JQuery

Now go to the [JQuery website](#) and download the latest 2.x version of JQuery. Put the `js` file in to the `static/js/` directory.

11.1.3 Including CSS/JS in The Base Template

If you download or look at the source for <http://getbootstrap.com/2.3.2/examples/fluid.html>, you'll notice that in the `<head>` section there is some additional `<style>` code. Copy the CSS inside these style tags and create a new CSS file in `/static/css/` called `bootstrap-fluid-adj.css`. Your code should look similar to the example shown below.

```
body {
padding-top: 60px;
padding-bottom: 40px;
}

.sidebar-nav {
padding: 9px 0;
}

@media (max-width: 980px) {
/* Enable use of floated navbar text */
.navbar-text.pull-right {
float: none;
padding-left: 5px;
padding-right: 5px;
}
}
```

By adding it to a file we can minimise the code in our template. Now update the <head> section of `base.html` as follows.

```
{% load static %}
<!DOCTYPE html>
<html>
<head>
    <meta name="viewport" content="width=device-width, initial-scale=1.0">
    <!-- Bootstrap -->
    <link href="{% static 'css/bootstrap-fluid-adj.css' %}" rel="stylesheet">
    <link href="{% static 'css/bootstrap.min.css' %}" rel="stylesheet" media="screen">
    <link href="{% static 'css/bootstrap-responsive.css' %}" rel="stylesheet">

    <title>Rango - {% block title %}How to Tango with Django!{% endblock %}</title>
</head>
```

Note how we are including all these files by externally linking them in. We will also need to include the JQuery and Bootstrap JavaScript files. Instead of adding them to the <head> section, we will add them at the bottom of `base.html` just before we close the <body> tag, just like in the example below.

```
<script src="{% static 'js/jquery-2.0.3.min.js' %}"></script>
<script src="{% static 'js/bootstrap.min.js' %}"></script>
</body>
</html>
```

The reason to add them here is so that the page can load up faster. The browser concentrates on rendering the elements of the webpage, and once this has been completed, the JavaScript files are loaded.

11.1.4 Structuring the Base Template

If you take a close look at the fluid html source, you'll notice it has a lot of structure in it created by a series of <div> tags. Essentially the is broken into three parts - the top navigation bar, the main pane (houses the side bar and the main content pane), and a footer.

In the body of base put in the navigation bar code, just like in the example below.

```
<div class="navbar navbar-inverse navbar-fixed-top">
    <div class="navbar-inner">
        <div class="container">
            <button type="button" class="btn btn-navbar" data-toggle="collapse" data-target=".nav-collapse">
```

```

</span>
</span>
</span>
</button>
<a class="brand" href="/rango/">Rango</a>

<div class="nav-collapse collapse">
    <ul class="nav pull-right">
        {% if user.is_authenticated %}
            <li class="navbar-text">Welcome, {{ user.username }}!</li>
            <li><a href="/rango/logout/">Logout</a></li>
        {% else %}
            <li><a href="/rango/register/">Register</a></li>
            <li><a href="/rango/login/">Login</a></li>
        {% endif %}
    </ul>
    <ul class="nav">
        {% if user.is_authenticated %}
            <li><a href="/rango/restricted/">Restricted</a></li>
            <li><a href="/rango/add_category/">Add Category</a></li>
        {% endif %}
            <li><a href="/rango/about/">About</a></li>
        </ul>
    </div>
    <!-- .nav-collapse -->
</div>
</div>
</div>

```

After this, you can add in the next `<div>` which will house the side bar navigation and the main content pane:

```

<div class="container-fluid">
    <div class="row-fluid">
        <div class="span3">
            <div class="well sidebar-nav">
                <!-- Empty for the timebeing -->
            </div>
            <!-- .well -->
        </div>
        <!-- /span -->
        <div class="span9">
            {% block body_block %}
            {% endblock %}
        </div>
        <!-- /span -->
    </div>
    <!-- /row -->
</div>
<!-- /fluid-container -->

<hr>

```

You can see that we have included the `body_block` in here. And now finally, below this add in a footer:

```

<footer>
    <div class="container">
        <p>&copy; Rango: How to Tango with Django 2013</p>
    </div>
</footer>

```

11.2 Quick Style Change

Now that we have the `base.html` all set up and ready to go, we can do a really quick face light to Rango by adding `<div class="hero-unit">` around the contents within each `body_block` on each page. For example, convert the `body_block` of the `about.html` template to be:

```
{% block body_block %}  
<div class="hero-unit">  
    <h1>About Rango</h1>  
    This is <strong>Rango's about page</strong>. <br />  
  
    You've visited the site on <strong>{{ visit_count }} occasion(s)</strong>. <br />  
  
    Here's a picture of Rango! <br />  
      
</div>  
{% endblock %}
```

With all pages fitted with hero unit `<div>` tags Rango you should be looking pretty good. However, you will notice that some of the pages still look pretty ugly, especially the pages with forms (see the figure below). Later in this chapter we will update these forms.

Also, you'll probably have noticed the sidebar is empty. In the next chapter we will sort that out with some handy navigation links.

11.2.1 The Index Page

Since we have only encapsulated the content into a hero unit `<div>`, we haven't really capitalised on the classes and styling that Bootstrap gives us. So here we have taken the columns from the fluid page and used them to house the top categories and top pages. Since the original page has three columns, we have taken two and made them slightly better by adjusting the class so that the `` is 6 for each instead of 4, so we can update the `index.html` template to look like the following.

```
{% block body_block %}  
<div class="hero-unit">  
    <h1>Ready to Rango</h1>  
    <p>Find, Add, Share and Rango useful links and resources.</p>  
</div>  
  
<div class="row-fluid">  
    <div class="span6">  
        <h2>Top Five Categories</h2>  
        {% if categories %}  
            <ul>  
                {% for category in categories %}  
                    <li><a href="/rango/category/{{ category.url }}">{{ category.name }}</a></li>  
                {% endfor %}  
            </ul>  
        {% else %}  
            <strong>No categories at present.</strong>  
        {% endif %}  
  
    </div>  
    <!--/span-->  
    <div class="span6">  
        <h2>Top Five Pages</h2>  
        {% if pages %}
```

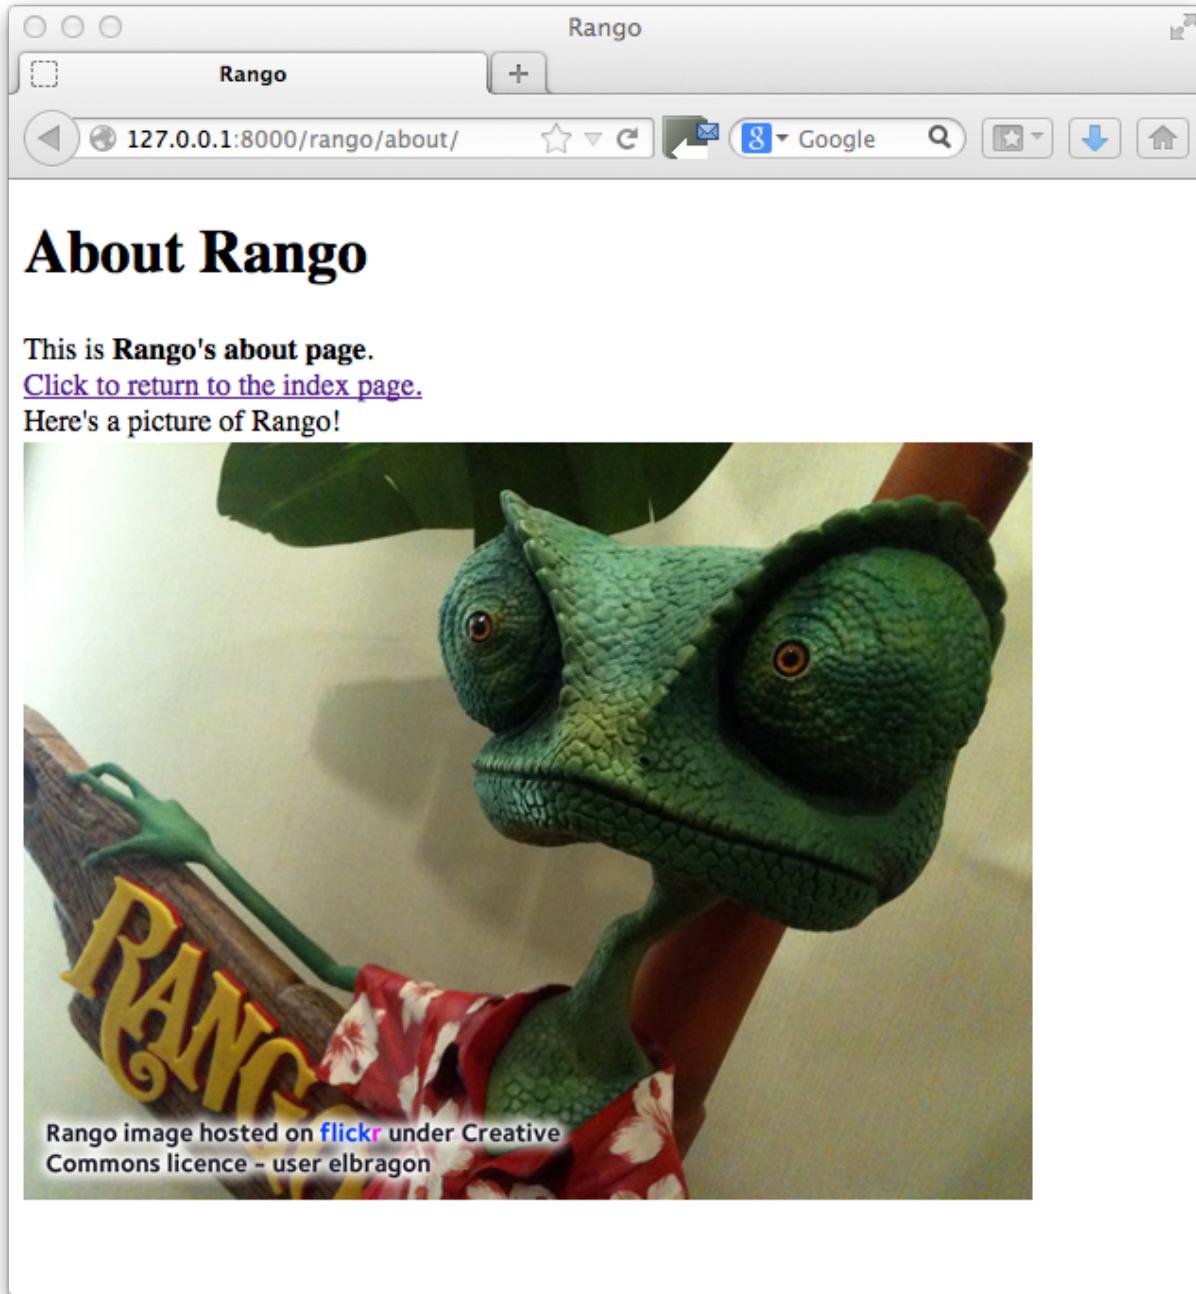


Fig. 11.1: A screenshot of the About page without style.

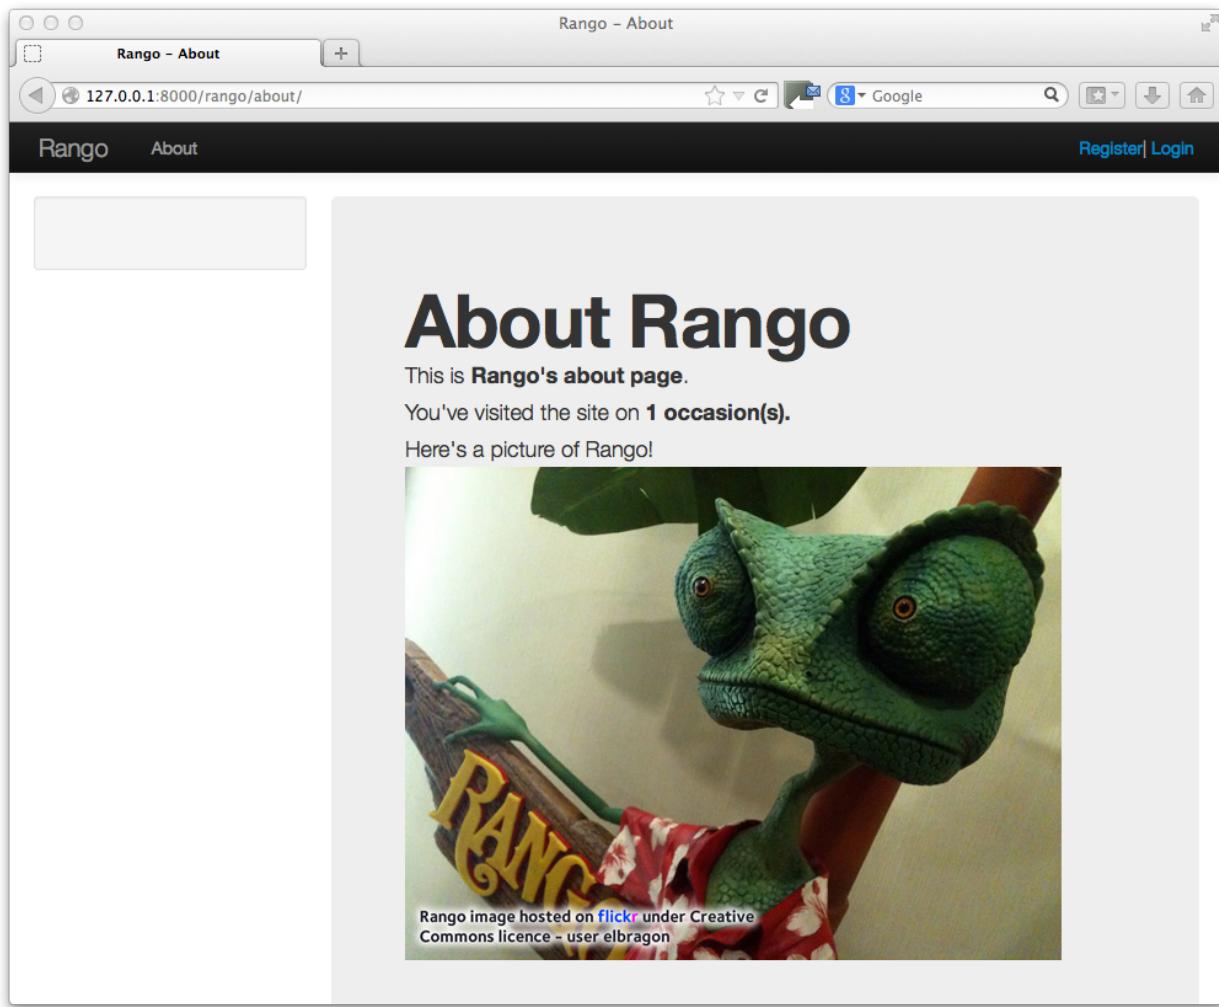


Fig. 11.2: A screenshot of the About page with Bootstrap Styling applied.

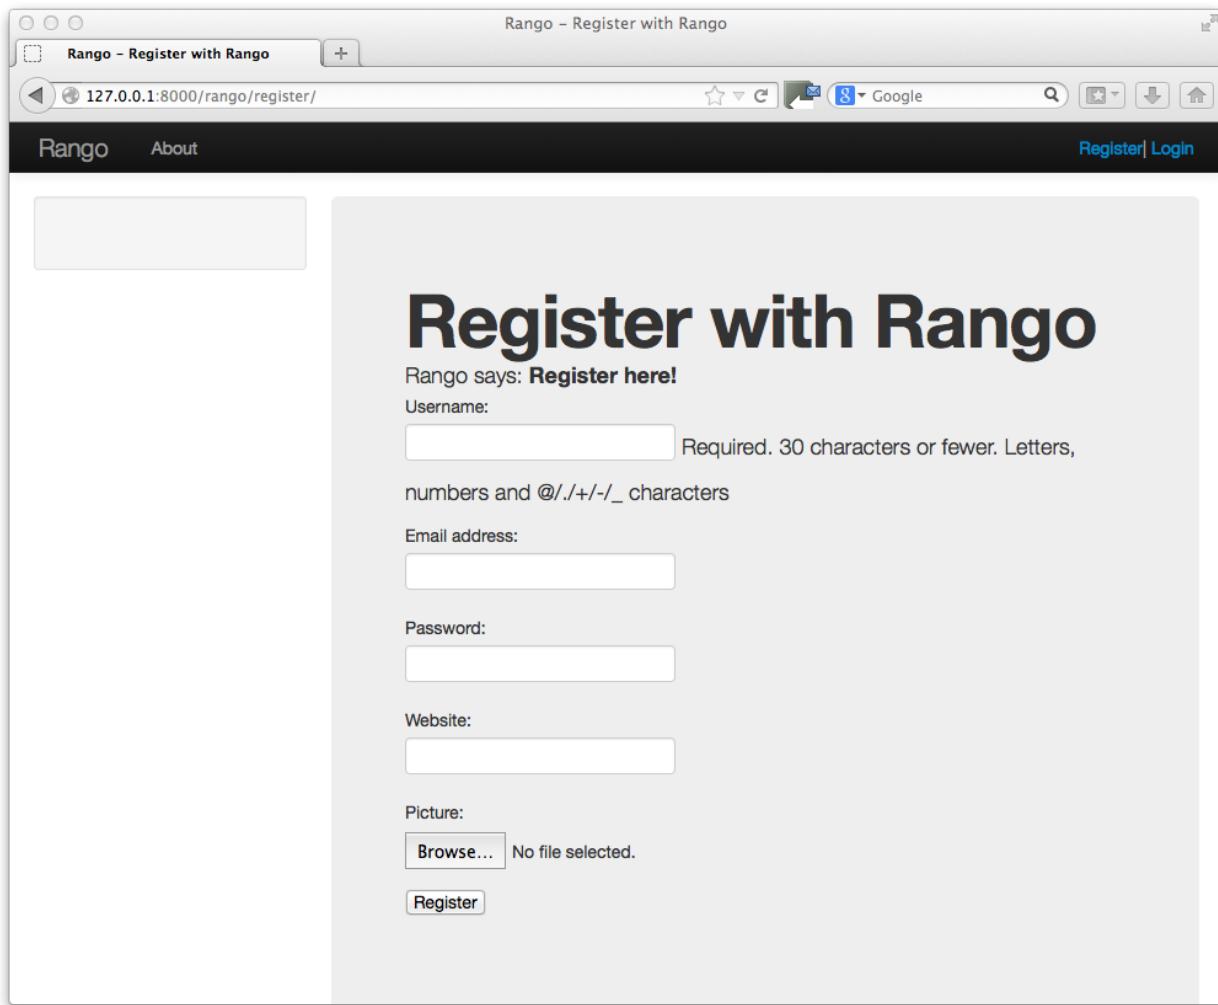


Fig. 11.3: A screenshot of the Registration page with Bootstrap Styling applied but not customised.

```
<ul>
    {% for page in pages %}
        <li><a href="{{ page.url }}">{{ page.title }}</a> - {{ page.category }} ({{ page.views }})
    {% endfor %}
</ul>
{% else %}
<strong>No pages at present.</strong>
{% endif %}

</div>
<!--/span-->
</div>
<!--/row-->
{% endblock %}
```

The page should look a lot better now.

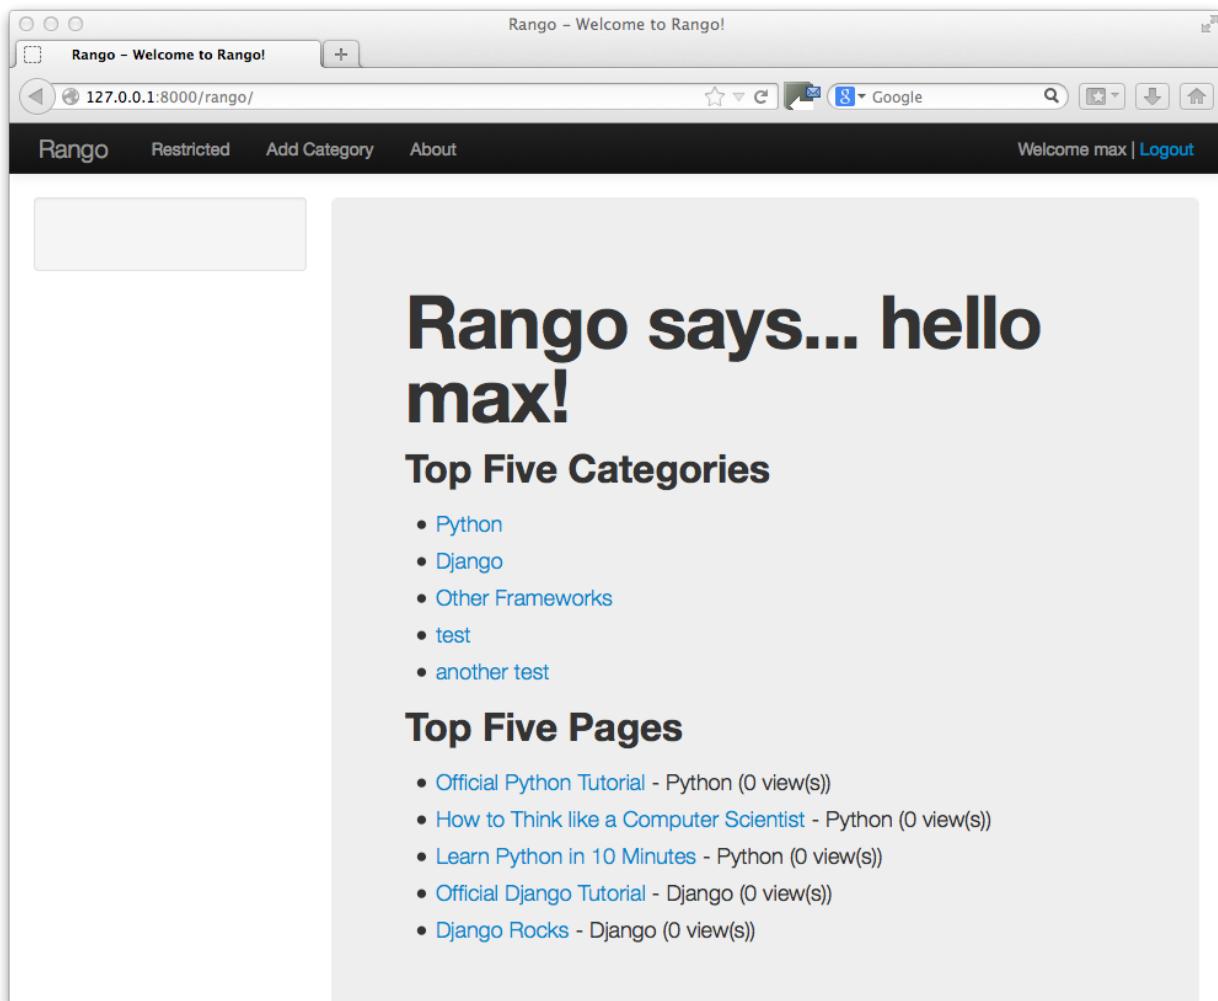


Fig. 11.4: A screenshot of the Index page with a Hero Unit.

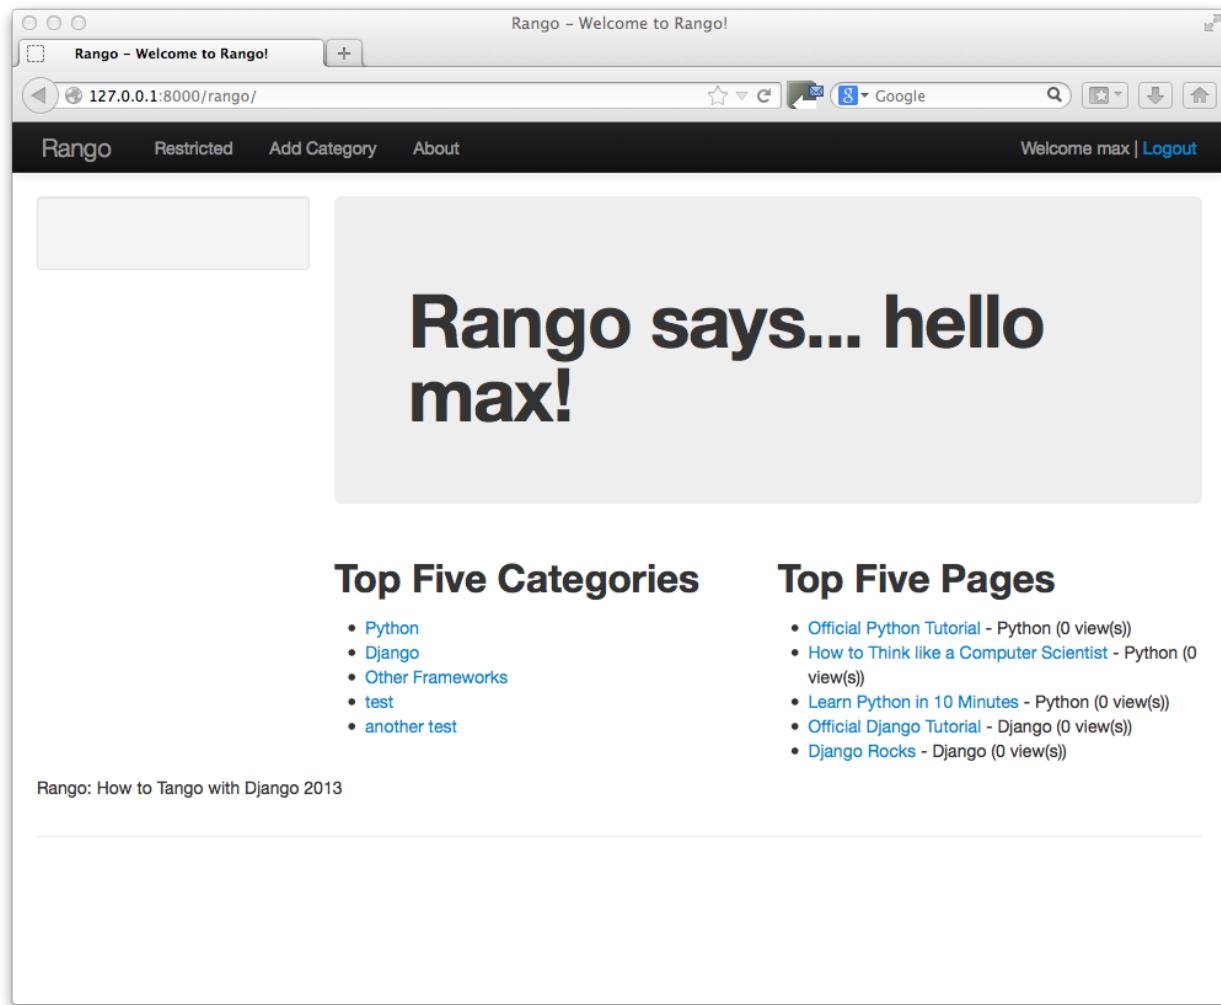


Fig. 11.5: A screenshot of the Index page with customised Bootstrap Styling.

11.3 The Login Page

Now let's turn our attention to the login page. On the Bootstrap website you can see they have already made a nice login form, see <http://getbootstrap.com/2.3.2/examples/signin.html>. If you take a look at the source, you'll notice that there are a number of classes that we need to include to pimp out the basic login form. Update the `login.html` template as follows:

```
{% block body_block %}  
<div class="hero-unit">  
    <h1>Login to Rango</h1>  
  
    <div class="container">  
        <form class="form-signin span4" id="login_form" method="post" action="/rango/login/">  
            <h2 class="form-signin-heading">Please sign in</h2>  
            {% csrf_token %}  
  
            {% if bad_details %}  
                <p><strong>Your username and/or password were incorrect!</strong></p>  
            {% elif disabled_account %}  
                <p><strong>Your Rango account is currently disabled; we can't log you in!</strong></p>  
            {% endif %}  
  
            Username: <input type="text" class="input-block-level" placeholder="Username" name="username"/>  
            <br />  
            Password: <input type="password" class="input-block-level" placeholder="Password" name="password"/>  
            <br />  
            <button class="btn btn-primary" type="submit">Sign in</button>  
        </form>  
  
    </div>  
    <!-- /container -->  
</div>  
{% endblock %}
```

We've made the following changes.

- `form-signin` and `span4` classes has been added to the form;
- `form-signin-heading` class as been put in the `<h2>` tag to head up the form;
- the input elements have had `input-block-level` classes added which control their width, along with placeholder text; and
- then the input element has been changed for a button element.

In the button, we have set the class to `btn` and `btn-primary`. If you check out the [Bootstrap toolkit](#) page on [Base CSS](#) you can see there are lots of different colours that can be assigned to buttons, see <http://getbootstrap.com/2.3.2/base-css.html#buttons>.

11.3.1 Other Form-based Templates

You can apply similar changes to `add_category.html` and `add_page.html` templates. For the `add_category.html` template, we can set it up as follows.

```
{% block body_block %}  
<div class="hero-unit">  
    <h1>Add a Category</h1>  
    <br/>
```

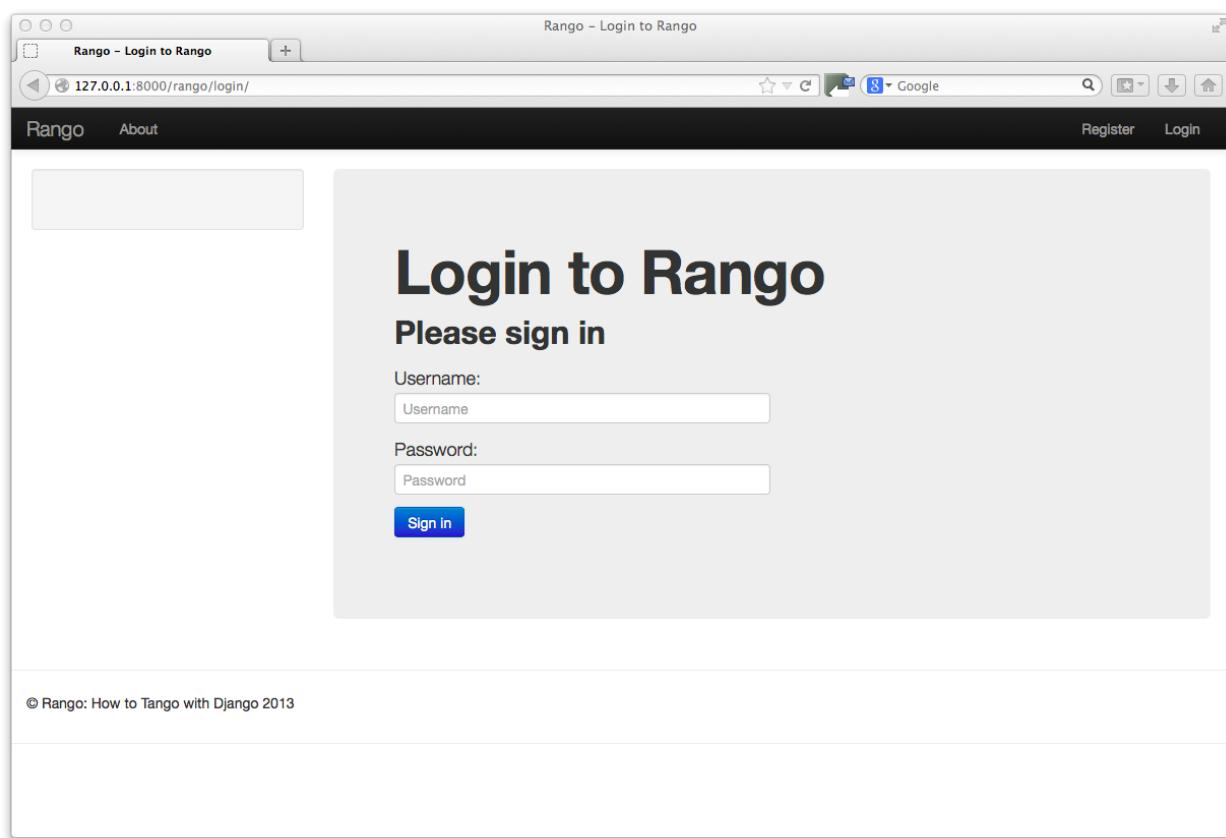


Fig. 11.6: A screenshot of the login page with customised Bootstrap Styling.

```
<div class="container">
    <form class="span6" id="category_form" method="post" action="/rango/add_category/">
        {% csrf_token %}
        {% for hidden in form.hidden_fields %}
        {{ hidden }}
        {% endfor %}

        {% for field in form.visible_fields %}
        {{ field.errors }}
        {{ field.help_text }}<br/>
        {{ field }}
        {% endfor %}

        <br/>
        <button class="btn btn-primary" type="submit" name="submit">Create Category</button>
    </form>
</div>
</div>
{% endblock %}
```

And similarly for the add_page.html template (not shown).

11.4 The Registration Template

The register.html template requires a bit more work. Currently, the template uses Django helper methods to convert the UserForm and UserProfileForm into HTML. However, we want a bit more control over the elements and how they are presented. This will require updating the UserForm and UserProfileForm as well as the register.html template.

Update the forms.py as follows.

```
class UserForm(forms.ModelForm):
    username = forms.CharField(help_text="Please enter a username.")
    email = forms.CharField(help_text="Please enter your email.")
    password = forms.CharField(widget=forms.PasswordInput(), help_text="Please enter a password.")

    class Meta:
        model = User
        fields = ['username', 'email', 'password']

class UserProfileForm(forms.ModelForm):
    website = forms.URLField(help_text="Please enter your website.", required=False)
    picture = forms.ImageField(help_text="Select a profile image to upload.", required=False)

    class Meta:
        model = UserProfile
        fields = ['website', 'picture']
```

Update the register.html template as follows.

```
{% block body_block %}
<div class="hero-unit">
    <h1>Register with Rango</h1>
    <br />

    <div class="container">
        {% if registered %}
```

```
<p>Thank you for registering.</p>

<p><a href="/rango/login/">Login</a> when you are ready to rango.</p>
{%
  else %
<form class="form-signin span8" id="user_form" method="post" action="/rango/register/"
      enctype="multipart/form-data">
  {% csrf_token %}
  <h2 class="form-signin-heading">Sign up Here</h2>
  <!-- Display each form here -->

  {% for field in user_form.visible_fields %}
  {{ field.errors }}
  {{ field.help_text }}<br />
  {{ field }}<br />
  {% endfor %}

  {% for field in profile_form.visible_fields %}
  {{ field.errors }}
  {{ field.help_text }}<br />
  {{ field }}<br />
  {% endfor %}

  <br />
  <!-- Provide a button to click to submit the form. -->
  <input class="btn btn-primary" type="submit" name="submit" value="Register"/>
</form>
{%
  endif %
</div>
</div>
{%
  endblock %}
```

Your registration form should be looking a lot better now and something like the figure below.

11.5 The End Result

Now that Rango is starting to look better we can go back and add in the extra functionality that will really pull the application together.

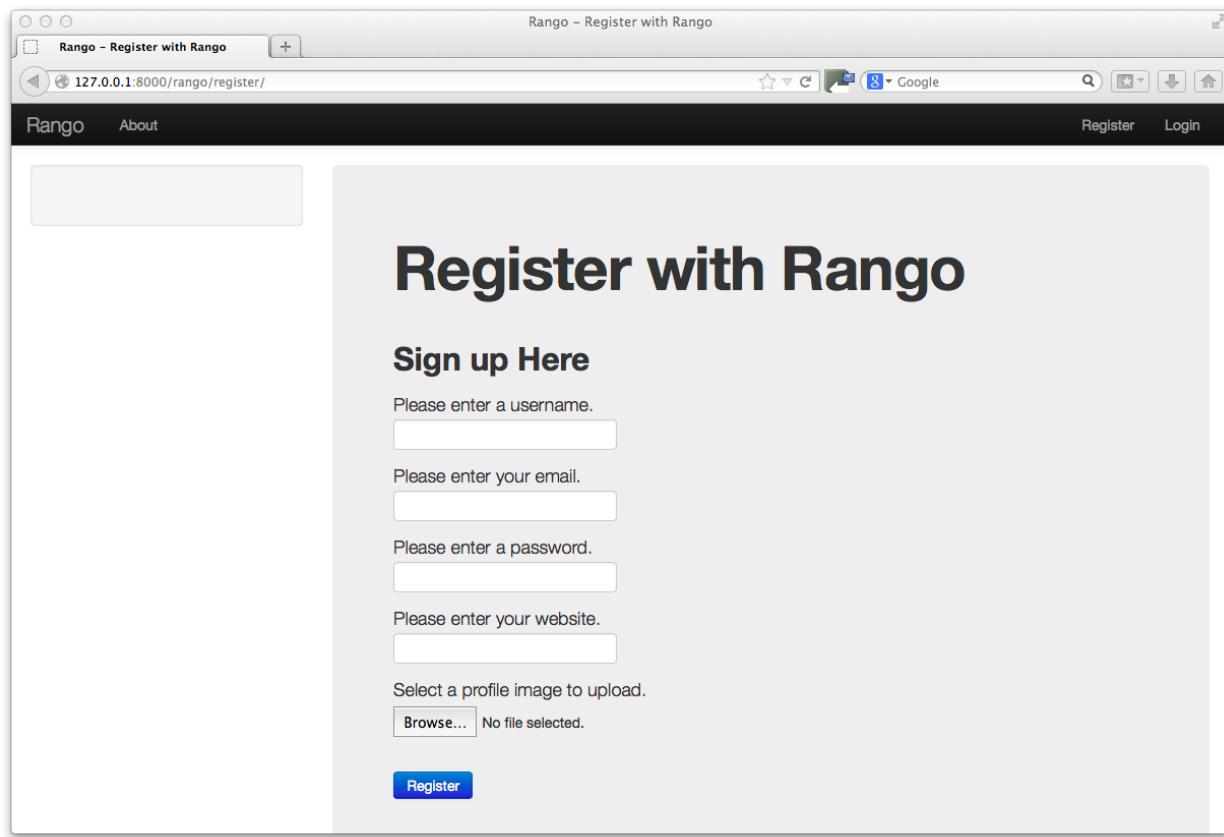


Fig. 11.7: A screenshot of the Registration page with customised Bootstrap Styling.

Adding External Search Functionality

At this stage, our Rango application is looking pretty good - a majority of our required functionality is implemented. In this chapter, we will connect Rango up to Bing's Search API so that users can also search for pages, rather than just used the categories. First let's get started by setting up an account to use Bing's Search API, then construct a wrapper to call Bing's web search functionality before integrating the search into Rango.

12.1 The Bing Search API

The Bing Search API provides you with the ability to embed search results from the Bing search engine within your own applications. Through a straightforward interface, you can request results from Bing's servers to be returned in either XML or JSON. The data returned can then be interpreted by a XML or JSON parser, with the results then rendered as part of a template within your application.

Although the Bing API can handle requests for different kinds of content, we'll be focusing on web search only for this tutorial - as well as handling JSON responses. To use the Bing Search API, you will need to sign up for an *API key*. The key currently provides subscribers with access to 5000 queries per month, which should be more than enough for our purposes.

12.1.1 Registering for a Bing API Key

To register for a Bing API key, you must first register for a free Microsoft account. The account provides you with access to a wide range of Microsoft services. If you already have a Hotmail account, you already have one! You can create your free account and login at <https://account.windowsazure.com>.

When you account has been created, jump to the [Windows Azure Marketplace Bing Search API page](#). At the top of the screen, you may first need to click the *Sign On* button - if you have already signed into your Microsoft account, you won't need to provide your account details again. If the text says *Sign Out*, you're already logged in.

Down the right hand side of the page is a list of transactions per month. At the bottom of the list is *5,000 Transactions/month*. Click the sign up button to the right - you should be subscribing for a free service. You should then read the *Publisher Offer Terms*, and if you agree with them click *Sign Up* to continue. You will then be presented with a page confirming that you have successfully signed up.

Once you've signed up, click the *Data* link at the top of the page. From there, you should be presented with a list of data sources available through the Windows Azure Marketplace. At the top of the list should be *Bing Search API* - it should also say that you are *subscribed* to the data source. Click the *use* link to the right of the page, and you will then be presented with a screen similar to that shown in Figure 12.1.

This page allows you to try out the Bing Search API by filling out the boxes to the left. For example, the *Query* box allows you to specify a query to send to the API. Ensure that at the bottom of the screen you select *Web* for web search

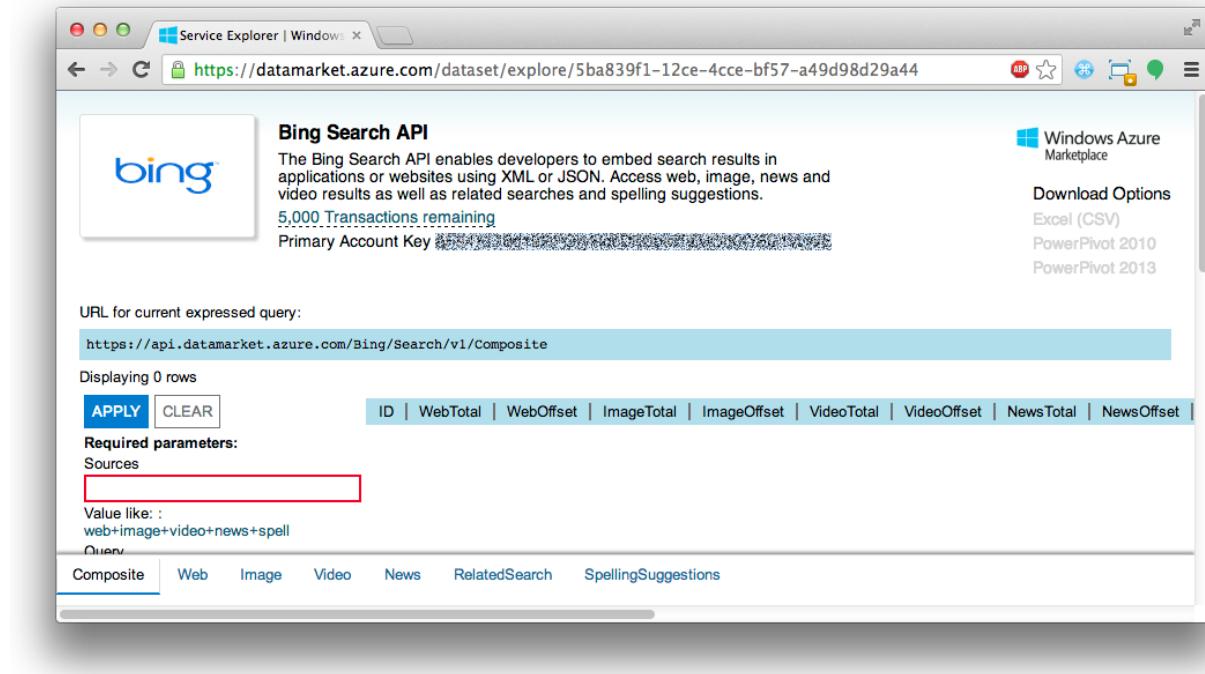


Fig. 12.1: The Bing Search API service explorer page. The *Primary Account Key* is deliberately obscured.

results only. Note the URL provided in the blue box at the top of the page changes as you alter the settings within the webpage. Take a note of the Web search URL. We'll be using part of this URL within our code later on. The following example is a URL to perform a web search using the query *rango*.

```
https://api.datamarket.azure.com/Bing/Search/v1/Web?Query=%27rango%27
```

We must also retrieve your API key so you can authenticate with the Bing servers when posing requests. To obtain your key, locate the text *Primary Account Key* at the top of the page and click the *Show* link next to it. Your key will then be exposed. We'll be using it later, so take a note of it - and keep it safe! If someone obtains your key, they'll be able to use your free query quota.

Note: The Bing API Service Explorer also keeps a tab of how many queries you have left of your monthly quota. Check out the top of the page to see!

12.2 Adding Search Functionality

To add search functionality to Rango, we first must write an additional function to query the Bing API. This code should take in the request from a particular user and return to the calling function a list of results. Any errors that occur during the API querying phase should also be handled gracefully within the function. Spinning off search functionality into an additional function also provides a nice abstraction between Django-related code and search functionality code.

To start, let's create a new Python module called `bing_search.py` within our `rango` application directory. Add the following code into the file. Check out the inline commentary for a description of what's going on throughout the function.

```

import json
import urllib, urllib2

def run_query(search_terms):
    # Specify the base
    root_url = 'https://api.datamarket.azure.com/Bing/Search/'
    source = 'Web'

    # Specify how many results we wish to be returned per page.
    # Offset specifies where in the results list to start from.
    # With results_per_page = 10 and offset = 11, this would start from page 2.
    results_per_page = 10
    offset = 0

    # Wrap quotes around our query terms as required by the Bing API.
    # The query we will then use is stored within variable query.
    query = "'{0}'".format(search_terms)
    query = urllib.quote(query)

    # Construct the latter part of our request's URL.
    # Sets the format of the response to JSON and sets other properties.
    search_url = "{0}{1}?format=json&$top={2}&$skip={3}&Query={4}".format(
        root_url,
        source,
        results_per_page,
        offset,
        query)

    # Setup authentication with the Bing servers.
    # The username MUST be a blank string, and put in your API key!
    username = ''
    bing_api_key = '<api_key>'

    # Create a 'password manager' which handles authentication for us.
    password_mgr = urllib2.HTTPPasswordMgrWithDefaultRealm()
    password_mgr.add_password(None, search_url, username, bing_api_key)

    # Create our results list which we'll populate.
    results = []

    try:
        # Prepare for connecting to Bing's servers.
        handler = urllib2.HTTPBasicAuthHandler(password_mgr)
        opener = urllib2.build_opener(handler)
        urllib2.install_opener(opener)

        # Connect to the server and read the response generated.
        response = urllib2.urlopen(search_url).read()

        # Convert the string response to a Python dictionary object.
        json_response = json.loads(response)

        # Loop through each page returned, populating out results list.
        for result in json_response['d']['results']:
            results.append({
                'title': result['Title'],
                'link': result['Url'],
                'summary': result['Description']})



```

```
# Catch a URLError exception - something went wrong when connecting!
except urllib2.URLError, e:
    print "Error when querying the Bing API: ", e

# Return the list of results to the calling function.
return results
```

The logic of the function above can be broadly split into six main tasks:

- First, the function prepares for connecting to Bing by preparing the URL that we'll be requesting.
- The function then prepares authentication, making use of your Bing API key. Make sure you replace <api_key> with your actual Bing API key, otherwise you'll be going nowhere!
- We then connect to the Bing API through the command `urllib2.urlopen(search_url)`. The results from the server are read and saved as a string.
- This string is then parsed into a Python dictionary object using the `json` Python package.
- We loop through each of the returned results, populating a `results` dictionary. For each result, we take the `title` of the page, the `link` or URL and a short summary of each returned result.
- The dictionary is returned by the function.

Notice that results are passed from Bing's servers as JSON. This is because we explicitly specify to use JSON in our initial request - check out the `search_url` variable which we define. If an error occurs when attempting to connect to Bing's servers, the error is printed to the terminal via the `print` statement within the `except` block.

Note: There are many different parameters that the Bing Search API can handle which we don't cover here. If you're interested in seeing how to tailor your results, check out the [Bing Search API Migration Guide and FAQ](#).

12.3 Putting Search into Rango

To add external search functionality, we will need to perform the following steps.

1. We must first create a `search.html` template which extends from our `base.html` template. The `search.html` template will include a HTML `<form>` to capture the user's query as well as template code to present any results.
2. We then create a view to handle the rendering of the `search.html` template for us, as well as calling the `run_query()` function we defined above.

12.3.1 Adding a Search Template

Let's first create our `search.html` template. Add the following HTML markup and Django template code.

```
{% extends "rango/base.html" %}

{% load static %}

{% block title %}Search{% endblock %}

{% block body_block %}
<div class="hero-unit">
    <h1>Search with Rango</h1>
    <br/>
```

```

<div class="container-fluid">
    <form class="form-signin span8" id="user_form" method="post" action="/rango/search/">
        {% csrf_token %}
        <!-- Display the search form elements here -->
        <input type="text" size="50" name="query" value="" id="query" />
        <input class="btn btn-primary" type="submit" name="submit" value="Search" />
        <br />
    </form>

    {% if result_list %}
        <!-- Display search results in an ordered list -->
        <div style="clear: both;">
            <ol>
                {% for result in result_list %}
                    <li>
                        <strong><a href="{{ result.link }}">{{ result.title }}</a></strong><br />
                        <em>{{ result.summary }}</em>
                    </li>
                {% endfor %}
            </ol>
        </div>
    {% endif %}
</div>
{% endblock %}

```

The template code above performs two key tasks:

1. In all scenarios, the template presents a search box and a search buttons within a HTML `<form>` for users to enter and submit their search queries.
2. If a `results_list` object is passed to the template's context when being rendered, the template then iterates through the object displaying the results contained within.

As you will see from our corresponding view code shortly, a `results_list` will only be passed to the template engine when there are results to return. There won't be results for example when a user lands on the search page for the first time - they wouldn't have posed a query yet!

12.3.2 Adding the View

With our search template added, we can then add the view which prompts the rendering of our template. Add the following `search()` view to Rango's `views.py` module.

```

def search(request):
    context = RequestContext(request)
    result_list = []

    if request.method == 'POST':
        query = request.POST['query'].strip()

        if query:
            # Run our Bing function to get the results list!
            result_list = run_query(query)

    return render_to_response('rango/search.html', {'result_list': result_list}, context)

```

By now, the code should be pretty self explanatory to you. The only major addition is the calling of the `run_query()` function we defined earlier in this chapter. To call it, we are required to also import the `bing_search.py` module,

too. Ensure that before you run the script that you add the following import statement at the top of the `views.py` module.

```
from rango.bing_search import run_query
```

You'll also need to ensure you do the following, too.

1. Add a mapping between your `search()` view and the `/rango/search/` URL.
2. Update the `base.html` navigation bar to include a link to the search page.

Note: According to the [relevant article on Wikipedia](#), an *Application Programming Interface (API)* specifies how software components should interact with one another. In the context of web applications, an API is considered as a set of HTTP requests along with a definition of the structures of response messages that each request can return. Any meaningful service that can be offered over the Internet can have its own API - we aren't limited to web search. For more information on web APIs, [Luis Rei provides an excellent tutorial on APIs](#).

12.4 Exercises

Taking the basic Bing Search API function we added above as a baseline, try out the following exercises.

- Add a `main()` function to the `bing_search.py` to test out the BING Search API when you run `python bing_search.py`.
- The main function should ask a user for a query (from the command line), and then issue the query to the BING API via the `run_query` method and print out the top ten results returned.
- Print out the rank, title and URL for each result.

Making Rango Tango! Exercises

So far we have been adding in different pieces of functionality to Rango. We've been building up the application in this manner to get you familiar with the Django Framework, and to learn about how to construct the various parts of a website that you are likely to make in your own projects. Rango however at the present moment is not very cohesive. In this chapter, we challenge you to improve the application and its user experience by bringing together functionality that we've already implemented alongside some awesome new additions.

To make Rango more coherent and integrated it would be nice to add the following functionality.

- **Integrate the browsing and searching within categories, i.e.:**

- provide categories on every page;
- provide some way to search through categories (see Chapter 15); and
- instead of have a disconnected search page, let users search for pages within a category so that they can then add these pages to the category (see Chapter 15)

- **Provide services for Registered Users, i.e.:**

- let users view their profile;
- let users edit their profile; and
- let users see the list of users and their profiles.

- **Track the click throughs of Categories and Pages, i.e.:**

- count the number of times a category is viewed;
- count the number of times a page is viewed via Rango; and
- collect likes for categories (see Chapter 15).

Note: We won't be working through all of these tasks right now. Some will be taken care of in Chapter 15, while some will be left to you to complete as additional exercises.

Before we start to add this additional functionality we will make a todo list to plan our workflow for each task. Breaking tasks down into sub-tasks will greatly simplify the implementation and mean that we are attacking each one with a clear plan. In this chapter, we will provide you with the workflow for a number of the above tasks. From what you have learnt so far, you should be able to fill in the gaps and implement most of it on your own. The following chapter however includes the code walkthrough, along with notes on how we have implemented each task.

13.1 Providing Categories on Every Page

It would be nice to show the different categories that users can browse through on any page, and not just the index page. There are a number of ways that this can be achieved, but given what we have learnt so far, we could undertake the following workflow:

- Create a `category_list.html` template and migrate the template code for presenting the categories list from `index.html` to this new template. This template won't be a full HTML document, but a portion of one which we can include in other templates.
- Use the `{% include "rango/category_list.html" %}` to now include this template code into the `base.html` template:
 - To handle the scenario where `cat_list` isn't available, add the conditional `{% if cat_list %}` to ensure that only templates providing `cat_list` attempt to render this component.
- Migrate the code that gathers the list of categories from the `index()` view and place it into your `category()` view. Why?
 - Create a helper function called `get_category_list()` within `views.py` that returns the list of categories.
 - We can then call this function whenever we want to get the category list to be presented in the sidebar. This saves a lot of code repetition!
- Pass the category list data (`cat_list`) through to the template to complete the process.

13.2 Searching Within a Category Page

Rango aims to provide users with a helpful directory of page links. At the moment, the search functionality is essentially independent of the categories. It would be nicer however to have search integrated into category browsing. Let's assume that a user will first browse their category of interest first. If they can't find the page that they want, they can then search for it. If they find a page that is suitable, then they can add it to the category that they are in. Let's tackle the first part of this description here.

We first need to remove the global search functionality and only let users search within a category. This will mean that we essentially decommission the current search page and search view. After this, we'll need to perform the following steps.

- Remove the generic *Search* link from the menu bar.
- Take the search form and results template markup from `search.html` and place it into `category.html`.
- Update the category view to handle a HTTP POST request. The view must then include any search results in the context dictionary for the template to render.

13.3 View Profile

Another useful feature to add is a profile page, where users can view details of their Rango profile. Undertake the following steps to add this functionality.

- First, create a template called `profile.html`. In this template, add in the fields associated with the user profile and the user (i.e. username, email, website and picture).
- Create a view called `profile()`. This view will obtain the data required to render the user profile template.
- Map the URL `/rango/profile/` to your new `profile()` view.

- In the base template add a link called *Profile* into the menu bar, preferably on the right-hand side with other user-related links. This should only be available to users who are logged in (i.e. `{% if user.is_authenticated %}`).

13.4 Track Page Click Throughs

Currently, Rango provides a direct link to external pages. This is not very good if you want to track the number of times each page is clicked and viewed. To count the number of times a page is viewed via Rango you will need to perform the following steps.

- Create a new view called `track_url()`, and map it to URL `/rango/goto/`.
- The `track_url()` view will examine the HTTP GET request parameters and pull out the `page_id`. The HTTP GET request parameters are passed in the `request.GET` dictionary.
 - Your view should then be able to find the relevant `Page` model for the selected page, and add 1 to the associated `views` field.
 - The view will then redirect the user to the specified URL using Django's `redirect` method.
 - In the scenario where no parameters are in the HTTP GET request for `page_id`, or the parameters do not return a `Page` object, redirect the user to Rango's homepage.
- Update the `category.html` so that it uses `/rango/goto/?page_id=XXX` instead of using the direct URL.

13.4.1 Hint

If you're unsure of how to retrieve the `page_id` querystring from the HTTP GET request, the following code sample should help you.

```
if request.method == 'GET':
    if 'page_id' in request.GET:
        page_id = request.GET['page_id']
```

Always check the request method is of type GET first, then you can access the dictionary `request.GET` which contains values passed as part of the request. If `page_id` exists within the dictionary, you can pull the required value out with `request.GET['page_id']`.

Doing the Tango with Rango!

Hopefully, you will have been able to complete the exercises given the workflows we provided, if not, or if you need a little help checkout snippets of code and use them within your version of Rango.

14.1 List Categories on Each Page

14.1.1 Creating a Category List Template

Create a `category_list.html` template so that it looks like the code below:

```
<ul class="nav nav-list">
    {%
        if cat_list %
            {%
                for cat in cat_list %
                    <li><a href="/rango/category/{{ cat.url }}">{{ cat.name }}</a></li>
                {%
                    endfor %
                }
            {%
                else %
                    <li>No categories at present.</li>
            {%
                endif %
            }
    </ul>
```

14.1.2 Updating the Base Template

In the sidebar `<div>`, add `{% include "rango/category_list.html" %}` so that the `base.html` contains:

```
<div class="well sidebar-nav">
    {% block sidebar %}
    {% endblock %}
    <div id="cats">
        {% if cat_list %}
            <ul class="nav nav-list"><li>Category List</li></ul>
            {% include 'rango/category_list.html' with cat_list=cat_list %}
        {% endif %}
    </div>
</div>
```

14.1.3 Creating Get Category List Function

In `rango/views.py`, create a function called `get_category_list()` that returns the list of categories.

```
def get_category_list():
    cat_list = Category.objects.all()

    for cat in cat_list:
        cat.url = encode_url(cat.name)

    return cat_list
```

14.1.4 Updating Views

Then call this function in each of the views that you want to display the category list in the sidebar, and pass the list into the context dictionary. For example, to have the categories showing on the index page, alter the `index()` view as follows:

```
def index(request):
    context = RequestContext(request)
    cat_list = get_category_list()

    context_dict['cat_list'] = cat_list

    .....

    return render_to_response('rango/index.html', context_dict, context)
```

Note: to add the category list to all the other pages, you will need to do some refactoring to pass in all the context variables.

14.2 Searching Within a Category Page

Rango aims to provide users with a helpful directory of page links. At the moment, the search functionality is essentially independent of the categories. It would be nicer however to have search integrated into category browsing. Let's assume that a user will first browse their category of interest first. If they can't find the page that they want, they can then search for it. If they find a page that is suitable, then they can add it to the category that they are in. Let's tackle the first part of this description here.

We first need to remove the global search functionality and only let users search within a category. This will mean that we essentially decommission the current search page and search view. After this, we'll need to perform the following.

14.2.1 Decommissioning Generic Search

Remove the generic `Search` link from the menu bar by editing the `base.html` template. You can also remove or comment out the URL mapping in `rango/urls.py`.

14.2.2 Creating a Search Form Template

Take the search form from `search.html` and put it into the `category.html`. Be sure to change the action to point to the `category()` view as shown below.

```
<div class="container-fluid">
    <p>Search for a page.</p>
    <form class="span8 form-search" id="search_form" method="post" action="/rango/category/{{ category }}>
        {% csrf_token %}</form>
```

```

<input type="text" class="input-long search-query" name="query" value="{{ category_name }}"/>
<button type="submit" class="btn btn-success" name="submit" value="Search">Search</button>
</form>
</div>

```

Also include a <div> to house the results underneath.

```

<div class="container-fluid">
    {% if result_list %}
        <!-- Display search results in an ordered list -->
        <ol>
            {% for result in result_list %}
                <li>
                    <strong><a href="{{ result.link }}">{{ result.title }}</a></strong><br />
                    <p>{{ result.summary }}</p>
                </li>
            {% endfor %}
        </ol>
    {% else %}
        <br/>
        <p>No results found</p>
    {% endif %}
</div>

```

14.2.3 Updating the Category View

Update the category view to handle a HTTP POST request (i.e. when the user submits a search) and inject the results list into the context. The following code demonstrates this new functionality.

```

def category(request, category_name_url):
    # Request our context
    context = RequestContext(request)

    # Change underscores in the category name to spaces.
    # URL's don't handle spaces well, so we encode them as underscores.
    category_name = decode_url(category_name_url)

    # Build up the dictionary we will use as our template context dictionary.
    context_dict = {'category_name': category_name, 'category_name_url': category_name_url}

    cat_list = get_category_list()
    context_dict['cat_list'] = cat_list

    try:
        # Find the category with the given name.
        # Raises an exception if the category doesn't exist.
        # We also do a case insensitive match.
        category = Category.objects.get(name__iexact=category_name)
        context_dict['category'] = category
        # Retrieve all the associated pages.
        # Note that filter returns >= 1 model instance.
        pages = Page.objects.filter(category=category).order_by('-views')

        # Adds our results list to the template context under name pages.
        context_dict['pages'] = pages
    except Category.DoesNotExist:
        # We get here if the category does not exist.

```

```
# Will trigger the template to display the 'no category' message.
pass

if request.method == 'POST':
    query = request.POST['query'].strip()
    if query:
        result_list = run_query(query)
        context_dict['result_list'] = result_list

# Go render the response and return it to the client.
return render_to_response('rango/category.html', context_dict, context)
```

14.3 View Profile

To add the view profile functionality, undertake the following steps.

14.3.1 Creating the Profile Template

First, create a new template called `profile.html`. In this template, add the following code.

```
{% extends "rango/base.html" %}

{% block title %}Profile{% endblock %}

{% block body_block %}
<div class="hero-unit">
    <h1> Profile <h1> <br/>
    <h2>{{ user.username }}</h2>
    <p>Email: {{ user.email }}</p>

    {% if userprofile %}
        <p>Website: <a href="{{ userprofile.website }}">{{ userprofile.website }}</a></p>
        <br/>
        {% if userprofile.picture %}
            
        {% endif %}
        {% endif %}
    </div>
{% endblock %}
```

14.3.2 Creating Profile View

Create a view called `profile` and add the following code.

```
from django.contrib.auth.models import User

@login_required
def profile(request):
    context = RequestContext(request)
    cat_list = get_category_list()
    context_dict = {'cat_list': cat_list}
    u = User.objects.get(username=request.user)
```

```

try:
    up = UserProfile.objects.get(user=u)
except:
    up = None

context_dict['user'] = u
context_dict['userprofile'] = up
return render_to_response('rango/profile.html', context_dict, context)

```

14.3.3 Mapping the Profile View and URL

Create a mapping between the URL /rango/profile and the profile() view. Do this by updating the urlpatterns tuple in rango/urls.py so that it includes the following entry.

```
url(r'^profile/$', views.profile, name='profile'),
```

14.3.4 Updating the Base Template

In the base.html template, update the code to put a link to the profile page in the menu bar.

```

{% if user.is_authenticated %}
    <li><a href="/rango/profile">Profile</a></li>
{% endif %}

```

14.4 Track Page Click Throughs

Currently, Rango provides a direct link to external pages. This is not very good if you want to track the number of times each page is clicked and viewed. To count the number of times a page is viewed via Rango you will need to perform the following steps.

14.4.1 Creating a URL Tracking View

Create a new view called track_url() in /rango/views.py which takes a parameterised HTTP GET request (i.e. rango/goto/?page_id=1) and updates the number of views for the page. The view should then redirect to the actual URL.

```

def track_url(request):
    context = RequestContext(request)
    page_id = None
    url = '/rango/'
    if request.method == 'GET':
        if 'page_id' in request.GET:
            page_id = request.GET['page_id']
            try:
                page = Page.objects.get(id=page_id)
                page.views = page.views + 1
                page.save()
                url = page.url
            except:
                pass
    return redirect(url)

```

Be sure that you import the `redirect()` function to `views.py` if it isn't included already!

```
from django.shortcuts import redirect
```

14.4.2 Mapping URL

In `/rango/urls.py` add the following code to the `urlpatterns` tuple.

```
url(r'^goto/$', views.track_url, name='track_url'),
```

14.4.3 Updating the Category Template

Update the `category.html` template so that it uses `rango/goto/?page_id=XXX` instead of providing the direct URL for users to click.

```
{% if pages %}
<ul>
    {% for page in pages %}
        <li>
            <a href="/rango/goto/?page_id={{page.id}}">{{page.title}}</a>
            {% if page.views > 1 %}
                - {{page.views}} views
            {% elif page.views == 1 %}
                - {{page.views}} view
            {% endif %}
        </li>
    {% endfor %}
</ul>
{% else %}
<strong>No pages currently in category.</strong><br/>
{% endif %}
```

Here you can see that in the template we have added some control statements to display `view`, `views` or nothing depending on the value of `page.views`.

14.4.4 Updating Category View

Since we are tracking the number of click throughs you can now update the `category()` view so that you order the pages by the number of views. To confirm this works, click on a link and refresh the category view - the link you clicked should jump up the rankings.

AJAX, Django and JQuery

To make the interaction with the Rango application more seamless let's add in a number of features that use AJAX, such as:

- Add a “Like Button” to let registered users “like” a particular category
- Add inline category suggestions - so that when a user types they can quickly find a category
- Add an “Add Button” to let registered users quickly and easily add a Page to the Category

AJAX essentially is a combination of technologies that are integrated together to reduce the number of page loads. Instead of reloading the full page, only part of the page or the data in the page is reloaded. If you haven't used AJAX before or would like to know more about it before using it, check out the AJAX tutorial provided by the W3C Schools: <http://www.w3schools.com/ajax/default.ASP>.

To simplify the AJAX components you can use a library like JQuery. If you are using the Twitter CSS Bootstrap toolkit then JQuery will already be added in. Otherwise, download the latest version of JQuery and include it within your application.

To include JQuery within your application, in the static folder create a *js* folder and plonk the JQuery javascript file (*jquery.js*) here along with an file called *rango-ajax.js*, which will house our javascript code. In *rango-ajax.js*, add the following javascript:

```
$ (document) .ready (function () {  
    // JQuery code to be added in here.  
});
```

Then in your *base* template include:

```
<script src="{% static "/js/jquery.js" %}"></script>  
<script src="{% static "/js/rango-ajax.js" %}"></script>
```

If you aren't familiar with JQuery it is worth checking out <http://jquery.com> and going through some examples in the documentation. The documentation provides numerous worked examples of the different functionality that the JQuery API provides.

15.1 Add a “Like Button”

It would be nice to let user, who are registered, denote that they “like” a particular category. In the following workflow, we will let users “like” categories, but we will not be keeping track of what categories they have “liked”, we'll be trusting them not to click the like button multiple times.

15.1.1 Workflow

To let users “like” certain categories undertake the following workflow:

- In the `category.html` template:
 - Add in a “Like” button with `id="like"`.
 - Add in a template tag to display the number of likes: `{% category.likes %}`
 - Place this inside a `div` with `id="like_count"`, i.e. `<div id="like_count">{{ category.likes }}</div>`
 - This sets up the template to capture likes and to display likes for the category.
- Update the `category` view to pass through the number of likes for the category.
- Create a view called, `like_category` which will examine the request and pick out the `category_id` and then increment it.
 - Don’t forgot to add in a the url mapping; so map the view to `rango/like_category/` i.e. the GET request will then be `rango/like_category/?category_id=XXX`
 - Instead of return a HTML page have this view will return the new total number of likes for that category.
- Now in “`rango-ajax.js`” add the JQuery code to perform the AJAX GET request.
 - If the request is successful, then update the `#like_count` element, and hide the like button.

15.1.2 Updating Category Template

To prepare the template we will need to add in the “Like” button with `id="like"` and create a `<div>` to display the number of likes `''{{ category.likes }}''`. To do this, add the following `<div>` to the `category.html` template:

```
<p>
<b id="like_count">{{ category.likes }}</b> people like this category

{% if user.is_authenticated %}
    <button id ="likes" data-catid="{{category.id}}" class="btn btn-mini btn-primary" type="button">
        Like
    </button>
{% endif %}

</p>
```

15.1.3 Update the Category View

To display the number of likes and the “Likes” button modify the `category` view:

```
def category(request, category_name_url):
    context = RequestContext(request)
    cat_list = get_category_list()
    category_name = decode_url(category_name_url)

    context_dict = {'cat_list': cat_list, 'category_name': category_name}

    try:
        category = Category.objects.get(name=category_name)
```

```

# Add category to the context so that we can access the id and likes
context_dict['category'] = category

pages = Page.objects.filter(category=category)
context_dict['pages'] = pages
except Category.DoesNotExist:
    pass

return render_to_response('rango/category.html', context_dict, context)

```

15.1.4 Create a Like Category View

Create a view called, `like_category` in `rango/views.py` which will examine the request and pick out the `category_id` and then increment the number of likes for that category.

```

@login_required
def like_category(request):
    context = RequestContext(request)
    cat_id = None
    if request.method == 'GET':
        cat_id = request.GET['category_id']

    likes = 0
    if cat_id:
        category = Category.objects.get(id=int(cat_id))
        if category:
            likes = category.likes + 1
            category.likes = likes
            category.save()

    return HttpResponse(likes)

```

On examining the code, you will see that we are only allowing authenticated users to denote that they like a category. The view assumes that a variable `category_id` has been passed through via a GET or POST so that we can identify the category to update. In this view, we could also track and record that a particular user has “liked” this category if we wanted - but here we are keeping it simple to focus on the AJAX mechanics.

Don’t forget to add in the URL mapping, into `rango/urls.py`. Update the `urlpatterns` by adding in:

```
url(r'^like_category/$', views.like_category, name='like_category'),
```

15.1.5 Making the AJAX request

Now in “`rango-ajax.js`” you will need to add some JQuery code to perform an AJAX GET request. Add in the following code:

```

$('#likes').click(function() {
    var catid;
    catid = $(this).attr("data-catid");
    $.get('/rango/like_category/', {category_id: catid}, function(data) {
        $('#like_count').html(data);
        $('#likes').hide();
    });
});

```

This piece of JQuery/Javascript will add and event handler to the element with id `#likes`, i.e. the button. When clicked, it will extract the category id from the button element, and then make an AJAX GET request which will make a call to `/rango/like_category/` encoding the `category_id` in the request. If the request is successful, then the HTML element with id `like_count` (i.e. the ``) is updated with the data returned by the request, and the HTML element with id `likes` (i.e. the `<button>`) is hidden.

There is a lot going on here and getting the mechanics right when constructing pages with AJAX can be a bit tricky. Essentially here, when the button is clicked an AJAX request is made, given our url mapping, this invokes the `like_category` view which updates the category and returns a new number of likes. When the AJAX request receives the response it update part of the page i.e. the text and the button.

15.2 Adding Inline Category Suggestions

It would be really neat if we could provide a fast way for users to find a category, rather than browsing through a long list. To do this we can create a suggestion component which lets users type in a letter or part of a word, and then the system responds by providing a list of suggested categories, that the user can then select from. As the user types a series of requests will be made to the server to fetch the suggested categories relevant to what the user has entered.

15.2.1 Workflow

To do this you will need to do the following.

- Create a parameterised function called `get_category_list(max_results=0, starts_with='')` that returns a list of categories.
 - The function returns a list of category objects annotated with the encoded category denoted by the attribute, `url`
 - Create a view called `suggest_category` which will examine the request and pick out the category query string.
 - Assume that a GET request is made and attempt to get the `query` attribute.
 - If the query string is not empty, ask the Category model to get the top 8 categories that start with the query string.
 - The list of category objects will then be combined into a piece of HTML via template.
 - Instead of creating a template called `suggestions.html` re-use the `category_list.html` as it will be displaying data of the same type (i.e. categories).
 - To let the client ask for this data, you will need to create a URL mapping lets call it `category_suggest`
- With the mapping, view, and template for this view in place, you will need to update the `base.html` template and add in some javascript so that the categories can be displayed as the user types.
- In the `base.html` template modify the sidebar block so that a div with an id="cats" encapsulates the categories being presented.
 - Above this `<div>` add an input box for a user to enter the letters of a category, i.e.:

```
<input class="input-medium search-query" type="text" name="suggestion" value="" id="suggestion" />
```

 - With these elements added into the templates, you can add in some JQuery to update the categories list as the user types.
 - Associate an on keypress event handler to the `input` with `id="suggestion"`

- `$('#suggestion').keyup(function(){ ... })`
- On keyup, issue an ajax call to retrieve the updated categories list
- Then use the JQuery `.get()` function i.e. `$(this).get(...)`
- If the call is successful, replace the content of the `<div>` with id="cats" with the data received.
- Here you can use the JQuery `.html()` function i.e. `$('#cats').html(data)`

15.2.2 Parameterise the Get Category List Function

In this helper function we use a filter to find all the categories that start with the string supplied.

```
def get_category_list(max_results=0, starts_with=''):
    cat_list = []
    if starts_with:
        cat_list = Category.objects.filter(name__startswith=starts_with)
    else:
        cat_list = Category.objects.all()

    if max_results > 0:
        if len(cat_list) > max_results:
            cat_list = cat_list[:max_results]

    for cat in cat_list:
        cat.url = encode_url(cat.name)

    return cat_list
```

15.2.3 Create a Suggest Category View

Using the `get_category_list` function we can now create a view that returns the top 8 matching results as follows:

```
def suggest_category(request):
    context = RequestContext(request)
    cat_list = []
    starts_with = ''
    if request.method == 'GET':
        starts_with = request.GET['suggestion']
    else:
        starts_with = request.POST['suggestion']

    cat_list = get_category_list(8, starts_with)

    return render_to_response('rango/category_list.html', {'cat_list': cat_list}, context)
```

Note here we are re-using the `rango/category_list.html` template :-).

15.2.4 Map View to URL

Add the following code to `urlpatterns` in `rango/urls.py`:

```
url(r'^suggest_category/$', views.suggest_category, name='suggest_category'),
```

15.2.5 Update Base Template

In the base template in the sidebar div add in the following HTML code:

```
<ul class="nav nav-list">
    <li class="nav-header">Find a Category</li>
    <form>
        <label></label>
        <li><input class="search-query span10" type="text" name="suggestion" value="" id="suggestion">
    </form>
</ul>
{%- if cat_list %}
    <div id="cats">
        {%- include 'rango/category_list.html' with cat_list=cat_list %}
    </div>
{%- endif %}
```

Here we have added in an input box with `id="suggestion"` and div with `id="Cats"` in which we will display the response. We don't need to add a button as we will be adding an event handler on keyup to the input box which will send the suggestion request.

15.2.6 Add AJAX to Request Suggestions

Add the following JQuery code to the `js/rango-ajax.js`:

```
$('#suggestion').keyup(function() {
    var query;
    query = $(this).val();
    $.get('/rango/suggest_category/', {suggestion: query}, function(data) {
        $('#cats').html(data);
    });
});
```

Here, we attached an event handler to the HTML input element with `id="suggestion"` to trigger when a keyup event occurs. When it does the contents of the input box is obtained and placed into the `query` variable. Then a AJAX GET request is made calling `/rango/category_suggest/` with the `query` as the parameter. On success, the HTML element with `id="cats"` i.e. the div, is updated with the category list html.

15.3 Exercises

To let registered users quickly and easily add a Page to the Category put an “Add” button next to each search result.

- **Update the `category.html` template:**
 - Add a mini-button next to each search result (if the user is authenticated), garnish the button with the title and url data, so that the JQuery can pick it out.
 - Put a `<div>` with `id="page"` around the pages in the category so that it can be updated when pages are added.
 - Remove that link to add button, if you like.
- Create a view `auto_add_page` that accepts a parameterised GET request (`title`, `url`, `catid`) and adds it to the category
- Map an url to the view `url(r'^auto_add_page/$', views.auto_add_page, name='auto_add_page'),`

- Add an event handler to the button using JQuery - when added hide the button. The response could also update the pages listed on the category page, too.

15.3.1 Hints

HTML Template code:

```
{% if user.is_authenticated %}
    <button data-catid="{{category.id}}" data-title="{{ result.title }}" data-url="{{ result.link }}"
{% endif %}
```

JQuery code:

Note here we are assigned the event handler to all the buttons with class rango-add.

View code:

```
@login_required
def auto_add_page(request):
    context = RequestContext(request)
    cat_id = None
    url = None
    title = None
    context_dict = {}
    if request.method == 'GET':
        cat_id = request.GET['category_id']
        url = request.GET['url']
        title = request.GET['title']
        if cat_id:
            category = Category.objects.get(id=int(cat_id))
            p = Page.objects.get_or_create(category=category, title=title, url=url)

            pages = Page.objects.filter(category=category).order_by('-views')

            # Adds our results list to the template context under name pages.
            context_dict['pages'] = pages

    return render_to_response('rango/page_list.html', context_dict, context)
```

Deploying Your Project

This chapter provides a step-by-step guide on how to deploy your Django applications. We'll be looking at deploying applications on [PythonAnywhere](#), an online IDE and web hosting service. The service provides in-browser access to the server-based Python and Bash command line interfaces, meaning you can interact with PythonAnywhere's servers just like you would with a regular terminal instance on your own computer. Currently, PythonAnywhere are offering a free account which sets you up with an adequate amount of storage space and CPU time to get a Django application up and running.

16.1 Creating a PythonAnywhere Account

First sign up for a Beginner PythonAnywhere account. If your application takes off and becomes popular, you can always upgrade your account at a later stage to gain more storage space and CPU time along with a number of other benefits (like hosting specific domains and ssh abilities).

Once your account has been created, you will have your own little slice of the World Wide Web at <http://<username>.pythonanywhere.com>, where <username> is your PythonAnywhere username. It is from this URL that your hosted application will be available from.

16.2 The PythonAnywhere Web Interface

The PythonAnywhere web interface contains a *dashboard* which in turn provides a series of tabs allowing you to manage your application. The tabs as illustrated in Figure 16.1 include:

- a *consoles* tab, allowing you to create and interact with Python and Bash console instances;
- a *files* tab, which allows you to upload to and organise files within your disk quota;
- a *web* tab, allowing you to configure settings for your hosted web application;
- a *schedule* tab, allowing you to setup tasks to be executed at particular times; and
- a *databases* tab, which allows you to configure a MySQL instance for your applications should you require it.

Of the the five tabs provided, we'll be working primarily with the *consoles* and *web* tabs. The [PythonAnywhere](#) wiki provides a series of detailed explanations on how to use the other tabs.

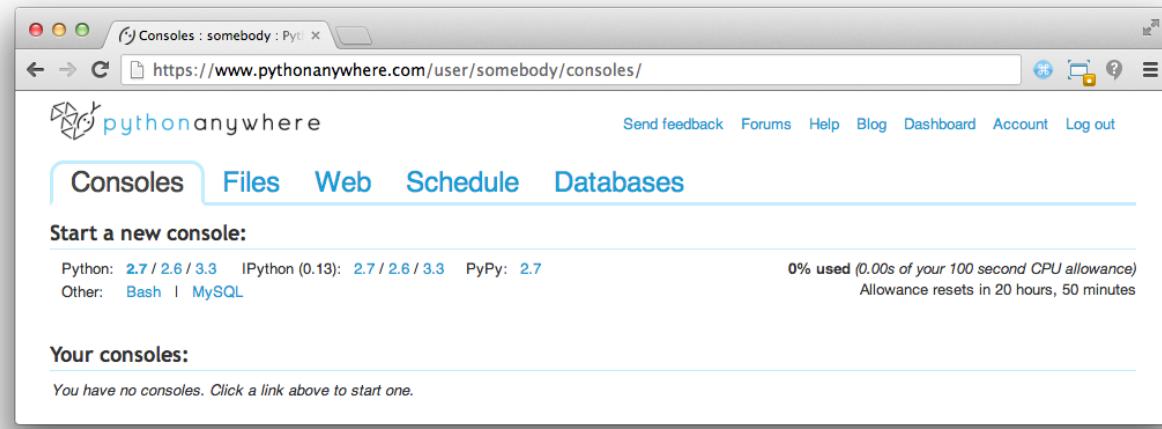


Fig. 16.1: The PythonAnywhere dashboard, showing the *Consoles* tab.

16.3 Creating a Virtual Environment

As part of its standard installation, PythonAnywhere comes with a number of packages pre-installed (i.e Python 2.7.4 and Django 1.3). However, since we are using a different setup, we need to setup a virtual environment and run our application inside that. Since our code base is compatible with Python 2.7.4 we can continue use that, but we will have to set up our virtual environment to use Django 1.5.4.

First, open a Bash console from the PythonAnywhere *Consoles* tab by clicking the *Bash* link. When the terminal is ready for you to interact, enter the following commands.

```
$ source virtualenvwrapper.sh  
$ mkvirtualenv rango
```

The first command imports the virtual environment wrapper. The wrapper provides a series of extensions by [Doug Hellman](#) to the original `virtualenv` tool, making it easier for us to create, delete and use virtual environments. The second command creates a new virtual environment called `rango`. This process should take a short while to create, after which you will be presented with a slightly different prompt.

```
(rango)16:38 ~ $
```

Note the inclusion of `(rango)` compared to your previous command prompt. This signifies that we have activated the `rango` virtual environment, so any package installations will be done within that environment, leaving the wider system setup alone. If you then issue the command `ls -la`, you will see that a directory called `.virtualenvs` has been created. This is the directory in which all of your virtual environments and associated packages will be stored. To confirm the setup, issue the command `which pip`. This will print the location in which the active `pip` binary is located - hopefully within `.virtualenvs` and `rango`, as shown in the example below.

```
/home/<username>/virtualenvs/test/bin/pip
```

Now we can customise our virtual environment by installing the required packages for our Rango application. Installing may take a considerable amount of time to install as your CPU time is limited - so have some patience. Issue the following two commands within your console:

```
$ pip install -U django==1.5.4  
$ pip install pil
```

After these have executed, check if Django has been installed with the command `which django-admin.py`. You

should receive output similar to the following example. We'll be using this virtual environment to work on the Rango application - whether we're synchronising the database or running population scripts.

```
/home/<username>/virtualenvs/rango/bin/django-admin.py
```

Note: PythonAnywhere also provides instructions on how to setup virtual environments, see <https://www.pythonanywhere.com/wiki/VirtualEnvForNewerDjango>.

16.3.1 Virtual Environment Switching

Moving between virtual environments can be done pretty easily. First you need to make sure that `virtualenvwrapper.sh` has been loaded by running `source virtualenvwrapper.sh`.

Rather than doing this each time you open up a console, you can add it to your `.bashrc` profile which is located in your home directory. Doing so will ensure the command is executed automatically for you every time you start a new Bash console instance. Any Bash consoles active will need to be closed for the changes to take effect.

With this done, you can then launch into a pre-existing virtual environment with the `workon` command. To load up the `rango` environment, enter:

```
16:48 ~ $ workon rango
```

where `rango` can be replaced with the name of the virtual environment you wish to use. Your prompt should then change to indicate you are working within a virtual environment.

```
(rango) 16:49 ~ $
```

You can then leave the virtual environment using the `deactivate` command. Your prompt should then be missing the `(rango)` prefix, with an example shown below.

```
(rango) 16:49 ~ $ deactivate  
16:51 ~ $
```

16.4 Cloning your Git Repository

Now that your virtual environment for Rango is all setup, you can now clone your Git repository to obtain a copy of your project's files. Clone your repository by issuing the following command from your home directory:

```
$ git clone https://<USERNAME>:<PASSWORD>@github.com/<OWNER>/<REPO_NAME>.git
```

where you replace - `<USERNAME>` with your GitHub username; - `<PASSWORD>` with your GitHub password; - `<OWNER>` with the username of the person who owns the repository; and - `<REPO_NAME>` with the name of your project's repository.

If you haven't put your code in a Git repository, you can clone the version we have made, by issuing the following command:

```
16:54 ~ $ git clone https://github.com/leifos/tango_with_django.git
```

Note: It doesn't matter if you clone your Git repository within your new virtual environment or not. You're only creating files within your disk quota, which doesn't require your special Python setup.

16.4.1 Setting Up the Database

With your files cloned, you must then prepare your database. We'll be using the `populate_rango.py` module that we created earlier in the book. As we'll be running the module, you must ensure that you are using the `rango` virtual environment (workon `rango`). From your home directory, move into the `tango_with_django` directory, and issue the following commands

```
(rango) 16:55 ~/tango_with_django $ python manage.py syncdb  
(rango) 16:56 ~/tango_with_django $ python populate_rango.py
```

As discussed earlier in the book, the first command synchronises your database with your project's installed models, and the second populates the database with some sample data.

16.5 Setting up your Web Application

Now that the database is setup, we need to configure the PythonAnywhere NGINX webserver to serve up your application. Within PythonAnywhere's web interface, navigate to your *dashboard* and click on the *Web* tab. On the left of the page that appears, click *Add a new web app*.

A popup box will then appear. Follow the instructions on-screen, and when the time comes, select the *manual configuration* option and complete the wizard.

Then, navigate to the PythonAnywhere subdomain at `http://<username>.pythonanywhere.com` in a new browser tab. You should be presented with the default `Hello, World!` webpage. This is because the WSGI script is currently serving up this page and not your Django application.

16.5.1 Configuring the WSGI Script

The [Web Server Gateway Interface](#), a.k.a. WSGI provides a simple and universal interface between web servers and web applications. PythonAnywhere uses WSGI to bridge the server-application link and map incoming requests to your subdomain to your web application.

To configure the WSGI script, navigate to the *Web* tab in PythonAnywhere's dashboard. From there, click the WSGI link located at the top of the page. The link should be preceded with text similar to `It is configured via a WSGI file stored at:`. The good people at PythonAnywhere have set up a sample WSGI file for us with several possible configurations. For your web application, you'll need to configure the Django section of the file. The example below demonstrates a possible configuration for your application.

```
# TURN ON THE VIRTUAL ENVIRONMENT FOR YOUR APPLICATION  
activate_this = '/home/<username>/.virtualenvs/rango/bin/activate_this.py'  
execfile(activate_this, dict(__file__=activate_this))  
  
import os  
import sys  
  
# ADD YOUR PROJECT TO THE PYTHONPATH FOR THE PYTHON INSTANCE  
path = '/home/<username>/tango_with_django/tango_with_django_project'  
  
if path not in sys.path:  
    sys.path.append(path)  
  
os.chdir(path)  
  
# TELL DJANGO WHERE YOUR SETTINGS MODULE IS LOCATED  
os.environ['DJANGO_SETTINGS_MODULE'] = 'tango_with_django_project.settings'
```

```
# IMPORT THE DJANGO WSGI HANDLER TO TAKE CARE OF REQUESTS
import django.core.handlers.wsgi
application = django.core.handlers.wsgi.WSGIHandler()
```

Ensure that you replace <username> with your username, and update any other path settings to suit your application. You should also remove all other code from the WSGI configuration script to ensure no conflicts take place.

The code sample above begins by activating your virtual environment `rango` as this has been configured with all the required packages. The script then adds your project's directory to the `PYTHONPATH` for the Python instance that runs your web application. This allows Python to access your project's modules. If you have additional paths to add, you can easily insert them here. You can then specify the location of your project's `settings.py` module. The final step is to include the Django WSGI handler and invoke it for your application.

When you have completed the WSGI configuration, click the *Save* button at the top-right of the webpage. Navigate back to the *Web* tab within the PythonAnywhere dashboard, and click the *Reload* button at the top of the page. When the application is reloaded, visiting `http://<username>.pythonanywhere.com` should present you with your Django application, all ready to go!

Note: During testing, we noted that you can sometimes receive HTTP 502 – Bad Gateway errors instead of your application. Try reloading your application again, and then waiting a longer. If the problem persists, try reloading again. If the problem still persists, check out your log files to see if any accesses/errors are occurring, before contacting the PythonAnywhere support.

16.5.2 Assigning Static Paths

We're almost there. One issue which we still have to address is to sort out paths for our application. Doing so will allow PythonAnywhere's servers to serve your static content, for example From the PythonAnywhere dashboard, click the *Web* tab and choose the subdomain hosting your application from the list on the left.

Underneath the *Static files* header, perform the following.

1. Click `Enter URL` and enter `/static/admin`, followed by return.
2. Click the corresponding `Enter path` text. Set this to `/home/<username>/.virtualenvs/rango/lib/python2.7/` where <username> should be replaced with your PythonAnywhere username. You may also need to change `rango` if this is not the name of your application's virtual environment. Remember to hit return to confirm the path.
3. Repeat the two steps above for the URL `/static/` and path `/home/<username>/tango_with_django/tango_with_django/static/` with the path setting pointing to the `static` directory of your web application.

With these changes saved, reload your web application by clicking the *Reload* button at the top of the page. Don't forget about potential HTTP 502 – Bad Gateway errors!

16.5.3 Bing API Key

Update `bing_search.py` with your own BING API Key to use the search functionality in Rango. Again, you will have to hit the *Reload* button for the changes to take effect.

16.5.4 Turning off DEBUG Mode

When your application is ready to go, it's a good idea to instruct Django that your application is now hosted on a production server. To do this, open your project's `settings.py` file and change `DEBUG = True` to `DEBUG = False`. This disables [Django's debug mode](#), and removes explicit error messages.

Changing the value of DEBUG also means you should set the ALLOWED_HOSTS property. Failing to perform this step will make Django return HTTP 400 Bad Request errors. Alter ALLOWED_HOSTS so that it includes your PythonAnywhere subdomain like in the example below.

```
ALLOWED_HOSTS = ['<username>.pythonanywhere.com']
```

Again, ensure <username> is changed to your PythonAnywhere username. Once complete, save the file and reload the application via the PythonAnywhere web interface.

16.6 Log Files

Deploying your web application to another environment introduces another layer of complexity to your setup. Unfortunately, it most likely won't be all plain sailing and will possibly result in new errors or unsuspecting problems. In order for you to diagnose and rectify the issues you may encounter, PythonAnywhere provides you with three log files that can help provide vital clues.

Log files can be viewed via the PythonAnywhere web interface by clicking on the *Web* tab, or by viewing the files in /var/log/ within a Bash console instance. The files provided are:

- `access.log`, which provides a log of requests made to your subdomain;
- `error.log`, which logs any error messages produced by your web application; and
- `server.log`, providing log details for the UNIX processes running your application.

Note that the names for each log file are prepended with your subdomain. For example, `access.log` will have the name `<username>.pythonanywhere.com.access.log`.

When debugging, you may find it useful to delete or move the log files so that you don't have to scroll through a huge list of previous attempts. If the files are moved or deleted, they will be recreated automatically when a new request or error arises.

16.7 Exercises

Congratulations, you've successfully deployed Rango! Why not tweet a link of your application to @tangowithdjango. We'd love to know you've succeeded!

Summary

In this book, we have gone through the process of web development from specification to deployment. Along the way we have shown how to use the Django framework to construct the models, views and templates associated with a web application. We have also demonstrated how toolkits and services like Bootstrap, JQuery, Bing Search, PythonAnywhere, etc can be integrated within an application. However, the road doesn't stop here. While, as we have only painted the broad brush strokes of a web application - as you have probably noticed there are lots of improvements that could be made to Rango - and these finer details often take a lot more time to complete as you polish the application. By developing your application on a firm base and good setup you will be able to construct up to 80% of your site very rapidly and get a working version up to demo fast.

In future versions of this book we intend to provide some more details on various aspects of the framework, along with covering the basics of some of the other fundamental technologies associated with web development. If you have any suggestions or comments about how to improve the book please get in touch.

A Git Crash Course

We strongly recommend that you spend some time familiarising yourself with a version control system. For your benefit, this section provides you with a crash course in how to use [Git](#), one of the many version control systems available. Originally developed by [Linus Torvalds](#), Git is today very popular, and is used by open-source and closed-source projects alike.

This tutorial demonstrates at a high level how Git works, explains the basic commands that you can use, and provides an explanation of Git's workflow.

18.1 Why Use Version Control?

As your software engineering skills develop, you will find that you are able to plan and implement ever more complex solutions to ever more complex problems. As a rule of thumb, the larger the problem specification, the more code you have to write. The more code you write, the greater the emphasis you should put on software engineering practices. Such practices include the use of design patterns and the DRY (don't repeat yourself) principle.

Most projects have many files and many people working on those files. This is a recipe for chaos. Have you ever encountered one or more of the following situations:

- Made a change to code, realised it was a mistake and wanted to go back?
- Lost code (through a faulty drive), or had a backup that was too old?
- Had to maintain multiple versions of a product (perhaps for different organisations)?
- Wanted to see the difference between two (or more) versions of your codebase?
- Wanted to show that a particular change broke or fixed a piece of code?
- Wanted to submit a change (patch) to someone else's code?
- Wanted to see how much work is being done (where it was done, when it was done, or who did it)?
- Wanted to experiment with a new feature without interfering with working code?

Using a version control system makes your life easier in *all* of the above cases. While using version control systems at the beginning may seem like a hassle it will pay off later - so get into the habit now.

18.2 Git on Windows

Like Python, Git doesn't come as part of a standard Windows installation. However, Windows implementations of the version control system can be downloaded and installed. You can download the official Windows Git client from the [Git website](#). The installer provides the `git` command line program, which we use in this crash course.

You can also download a program called *TortoiseGit*, a graphical extension to the Windows Explorer shell. The program provides a really nice right-click Git context menu for files. This makes version control really easy to use. You can [download TortoiseGit from here](#). Although we do not cover TortoiseGit in this crash course, many tutorials exist online for it. Check [this TortoiseGit tutorial](#) if you are interested in using it.

18.3 The Git System

Essentially, Git comprises of four separate storage locations: your *workspace*, the local *index*, the *local repository* and the *remote repository*. As the name may suggest, the remote repository is stored on some remote server - and is the only part of Git stored outwith your computer. This is considered a huge advantage of Git - you can make changes to your local repository when you may not have Internet access, and then apply those changes to the remote repository at a later stage.

Note: We keep repeating the word *repository*, but what do we actually mean by that? In the context of version control systems, consider a repository as a form of data structure that contains a set of *commit objects*, and a set of references to commit objects, called *heads*. You can find out more about what these are on [this Git tutorial](#) - and we will be explaining what the terminology for head means later on.

For now though, the following bullet points provide an explanation of each part of the Git system.

- As already explained, the *remote repository* is the copy of your project's repository stored on some remote server. This is particularly important for Git projects that have more than one contributor - you require a central place to store all the work that your team members produce. If you're feeling adventurous, you can set up a Git server on a computer with Internet access and a properly configured firewall (check out [this Git server tutorial](#), for example), or use one of many services providing free Git repositories. One of the most widely-used services available today is [GitHub](#). In fact, this book has a Git [repository](#) on GitHub!
- The *local repository* is a copy of the remote repository. The key difference however is that the local repository is stored on your own computer. It is to this repository you make all your additions, changes and deletions. When you reach a particular milestone, you can then push all your local changes to the remote repository. From there, you can instruct your team members to retrieve your changes. This concept is known as *pulling* from the remote repository, and we will explain that in a bit more detail later.
- The *index* is technically part of the local repository. The index stores a list of files that you want to be managed with version control. This is explained in more detail in the commands and workflow section. You can have a look [here](#) to see a discussion on what exactly a Git index contains.
- The final aspect of Git is your *workspace*. Think of this folder or directory as the place on your computer where you make changes to your version controlled files. From within your workspace, you can add new files or modify or remove previously existing ones. From there, you then instruct Git to update the repositories to reflect the changes you make in your workspace. This is important - *don't modify code inside the local repository - only ever edit files in your workspace*. The local repository contains a load of files that Git uses to keep track of your version controlled content. If you start messing around with these files, you'll more than likely break something!

Next, we'll be looking at how to get your Git workspace set up. We'll also discuss the basic workflow you should use when using Git.

18.4 Setting up Git

Setting up your Git workspace is a straightforward process. Once everything is set up, you will begin to make sense of the directory structure that Git uses. Assume that you have signed up for a new account on [GitHub](#) and [created a new repository on the service](#) for your project. With your remote repository setup, follow these steps to get your local

repository and workspace setup on your computer. We'll assume you will be working from your <workspace> directory.

1. Open a terminal and navigate to your home directory (e.g. `cd ~`).
2. *Clone* the remote repository - or in other words, make a copy of it. Check out how to do this below.
3. Navigate into the newly created directory. That's your workspace in which you can add files to be version controlled!

18.4.1 How to Clone a Remote Repository

Cloning your repository is a straightforward process with the `git clone` command. Supplement this command with the URL of your remote repository - and if required, authentication details, too. The URL of your repository varies depending on the provider you use. If you are unsure of the URL to enter, it may be worth querying it with your search engine or asking someone in the know.

For GitHub, try the following command, replacing the parts below as appropriate:

```
$ git clone https://<USERNAME>:<PASSWORD>@github.com/<OWNER>/<REPO_NAME>.git  
<workspace>
```

where you replace

- <USERNAME> with your GitHub username;
- <PASSWORD> with your GitHub password;
- <OWNER> with the username of the person who owns the repository;
- <REPO_NAME> with the name of your project's repository; and
- <workspace> with the name for your workspace directory. Although optional, we will specify it here to create the <workspace> directory.

If all is successful, you should see some positive messages in your terminal or Command Prompt alerting you to the fact that the clone has been successful.

18.4.2 The Directory Structure

Once you have cloned your remote repository onto your local computer, navigate into the directory with your terminal, Command Prompt or GUI file browser. If you have cloned an empty repository the workspace directory should appear empty. This directory is therefore your blank workspace with which you can begin to add files for your project.

However, the directory isn't blank at all! On closer inspection, you will notice a hidden directory called `.git`. Stored within this directory are both the local repository and index. Do not alter the contents of the `.git` directory. Doing so could damage your Git setup - and break version control functionality. *Your newly-created workspace directory therefore contains the workspace, local repository and index.*

18.4.3 Final Tweaks

With your workspace setup, now would be a good time to make some final tweaks. Here, we discuss two cool features you can try which could make your life (and your team members') a little bit easier.

When using your Git repository as part of a team, any changes you make will be associated with the username you use to access your remote Git repository. However, you can also specify your full name and e-mail address to be included with changes that are made by you on the remote repository. This is really easy to do. Simply open a Command

Prompt/terminal and navigate to your workspace. From there, issue two commands: one to tell Git your full name, and the other to tell Git your e-mail address.

```
$ git config user.name "John Doe"  
$ git config user.email "johndoe123@me.com"
```

Obviously, replace the example name and e-mail address with your own. We don't want random commits from some guy called John Doe! How unlucky would it be if you were actually called John Doe?

Anyway, moving on to the second feature. Git provides you with the capability to stop - or ignore - particular files from being added to version control. For example, you may not wish a file containing unique keys to access web services from being added to version control. If the file were to be added to the remote repository, anyone could theoretically access the file by cloning the repository.

With Git, files can be ignored by including them in the `.gitignore` file. This file which should reside in the root of your workspace. When adding files to version control, Git parses this file. If a file that is being added to version control is listed within `.gitignore`, the file is ignored. Each line of `.gitignore` should be a separate file entry. Check out the following example:

```
config/api_keys.py  
*.pyc
```

In this example file, there are two entries. The first one prompts git to ignore the file `api_keys.py` residing within the `config` directory. The second entry prompts Git to ignore *all* instances of files with a `.pyc` extension. This is really cool: you can use *wildcards* to make generic entries if you need to!

18.5 Basic Commands and Workflow

With your repository cloned and ready to go on your local computer, you're ready to get to grips with the Git workflow. This section shows you the basic Git workflow - and the associated Git commands you can issue.

Fig. 18.1: A diagram depicting the basic workflow and associated commands of interacting with a Git repository.

We have provided a pictorial representation of the basic Git workflow in Figure 18.1. Match each of the numbers in the blue circles to the numbered descriptions below.

18.5.1 1. Starting Off

Before you can start work on your project, you must prepare Git for your forthcoming geek session. If you haven't yet sorted out your project's Git workspace, you'll need to `clone` the repository to obtain a copy of all of its files. Check out Section 18.4.1 for more information on how to achieve this.

If you have previously made a clone of the remote repository, it's good practice to get into the habit of updating your local copy by using the `git pull` command. This 'pulls' changes from the remote repository. By doing this, you'll be working from the same page as your team members, which will help keep the issue of conflicting file contents from making your life a nightmare.

18.5.2 2. Doing Some Work!

Once your workspace has been updated with the latest changes, the onus is on you to do some work! Within your workspace, you can take existing files and modify them. You can delete them too, or add new files to be version controlled.

It's not all plain sailing, however. You must be aware that as you work away, you need to keep Git up-to-date on the list of files you have added, removed or updated by modifying the *local index*. The list of files stored within the local index are then used to perform your next *commit*, which we'll be discussing in the next step. To keep Git informed, there are several Git commands which let you update the local index. Three of the commands are near-identical to those that were discussed in Chapter 2, with the addition of a `git` prefix.

- The first command `git add` allows you to request Git to add a particular file to the next commit for you. A common newbie mistake is to assume that `git add` is used for adding new files to your repository only - *this is not the case! You must tell Git what modified files you wish to commit, too!* The command can be used in the fashion `git add <filename>`, where `<filename>` is the name of the file you wish to add to your next commit. Multiple files and directories can be added with the command `git add .` - *but be careful with this!*
- `git mv` performs the same function as the Unix `mv` command - it moves files. The only difference between the two is that `git mv` updates the local index for you before moving the file. Specify the filename with the syntax `git mv <filename>`. For example, with this command you can move files to a different directory within your repository. This will be reflected in your next commit.
- `git cp` allows you to make a copy of a file or directory while adding references to the new files into the local index for you. The syntax is the same as `git mv` above where the filename or directory name is specified thus: `git cp <filename>`.
- The command `git rm` adds a file or directory delete request into the local index. While the `git rm` command does not delete the file straight away, the requested file or directory is removed from your filesystem and the Git repository upon the next commit. The syntax is the same as the above commands, where a filename can be specified thus: `git rm <filename>`. Note that you can add a large number of requests to your local index in one go, rather than removing each file manually. For example, `git rm -rf media/` creates delete requests in your local index for the `media/` directory. The `r` switch enables Git to *recursively* remove each file within the `media/` directory, while `f` allows Git to *forcibly* remove the files. Check out the [Wikipedia page](#) on the `rm` command for more information.

Lots of changes between commits can make things pretty confusing. You may easily forgot what files you've already instructed Git to remove, for example. Fortunately, you can run the `git status` command to see a list of files which have been modified from your current working directory, but haven't been added to the local index for processing. Check out typical output from the command below to get a taste of what you can see.

```
$ git status

# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
#       modified:   chapters/requirements.rst
#
# Changes not staged for commit:
#   (use "git add/rm <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
#       modified:   ../TODO.txt
#       modified:   chapters/deploy.rst
#       deleted:    chapters/index.rst
#       deleted:    images/css-font.png
#       modified:   images/git-sequence.pdf
#       modified:   omnigraffle/git-sequence.graffle
#
```

For further information on this useful command, check out the [official Git documentation](#).

18.5.3 3. Committing your Changes

We've mentioned *committing* several times in the previous step - but what on earth does it mean? In the world of Git, committing is when you save changes - which are listed in the local index - that you have made within your workspace. The more often you commit, the greater the number of opportunities you'll have to revert back to an older version of your code if things go disastrously wrong! Make sure you commit often - but don't commit an incomplete or broken version of a particular module or function! There's a lot of online discussion about when the ideal time to commit is - have a look on [this Stack Overflow page](#) for the opinions of several developers.

To commit, you issue the `git commit` command. Any changes to existing files that you have indexed will be saved to version control at this point. Additionally, any files that you've requested to be copied, removed, moved or added to version control via the local index will be undertaken at this point. When you commit, you are updating the *HEAD* of your local repository. The HEAD is essentially the *latest commit at the top of the pile* - have a look at [this Stack Overflow page](#) for more information.

As part of a commit, it's incredibly useful to your future self and others to explain why you committed when you did. You can supply an optional message with your commit if you wish to do so - though we highly recommend it. Instead of simply issuing `git commit`, run the following amended command.

```
$ git commit -m "Updated helpers.py to include a Unicode conversion function,  
str_to_unicode()."
```

From the example above, you can see that using the `-m` switch followed by a string provides you with the opportunity to append a message to your commit. Be as explicit as you can, but don't write too much. People want to see at a glance what you did, and do not want to be bored with a long essay. At the same time, don't be too vague. Simply specifying `Updated helpers.py` may tell a developer what file you modified, but they will require further investigation to see exactly what you changed.

Note: Although frequent commits may be a good thing, you will want to ensure that what you have written actually *works* before you commit. This may sound silly, but it's an incredibly easy thing to not think about. Committing code which doesn't actually work can be infuriating to your team members if they then rollback to a version of your project's codebase which is broken!

18.5.4 4. Synchronising your Repository

After you've committed your local repository and committed your changes, you're just about ready to send your commits to the remote repository by *pushing* your changes. However, what if someone within your group pushes their changes before you do? This means your local repository will be out of sync with the remote repository, making any `git push` command very difficult to do!

It's therefore always a good idea to check whether changes have been made on the remote repository before updating it. Running a `git pull` command will pull down any changes from the remote repository, and attempt to place them within your local repository. If no changes have been made, you're clear to push your changes. If changes have been made and cannot be easily rectified, you'll need to do a little bit more work.

In scenarios such as this, you have the option to *merge* changes from the remote repository. After running the `git pull` command, a text editor will appear in which you can add a comment explaining why the merge is necessary. Upon saving the text document, Git will merge the changes in the remote repository to your local repository.

Note: If you do see a text editor on your Mac or Linux installation, it's probably the `vi` text editor. If you've never used `vi` before, check out [this helpful page containing a list of basic commands](#) on the Colorado State University Computer Science Department website. If you don't like `vi`, you can change the default text editor that Git calls upon. Windows installations most likely will bring up Notepad.

18.5.5 5. Pushing your Commit(s)

Pushing is the phrase used by Git to describe the sending of any changes in your local repository to the remote repository. This is the way in which your changes become available to your other team members, who can then retrieve them by running the `git pull` command in their respective local workspaces. The `git push` command isn't invoked as often as committing - *you require one or more commits to perform a push*. You could aim for one push per day, when a particular feature is completed, or at the request of a team member who is desperately after your updated code.

To push your changes, the simplest command to run is:

```
$ git push origin master
```

As explained on [this Stack Overflow question and answer page](#), this command instructs the `git push` command to push your local master branch (where your changes are saved) to the `origin` (the remote server from which you originally cloned). If you are using a more complex setup involving [branching and merging](#), alter `master` to the name of the branch you wish to push.

If what you are pushing is particularly important, you can also optionally alert other team members to the fact they should really update their local repositories by pulling your changes. You can do this through a *pull request*. Issue one after pushing your latest changes by invoking the command `git request-pull master`, where `master` is your branch name (this is the default value). If you are using a service such as GitHub, the web interface allows you to generate requests without the need to enter the command. Check out [the official GitHub website's tutorial](#) for more information.

18.6 Recovering from Mistakes

This section presents a solution to a coder's worst nightmare: what if you find that your code no longer works? Perhaps a refactoring went terribly wrong, someone changed something, or everything is so terribly messed up you have no idea what happened. Whatever the reason, using a form of version control always gives you a last resort: rolling back to a previous commit. This section details how to do just that. We follow the information given from [this Stack Overflow question and answer page](#).

Warning: You should be aware that this guide will rollback your workspace to a previous iteration. Any uncommitted changes that you have made will be lost, with a very slim chance of recovery! Be wary. If you are having a problem with only one file, you could always view the different versions of the files for comparison. Have a look at [this Stack Overflow page](#) to see how to do that.

Rolling back your workspace to a previous commit involves two distinct steps:

- determining which commit to roll back to; and
- performing the rollback.

To determine what commit to rollback to, you can make use of the `git log` command. Issuing this command within your workspace directory will provide a list of recent commits that you made, your name and the date at which you made the commit. Additionally, the message that is stored with each commit is displayed. This is where it is highly beneficial to supply commit messages that provide enough information to explain what is going on. Check out the following output from a `git log` invocation below to see for yourself.

commit 88f41317640a2b62c2c63ca8d755feb9f17cf16e	<- Commit hash
Author: John Doe <someaddress@domain.com>	<- Author
Date: Mon Jul 8 19:56:21 2013 +0100	<- Date/time
Nearly finished initial version of the requirements chapter	<- Message

```
commit f910b7d557bf09783b43647f02dd6519fa593b9f
Author: John Doe <someaddress@domain.com>
Date:   Wed Jul 3 11:35:01 2013 +0100

    Added in the Git figures to the requirements chapter.

commit c97bb329259ee392767b87cfe7750ce3712a8bdf
Author: John Doe <someaddress@domain.com>
Date:   Tue Jul 2 10:45:29 2013 +0100

    Added initial copy of Sphinx documentation and tutorial code.

commit 2952efa9a24dbf16a7f32679315473b66e3ae6ad
Author: John Doe <someaddress@domain.com>
Date:   Mon Jul 1 03:56:53 2013 -0700

    Initial commit
```

From this list, you can choose a commit to rollback to. For the selected commit, you must take the commit hash - the long string of letters and numbers. To demonstrate, the top (or HEAD) commit hash in the example output above is 88f41317640a2b62c2c63ca8d755feb9f17cf16e. You can select this in your terminal and copy it to your computer's clipboard.

With your commit hash selected, you can now rollback your workspace to the previous revision. You can do this with the `git checkout` command. The following example command would rollback to the commit with hash 88f41317640a2b62c2c63ca8d755feb9f17cf16e.

```
$ git checkout 88f41317640a2b62c2c63ca8d755feb9f17cf16e .
```

Make sure that you run this command from the root of your workspace, and do not forget to include the dot at the end of the command! The dot indicates that you want to apply the changes to the entire workspace directory tree. After this has completed, you should then immediately commit with a message indicating that you performed a rollback. Push your changes and alert your team members. From there, you can start to recover from the mistake by putting your head down and getting on with your project.

18.7 Exercises

If you haven't undertaken what we've been discussing in this chapter already, you should go through everything now to ensure your system and repository is ready to go.

First, ensure that you have setup your environment correctly. Install all of the prerequisites, including Python 2.7.5 and Django 1.5.4. Django should be installed by Pip, the package manager.

Once that is complete, create a new Git repository on Github for your project. To try out the commands, you can create a new file `readme.md` in the root of your workspace. The file [will be used by GitHub](#) to provide information on your project's GitHub homepage.

- Create the file, and write some introductory text to your project.
- Add the file to the local index upon completion of writing, and commit your changes.
- Push the new file to the remote repository and observe the changes on the GitHub website.

Once you have completed these basic steps, you can then go back and edit the file some more. Add, commit and push - and then try to revert to the initial version to see if it all works as expected.

Upon completion of these exercises, all that is left for us to discuss is the environment you just setup. While all may be good just now, what if you have another Python application that requires a different version to run? This is where

the concept of [virtual environments](#) comes into play. Virtual environments allow multiple installations of Python and their relevant packages to exist in harmony, without disrupting one another. This is the generally accepted approach to configuring a Python setup nowadays. We don't go into much detail about them in this chapter now but you will be using a virtual environment when it comes to deploying your application. For now though, [check out this article](#) to read up on what they are, and how they can benefit you.

Note: There are many more advanced aspects of Git that we have not covered here, such as branching and merging. There are many fantastic tutorials available online if you are interested in taking your super-awesome version control skills a step further. For more details about such features take a look at this [tutorial on getting started with Git](#), the [Git Guide](#) or [Learning about Git Branching](#).

A CSS Crash Course

In web development, we use *Cascading Style Sheets (CSS)* to describe the presentation of a HTML document (i.e. its look and feel).

Each element within a HTML document can be *styled*. The CSS for a given HTML element describes how it is to be rendered on screen. This is done by ascribing *values* to the different *properties* associated with an element. For example, the `font-size` property could be set to `24pt` to make any text contained within the specified HTML element to appear at `24pt`. We could also set the `text-align` property to a value of `right` to make text appear within the HTML element on the right-hand side.

Note: There are many, many different CSS properties that you can use in your stylesheets. Each provides a different functionality. Check out the [W3C website](#) and [HTML Dog](#) for lists of available properties. [pageresource.com](#) also has a neat list of properties, with descriptions of what each one does. Check out Section [19.12](#) for a more comprehensive set of links.

CSS works by following a *select and apply pattern* - for a specified element, a set of styling properties are applied. Take a look at the following example in Figure 19.1 where we have some HTML containing `<h1>` tags. In the CSS code example, we specify that all `h1` are styled. We'll come back to [selectors](#) in Section [19.2](#). For now though, you can assume the CSS style defined will be applied to our `<h1>` tags. The style contains four properties:

- the first property (`font-size`) sets the size of the font to `16pt`;
- the second property (`font-style`) italicises the contents of all `<h1>` tags within the document;
- the third property (`text-align`) centres the text of the `<h1>` tags; and
- the final property (`color`) sets the colour of the text to red via [hexadecimal code #FF0000](#).

With all of these properties applied, the resultant page render can be seen in the browser in Figure 19.1.

Note: Due to the nature of web development, *what you see isn't necessarily what you'll get*. This is because different browsers have their own way of interpreting [web standards](#) and so the pages may be rendered differently. This quirk can unfortunately lead to plenty of frustration.

19.1 Including Stylesheets

Including stylesheets in your webpages is a relatively straightforward process, and involves including a `<link>` tag within your HTML's `<head>`. Check out the minimal HTML markup sample below for the attributes required within a `<link>` tag.

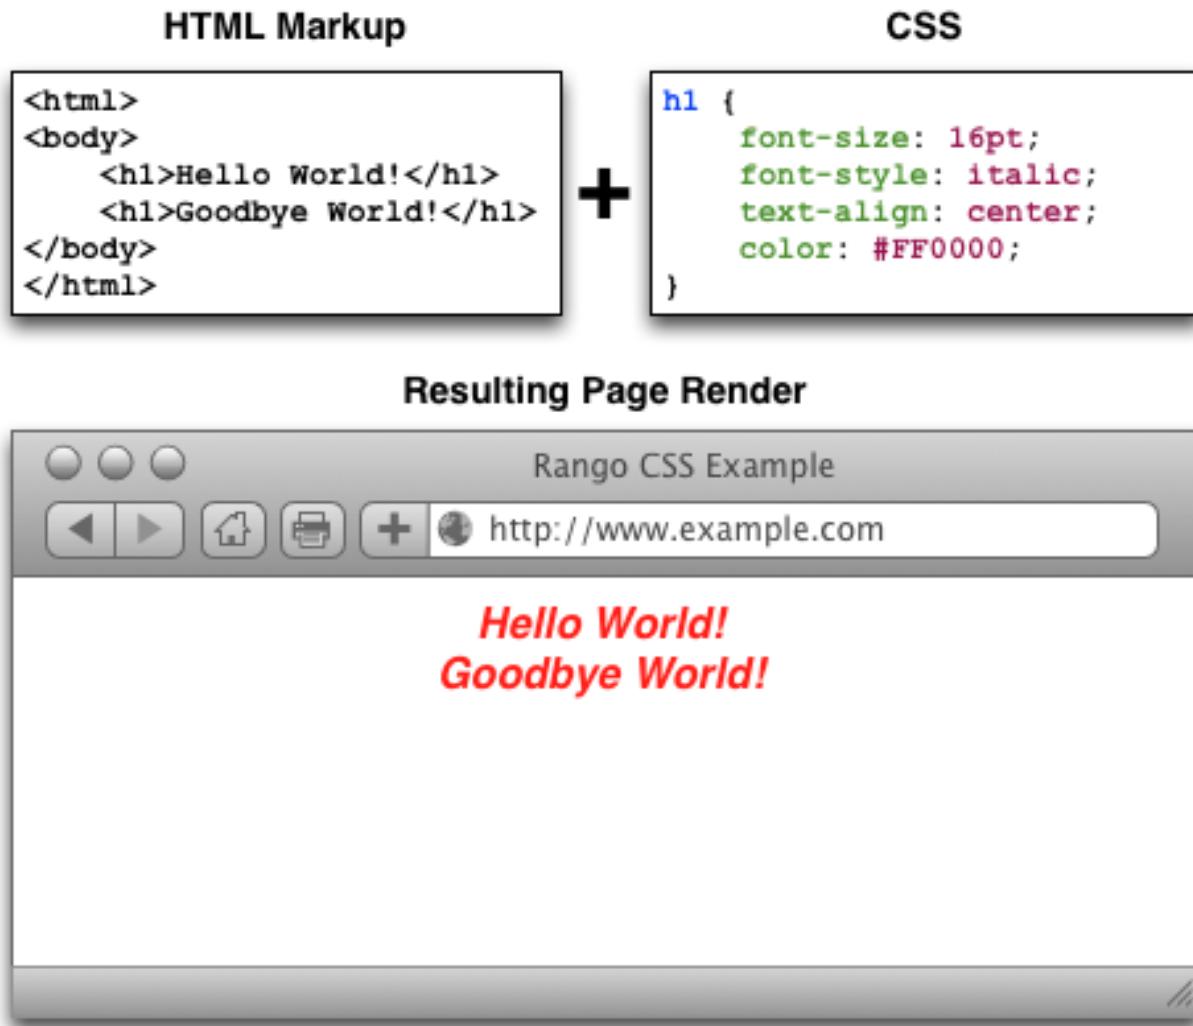


Fig. 19.1: Illustration demonstrating the rendered output of the sample HTML markup and CSS stylesheet shown. Pay particular attention to the CSS example - the colours are used to demonstrate the syntax used to define styles and the property/value pairings associated with them.

```
<!DOCTYPE html>
<html>
  <head>
    <link rel="stylesheet" type="text/css" href="URL/TO/stylesheet.css" />
    <title>Sample Title</title>
  </head>

  <body>
    <h1>Hello world!</h1>
  </body>
</html>
```

As can be seen from above, there are at minimum three attributes which you must supply to the `<link>` tag:

- `rel`, which allows you to specify the relationship between the HTML document and the resource you're linking to (i.e., a stylesheet);
- `type`, in which you should specify the **MIME type** for CSS; and
- `href`, the attribute which you should point to the URL of the stylesheet you wish to include.

With this tag added, your stylesheet should be included with your HTML page, and the styles within the stylesheet applied. It should be noted that CSS stylesheets are considered by Django as static media, meaning you should place them within your project's `static` directory.

Note: You can also add CSS to your HTML documents *inline*, meaning that the CSS is included as part of your HTML page. However, this isn't generally advised because it removes the nice abstraction between presentational semantics (CSS) and content (HTML).

19.2 Basic CSS Selectors

CSS selectors are used to map particular styles to particular HTML elements. In essence, a CSS selector is a *pattern*. Here, we cover three basic forms of CSS selector: *element selectors*, *id selectors* and *class selectors*. In Section 19.10, we also touch on what are known as *pseudo-selectors*.

19.3 Element Selectors

Taking the CSS example from Figure 19.1, we can see that the selector `h1` matches to any `<h1>` tag. Any selector referencing a tag like this can be called an *element selector*. We can apply element selectors to any HTML element such as `<body>`, `<h1>`, `<h2>`, `<h3>`, `<p>` and `<div>`. These can be all styled in a similar manner. However, using element selectors is pretty crude - styles are applied to *all* instances of a particular tag. We usually want a more fine-grained approach to selecting what elements we style, and this is where *id selectors* and *class selectors* come into play.

19.3.1 ID Selectors

The *id selector* is used to map to a unique element on your webpage. Each element on your webpage can be assigned a unique `id` via the `id` attribute, and it is this identifier that CSS uses to latch styles onto your element. This type of selector begins with a hash symbol (#), followed directly by the identifier of the element you wish to match to. Check out Figure 19.2 for an example.

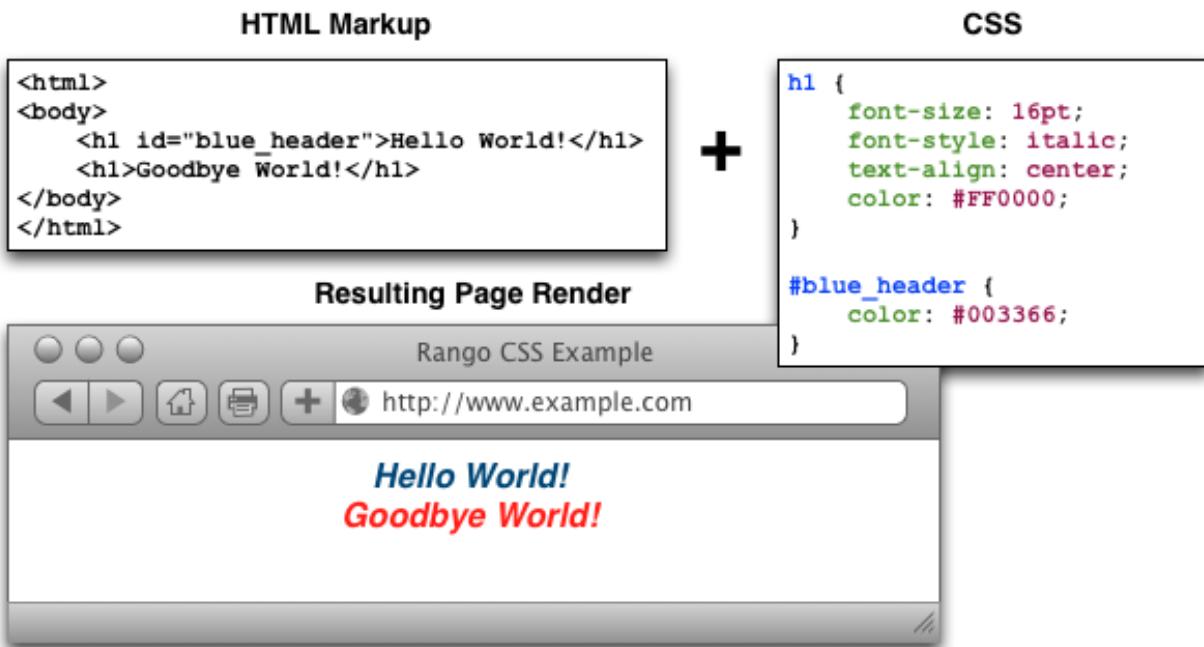


Fig. 19.2: An illustration demonstrating the use of an *id selector* in CSS. Note the blue header has an identifier which matches the CSS attribute `#blue_header`.

19.3.2 Class Selectors

The alternative option is to use *class selectors*. This approach is similar to that of *id selectors*, with the difference that you can legitimately target multiple elements with the same class. If you have a group of HTML elements that you wish to apply the same style to, use a class-based approach. The selector for using this method is to precede the name of your class with a period (.) before opening up the style with curly braces ({}). Check out Figure 19.3 for an example.

Warning: Try to use id selectors sparingly. Ask yourself: *do I absolutely need to apply an identifier to this element in order to target it?* If you need to apply it to more than one element, the answer will always be **no**. In cases like this, you should use a class or element selector.

19.4 Fonts

Due to the huge number available, using fonts has historically been a pitfall when it comes to web development. Picture this scenario: a web developer has installed and uses a particular font on his or her webpage. The font is pretty arcane - so the probability of the font being present on other computers is relatively small. A user who visits the developer's webpage subsequently sees the page rendered incorrectly as the font is not present on their system. CSS tackles this particular issue with the `font-family` property.

The value you specify for `font-family` can be a *list* of possible fonts - and the first one your computer or other device has installed is the font that is used to render the webpage. Text within the specified HTML element subsequently has the selected font applied. The example CSS shown below applies *Arial* if the font exists. If it doesn't, it looks for *Helvetica*. If that font doesn't exist, any available *sans-serif* font is applied.

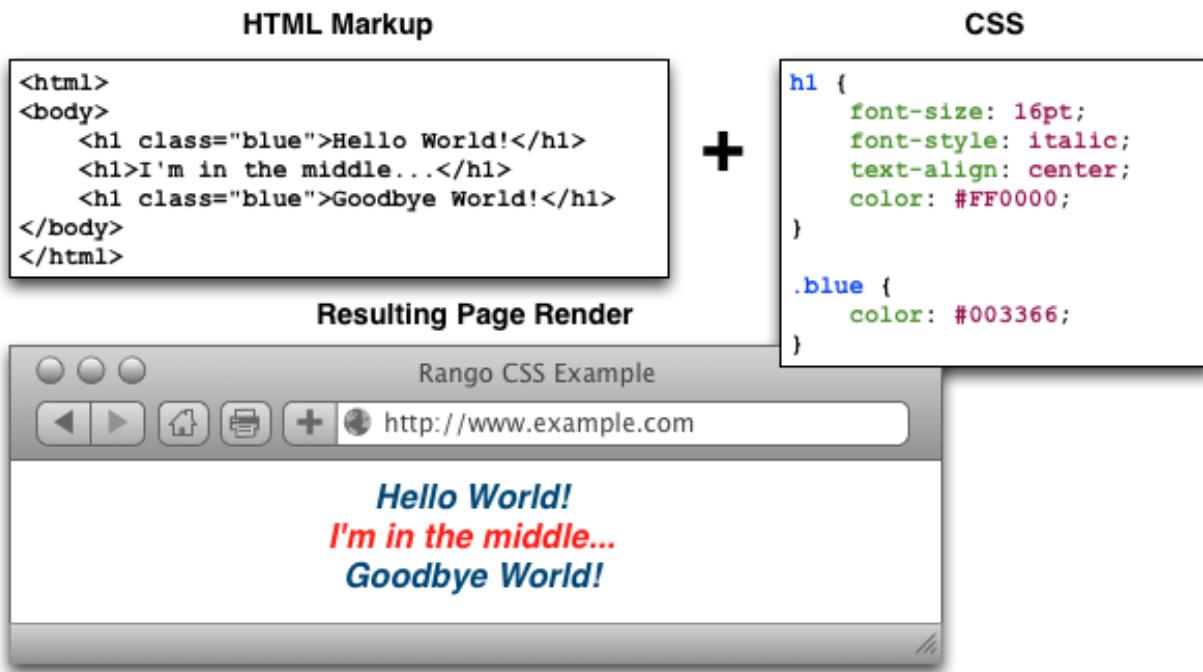


Fig. 19.3: An illustration demonstrating the use of a *class selector* in CSS. The blue headers employ the use of the `.blue` CSS style to override the red text of the `h1` style.

```
h1 {
  font-family: 'Arial', 'Helvetica', sans-serif;
}
```

In 1996, Microsoft started the [Core fonts for the Web](#) initiative with the aim of guaranteeing a particular set of fonts to be present on all computers. Today however, you can use pretty much any font you like - check out [Google Fonts](#) for examples of the typefaces you can use and this [Web Designer Depot](#) article on how to use such fonts.

19.5 Colours and Backgrounds

Colours are important in defining the look and feel of your website. You can change the colour of any element within your webpage, ranging from background colours to borders and text. In this book, we make use of words and *hexadecimal colour codes* to choose the colours we want. As you can see from the list of basic colours in Figure 19.4, you can supply either a *hexadecimal* or *RGB* (*red-green-blue*) value for the colour you want to use. You can also specify words to describe your colours, such as green, yellow or blue.

Warning: You must take great care when picking colours to use on your webpages. Don't select colours that don't contrast well - people simply won't be able to read them! There are many websites available that can help you pick out a good colour scheme - try [colorcombos.com](#) for starters.

Applying colours to your elements is a straightforward process. The property that you use depends on the aspect of the element you wish to change! The following subsections explain the relevant properties and how to apply them.

Fig. 19.4: Illustration of some basic colours with their corresponding hexadecimal and RGB values. Illustration adapted from [W3Schools](#).

There are many different websites which you can use to aid you in picking the right hexadecimal codes to enter into your stylesheets. You aren't simply limited to the nine examples above! Try out html-color-codes.com for a simple grid of colours and their associated six character hexadecimal code. You can also try sites such as color-hex.com which gives you fine-grain control over the colours you can choose.

Note: For more information on how colours are coded with hexadecimal, check out [this thorough tutorial](#).

Warning: As you may have noticed, CSS uses American/International English to spell words. As such, there are a few words which are spelt slightly differently compared to their British counterparts, like `color` and `center`. If you have grown up in Great Britain, double check your spelling and be prepared to spell it the *wrong way!* Hah!

19.5.1 Text Colours

To change the colour of text within an element, you must apply the `color` property to the element containing the text you wish to change. The following CSS for example changes all the text within an element using class `red` to...red!

```
.red {  
    color: #FF0000;  
}
```

You can alter the presentation of a small portion of text within your webpage by wrapping the text within `` tags. Assign a class or unique identifier to the element, and from there you can simply reference the `` tag in your stylesheet while applying the `color` property.

19.5.2 Borders

You can change the colour of an element's *borders*, too. We'll discuss what borders are in Section 19.8 - but for now, we'll show you how to apply colours to them to make everything look pretty.

Border colours can be specified with the `border-color` property. You can supply one colour for all four sides of your border, or specify a different colour for each side. To achieve this, you'll need to supply different colours, each separated by a space.

```
.some-element {  
    border-color: #000000 #FF0000 #00FF00  
}
```

In the example above, we use multiple colours to specify a different colour for three sides. Starting at the top, we rotate clockwise. Thus, the order of colours for each side would be top right bottom left.

Our example applies any element with class `some-element` with a black top border, a red right border and a green bottom border. No left border value is supplied, meaning that the left-hand border is left transparent. To specify a color for only one side of an element's border, consider using the `border-top-color`, `border-right-color`, `border-bottom-color` and `border-left-color` properties where appropriate.

19.5.3 Background Colours

You can also change the colour of an element's background through use of the CSS `background-color` property. Like the `color` property described above, the `background-color` property can be easily applied by specifying a single colour as its value. Check out the example below which applies a bright green background to the entire webpage. Yuck!

```
body {
    background-color: #00FF00;
}
```

19.5.4 Background Images

Of course, a colour isn't the only way to change your backgrounds. You can also apply background images to your elements, too. We can achieve this through the `background-image` property.

```
#some-unique-element {
    background-image: url('../images/filename.png');
    background-color: #000000;
}
```

The example above makes use of `filename.png` as the background image for the element with identifier `some-unique-element`. The path to your image is specified *relative to the path of your CSS stylesheet*. Our example above uses the [double dot notation to specify the relative path](#) to the image. *Don't provide an absolute path here; it won't work as you expect!* We also apply a black background colour to fill the gaps left by our background image - it may not fill the entire size of the element.

Note: By default, background images default to the top-left corner of the relevant element and are repeated on both the horizontal and vertical axes. You can customise this functionality by altering [how the image is repeated](#) with the `background-image` property. You can also specify [where the image is placed](#) by default with the `background-position` property.

19.6 Containers, Block-Level and Inline Elements

Throughout the crash course thus far, we've introduced you to the `` element but have neglected to tell you what it is. All will become clear in this section as we explain *inline* and *block-level* elements.

A `` is considered to be a so-called *container element*. Along with a `<div>` tag, these elements are themselves meaningless and are provided only for you to *contain* and *separate* your page's content in a logical manner. For example, you may use a `<div>` to contain markup related to a navigation bar, with another `<div>` to contain markup related to the footer of your webpage. As containers themselves are meaningless, styles are usually applied to help control the presentational semantics of your webpage.

Containers come in two flavours: *block-level elements* and *inline elements*. Check out Figure 19.5 for an illustration of the two kinds in action, and read on for a short description of each.

Fig. 19.5: Diagram demonstrating how block-level elements and inline elements are rendered by default. With block-level elements as green, note how a line break is taken between each element. Conversely, inline elements can appear on the same line beside each other. You can also nest block-level and inline elements within each other, but block-level elements cannot be nested within an inline element.

19.6.1 Block-Level Elements

In simple terms, a *block-level element* are by default rectangular in shape and spread across the entire width of the containing element. Block-level elements therefore by default appear underneath each other. The rectangular structure of each block-level element is commonly referred to as the *box model*, which we discuss in Section 19.8. A typical block-level element you will use is the `<div>` tag, short for *division*.

Block-level elements can be nested within other block-level elements to create a hierarchy of elements. You can also nest *inline elements* within block-level elements, but not vice-versa! Read on to find out why.

19.6.2 Inline Elements

An *inline element* does exactly what it says on the tin. These elements appear *inline* to block-level elements on your webpage, and are commonly found to be wrapped around text. You'll find that `` tags are commonly used for this purpose.

This text-wrapping application was explained in Section 19.5.1, where a portion of text could be wrapped in `` tags to change its colour. The corresponding HTML markup would look similar to the example below.

```
<div>
```

```
    This is some text wrapped within a block-level element. <span class="red">This text is wrapped w
```

```
</div>
```

Refer back to Figure 19.5 to refresh your mind about what you can and cannot nest before you move on.

19.7 Basic Positioning

An important concept that we have not yet covered in this CSS crash course regards the positioning of elements within your webpage. Most of the time, you'll be satisfied with inline elements appearing alongside each other, and block-level elements appearing underneath each other. These elements are said to be *positioned statically*.

However, there will be scenarios where you require a little bit more control on where everything goes. In this section, we'll briefly cover three important techniques for positioning elements within your webpage: *floats*, *relative positioning* and *absolute positioning*.

19.7.1 Floats

CSS *floats* are one of the most straightforward techniques for positioning elements within your webpage. Using floats allows us to position elements to the left or right of a particular container - or page.

Let's work through an example. Consider the following HTML markup and CSS code.

```
<div class="container">
    <span class="yellow">Span 1</span>
    <span class="blue">Span 2</span>
</div>
```

```
.container {
    border: 1px solid black;
}

.yellow {
    background-color: yellow;
    border: 1px solid black;
}

.blue {
    background-color: blue;
    border: 1px solid black;
}
```

This produces the output shown below.

We can see that each element follows its natural flow: the container element with class `container` spans the entire width of its parent container, while each of the `` elements are enclosed inline within the parent. Now suppose that we wish to then move the blue element with text `Span 2` to the right of its container. We can achieve this by modifying our CSS `.blue` class to look like the following example.

```
.blue {
    background-color: blue;
    border: 1px solid black;
    float: right;
}
```

By applying the `float: right;` property and value pairing, we should then see something similar to the example shown below.

Note how the `.blue` element now appears at the right of its parent container, `.container`. We have in effect disturbed the natural flow of our webpage by artificially moving an element! What if we then also applied `float: left` to the `.yellow` ``?

This would float the `.yellow` element, removing it from the natural flow of the webpage. In effect, it is not sitting on top of the `.container` container. This explains why the parent container does not now fill down with the `` elements like you would expect. You can apply the `overflow: hidden;` property to the parent container as shown below to fix this problem. For more information on how this trick works, have a look at [this QuirksMode.org online article](#).

```
.container {
    border: 1px solid black;
    overflow: hidden;
}
```

Applying `overflow: hidden` ensures that that our `.container` pushes down to the appropriate height.

19.7.2 Relative Positioning

Relative positioning can be used if you required a greater degree of control over where elements are positioned on your webpage. As the name may suggest to you, relative positioning allows you to position an element *relative to where it would otherwise be located*. We make use of relative positioning with the `position: relative;` property and value pairing. However, that's only part of the story.

Let's explain how this works. Consider our previous example where two `` elements are sitting within their container.

```
<div class="container">
    <span class="yellow">Span 1</span>
    <span class="blue">Span 2</span>
</div>
```

```
.container {
    border: 1px solid black;
    height: 200px;
}

.yellow {
    background-color: yellow;
    border: 1px solid black;
}
```

```
.blue {
    background-color: blue;
    border: 1px solid black;
}
```

This produces the following result - just as we would expect. Note that we have artificially increased the height of our container element to 150 pixels. This will allow us more room with which to play with.

Now let's attempt to position our `.blue ` element relatively. First, we apply the `position: relative;` property and value pairing to our `.blue` class, like so.

```
.blue {
    background-color: blue;
    border: 1px solid black;
    position: relative;
}
```

This has no effect on the positioning of our `.blue` element. What it does do however is change the positioning of `.blue` from `static` to `relative`. This paves the way for us to specify where - from the original position of our element - we now wish the element to be located at.

```
.blue {
    background-color: blue;
    border: 1px solid black;
    position: relative;
    left: 150px;
    top: 80px;
}
```

By applying the `left` and `top` properties as shown in the example above, we are wanting the `.blue` element to be *pushed* 150 pixels *from the left*. In other words, we move the element 150 pixels to the right. Think about that carefully! The `top` property indicates that the element should be pushed 80 pixels from the *top* of the element. The result our experimentation can be seen below.

From this behaviour, we can deduce that the properties `right` and `bottom` *push* elements from the right and bottom respectively. We can test this out by applying the properties to our `.yellow` class as shown below.

```
.yellow {
    background-color: blue;
    border: 1px solid black;
    float: right;
    position: relative;
    right: 10px;
    bottom: 10px;
}
```

This produces the following output. The `.yellow` container is pushed into the top left-hand corner of our container by pushing up and to the right.

Note: What happens if you apply both a `top` and `bottom` property, or a `left` and `right` property? Interestingly, the *first* property for the relevant axis is applied. For example, if `bottom` is specified before `top`, the `bottom` property is used.

We can even apply relative positioning to elements which are floated. Consider our earlier example where the two `` elements were positioned on either side of the container by floating `.blue` to the right.

We can then alter the `.blue` class to the following.

```
.blue {
    background-color: blue;
    border: 1px solid black;
    float: right;
    position: relative;
    right: 100px;
}
```

This therefore means that relative positioning works from the position at which the element would have otherwise been at - regardless of any other position-changing properties being applied. Neat!

19.7.3 Absolute Positioning

Our final positioning technique is *absolute positioning*. While we still modify the `position` parameter of a style, we use `absolute` as the value instead of `relative`. In contrast to relative positioning, absolute positioning places an element *relative to its first parent element that has a position value other than static*. This may sound a little bit confusing, but let's go through it step by step to figure out what exactly happens.

First, we can again take our earlier example of the two coloured `` elements within a `<div>` container. The two `` elements are placed side-by-side as they would naturally.

```
<div class="container">
    <span class="yellow">Span 1</span>
    <span class="blue">Span 2</span>
</div>
```

```
.container {
    border: 1px solid black;
    height: 70px;
}

.yellow {
    background-color: yellow;
    border: 1px solid black;
}

.blue {
    background-color: blue;
    border: 1px solid black;
}
```

This produces the output shown below. Note that we again set our `.container` height to an artificial value of 70 pixels to give us more room.

We now apply absolute positioning to our `.blue` element.

```
.blue {
    background-color: blue;
    border: 1px solid black;
    position: absolute;
}
```

Like with relative positioning, this has no overall effect on the positioning of our blue element in the webpage. We must apply one or more of `top`, `bottom`, `left` or `right` in order for a new position to take effect. As a demonstration, we can apply `top` and `left` properties to our blue element like in the example below.

```
.blue {
    background-color: blue;
```

```
border: 1px solid black;
position: absolute;
top: 0;
left: 0;
}
```

Wow, what happened here? Our blue element is now positioned outside of our container! You'll note that if you run this code within your own web browser window, the blue element appears in the top left-hand corner of the viewport. This therefore means that our `top`, `bottom`, `left` and `right` properties take on a slightly different meaning when absolute positioning is concerned.

As our container element's position is by default set to `position: static`, the blue and yellow elements are moving to the top left and bottom right of our screen respectively. Let's now modify our `.yellow` class to move the yellow `` to 5 pixels from the bottom right-hand corner of our page. The `.yellow` class now looks like the example below.

```
.yellow {
background-color: yellow;
border: 1px solid black;
position: absolute;
bottom: 5px;
right: 5px;
}
```

This produces the following result.

But what if we don't want our elements to be positioned absolutely in relation to the entire page? More often than not, we'll be looking to adjusting the positioning of our elements in relation to a container. If we recall our definition for absolute positioning, we will note that absolute positions are calculated *relative to the first parent element that has a position value other than static*. As our container is the only parent for our two `` elements, the container to which the absolutely positioned elements is therefore the `<body>` of our HTML page. We can fix this by adding `position: relative;` to our `.container` class, just like in the example below.

```
.container {
border: 1px solid black;
height: 70px;
position: relative;
}
```

This produces the following result. `.container` becomes the first parent element with a position value of anything other than `relative`, meaning our `` elements latch on!

Our elements are now absolutely positioned in relation to `.container`. Awesome! Let's adjust the positioning values of our two `` elements to move them around.

```
.yellow {
background-color: yellow;
border: 1px solid black;
position: absolute;
top: 20px;
right: 100px;
}

.blue {
background-color: blue;
border: 1px solid black;
position: absolute;
float: right;
bottom: 50px;
}
```

```
left: 40px;
}
```

Note that we also apply `float: right;` to our `.blue` element. This is to demonstrate that unlike relative positioning, absolute positioning *ignores any other positioning properties applied to an element.* `top: 10px` for example will always ensure that an element appears 10 pixels down from its parent (set with `position: relative;`), regardless of whether the element has been floated or not.

19.8 The Box Model

When using CSS, you're never too far away from using *padding*, *borders* and *margins*. These properties are some of the most fundamental styling techniques which you can apply to the elements within your webpages. They are incredibly important and are all related to what we call the *CSS box model*.

Each element that you create on a webpage can be considered as a box. The [CSS box model](#) is defined by the [W3C](#) as a formal means of describing the elements or boxes that you create, and how they are rendered in your web browser's viewport. Each element or box consists of *four separate areas*, all of which are illustrated in Figure 19.6. The areas - listed from inside to outside - are the *content area*, the *padding area*, the *border area* and the *margin area*.

Fig. 19.6: An illustration demonstrating the CSS box model, complete with key showing the four areas of the model.

For each element within a webpage, you can create a margin, apply some padding or a border with the respective properties `margin`, `padding` and `border`. Margins clear a transparent area around the border of your element, meaning margins are incredibly useful for creating a gap between elements. In contrast, padding creates a gap between the content of an element and its border. This therefore gives the impression that the element appears wider. If you supply a background colour for an element, the background colour is extended with the element's padding. Finally, borders are what you might expect them to be - they provide a border around your element's content and padding.

For more information on the CSS box model, check out [addedbytes](#) excellent explanation of the model. Heck, [why not even order a t-shirt with the box model on it?](#)

Warning: As you may gather from examining Figure 19.6, the width of an element isn't defined simply by the value you enter as the element's `width`. Rather, you should always consider the width of the border and padding on both sides of your element. This can be represented mathematically as:

$$\text{total_width} = \text{content_width} + \text{left padding} + \text{right padding} + \text{left border} + \text{left margin} + \text{right margin}$$

Don't forget this. You'll save yourself a lot of trouble if you don't!

19.9 Styling Lists

Lists are everywhere. Whether you're reading a list of learning outcomes for a course or a reading a list of times for the train, you know what a list looks like and appreciate its simplicity. If you have a list of items on a webpage, why not use a HTML list? Using lists within your webpages - according to [Brainstorm and Raves](#) - promotes good HTML document structure, allowing text-based browsers, screen readers and other browsers that do not support CSS to render your page in a sensible manner.

Lists however don't look particularly appealing to end-users. Take the following HTML list that we'll be styling as we go along trying out different things.

```
<ul class="sample-list">
  <li>Django</li>
  <li>How to Tango with Django</li>
  <li>Two Scoops of Django</li>
</ul>
```

Rendered without styling, the list looks pretty boring.

Let's make some modifications. First, let's get rid of the ugly bullet points. With our `` element already (and conveniently) set with class `sample-list`, we can create the following style.

```
.sample-list {
  list-style-type: none;
}
```

This produces the following result. Note the now lacking bullet points!

Let's now change the orientation of our list. We can do this by altering the `display` property of each of our list's elements (``). The following style maps to this for us.

```
.sample-list li {
  display: inline;
}
```

When applied, our list elements now appear on a single line, just like in the example below.

While we may have the correct orientation, our list now looks awful. Where does one element start and the other end? It's a complete mess! Let's adjust our list element style and add some contrast and padding to make things look nicer.

```
.example-list li {
  display: inline;
  background-color: #333333;
  color: #FFFFFF;
  padding: 10px;
}
```

When applied, our list looks so much better - and quite professional, too!

From the example, it is hopefully clear that lists can be easily customised to suit the requirements of your webpages. For more information and inspiration on how to style lists, you can check out some of the selected links below.

- Have a look at [this excellent tutorial on styling lists on A List Apart](#).
- Have a look at [this about.com article which demonstrates how to use your own bullets!](#)
- Check out [this advanced tutorial from Web Designer Wall](#) which uses graphics to make awesome looking lists. In the tutorial, the author uses Photoshop - you could try using a simpler graphics package if you don't feel confident with Photoshop.
- [This awesome site compilation from devsnippets.com](#) provides some great inspiration and tips on how you can style lists.

The possibilities of styling lists is endless! You could say it's a never-ending list...

19.10 Styling Links

CSS provides you with the ability to easily style hyperlinks in any way you wish. You can change their colour, their font or any other aspect that you wish - and you can even change how they look when you hover over them!

Hyperlinks are represented within a HTML page through the `<a>` tag, which is short for *anchor*. We can apply styling to all hyperlinks within your webpage as shown in following example.

```
a {
    color: red;
    text-decoration: none;
}
```

Every hyperlink's text colour is changed to red, with the default underline of the text removed. If we then want to change the `color` and `text-decoration` properties again when a user hovers over a link, we can create another style using the so-called **pseudo-selector** `:hover`. Our two styles now look like the example below.

```
a {
    color: red;
    text-decoration: none;
}

a:hover {
    color: blue;
    text-decoration: underline;
}
```

This produces links as shown below. Hover over them to see them change!

You may not however wish for the same link styles across the entire webpage. For example, your navigation bar may have a dark background while the rest of your page has a light background. This would necessitate having different link stylings for the two areas of your webpage. The example below demonstrates how you can apply different link styles by using a slightly more complex CSS style selector.

```
#dark {
    background-color: black;
}

#dark a {
    color: white;
    text-decoration: underline;
}

#dark a:hover {
    color: aqua;
}

.light {
    background-color: white;
}

.light a {
    color: black;
    text-decoration: none;
}

.light a:hover {
    color: olive;
    text-decoration: underline;
}
```

We can then construct some simple markup to demonstrate these classes.

```
<div id="dark">
  <a href="http://www.google.co.uk/">Google Search</a>
</div>

<div class="light">
  <a href="http://www.bing.co.uk/">Bing Search</a>
</div>
```

The resultant output looks similar to the example shown below. Again, hover over the links to see them change!

With a small amount of CSS, you can make some big changes in the way your webpages appear to end users.

19.11 The Cascade

It's worth pointing out where the *Cascading* in *Cascading Style Sheets* comes into play. You may have noticed in the example rendered output in Figure 19.1 that the red text is **bold**, yet no such property is defined in our `h1` style. This is a perfect example of what we mean by *cascading styles*. Most HTML elements have associated with them a *default style* which web browsers apply. For `<h1>` elements, the [W3C website provides a typical style that is applied](#). If you check the typical style, you'll notice that it contains a `font-weight: bold;` property and value pairing, explaining where the **bold** text comes from. As we define a further style for `<h1>` elements, typical property/value pairings *cascade* down into our style. If we define a new value for an existing property/value pairing (such as we do for `font-size`), we *override* the existing value. This process can be repeated many times - and the property/value pairings at the end of the process are applied to the relevant element. Check out 19.7 for a graphical representation of the cascading process.

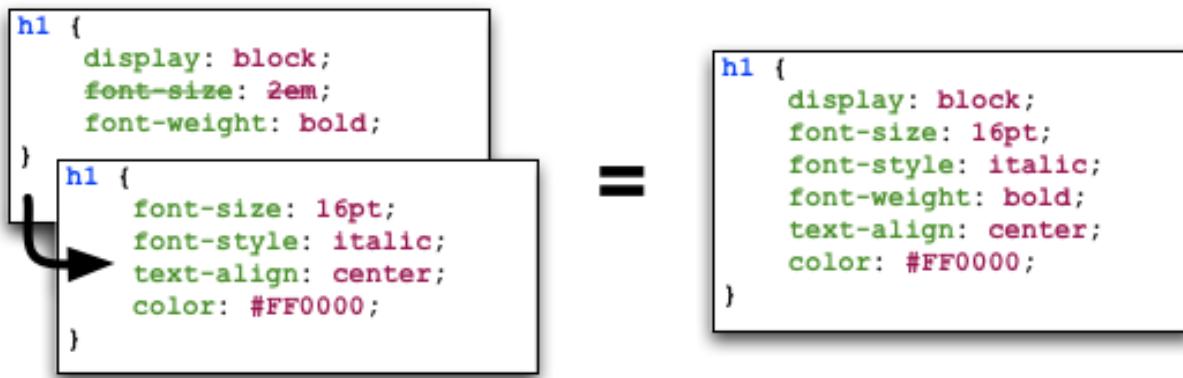


Fig. 19.7: Illustration demonstrating the *cascading* in *Cascading Style Sheets* at work. Take note of the `font-size` property in our `h1` style - it is overridden from the default value. The cascading styles produce the resultant style, shown on the right of the illustration.

19.12 Additional Reading

What we've discussed in this section is by no means a definitive guide to CSS. There are [300-page books](#) devoted to CSS alone! What we have provided you with here is a very brief introduction showing you the very basics of what CSS is and how you can use it.

As you develop your web applications, you'll undoubtedly run into issues and frustrating problems with styling web content. This is part of the learning experience, and you still have a bit to learn. We strongly recommend that you

invest some time trying out several online tutorials about CSS - there isn't really any need to buy a book (unless you want to).

- The [W3C provides a neat tutorial on CSS](#), taking you by the hand and guiding you through the different stages required. They also introduce you to several new HTML elements along the way, and show you how to style them accordingly.
- [W3Schools also provides some cool CSS tutorials](#). Instead of guiding you through the process of creating a webpage with CSS, *W3Schools* has a series of mini-tutorials and code examples to show you how to achieve a particular feature, such as setting a background image. We highly recommend that you have a look here.
- [html.net has a series of lessons on CSS](#) which you can work through. Like W3Schools, the tutorials on *html.net* are split into different parts, allowing you to jump into a particular part you may be stuck with.
- It's also worth having a look at [CSSeasy.com](#)'s collection of tutorials, providing you with the basics on how to develop different kinds of page layouts.

This list is by no means exhaustive, and a quick web search will indeed yield much more about CSS for you to chew on. Just remember: CSS can be tricky to learn, and there may be times where you feel you want to throw your computer through the window. We say this is pretty normal - but take a break if you get to that stage. We'll be tackling some more advanced CSS stuff as we progress through the tutorial in the next few sections.

Note: With an increasing array of devices equipped with more and more powerful processors, we can make our web-based content do more. To keep up, [CSS has constantly evolved](#) to provide new and intuitive ways to express the presentational semantics of our SGML-based markup. To this end, support for relatively new CSS properties may be limited on several browsers, which can be a source of frustration. The only way to reliably ensure that your website works across a wide range of different browsers and platforms is to [test, test and test some more!](#)

Javascript and JQuery

In the next edition we'll be adding in tutorials on how to use Javascript and JQuery within Web Application.

Test Driven Development

In the next edition, we'll be adding in a chapter about Test Driven Development - in the meantime we highly recommend checking out the [tutorial on test driven development by Harry Percival](#).

21.1 Exercises

- Undertake Part Five of the official Django Tutorial to learn about automated testing.

Acknowledgements

The first version of this book began because my students were struggling to pick up the Django Framework as part of the course I teach on Web Application Development. However, it became painfully clear that the official Django Tutorials made many assumptions regarding one's background, missed out crucial steps here and there that would trip up the students, and finally, it was clear that the students lacked the vocabulary to google for the right things to resolve any problems that did arise (and also that they didn't have hours to spend trying to figure out what was going on).

I realised that to go from woe to go required a lot of additional information and help to resolve common problems. For this I often consulted stackoverflow.com for advice, help and explanations, along with other sites - the folk who have contributed to these sites have been an invaluable source of knowledge which I have tried to distill and combine here. Thanks for your help it is very much appreciated.

Leif

In this book, we have pulled together the relevant information and details so that students and beginners to web development can go through the process in a seamless fashion and learn the basics of web application development. After realising the book we have had numerous people contact us with ways to improve the quality of the book. So we'd like to thank:

- Manoel Maria @xmadruga157
- Codenius @Codenius
- Ally Weir
- Can Ibanoglu @canlbanoglu
- Sadegh Kh @IR_Melbourne
- Jan Felix Trettow @JanFelixTrettow
- Rezha Julio @kimiamania
- Thomas Murphy @TMurphyMusic
- Saurabh Tandon @saurabhtand
- Pierre-Yves Mathieu <https://github.com/pywebdesign>
- Dhiraj Thakur <https://github.com/dhirajt>
- Svante Kvarnström @sjktje
- Ramdog <https://github.com/ramdog>
- Krace Kumar <https://github.com/kracekumar>
- James Yeo
- Michael Ho Chum <https://github.com/michaelchum>

- Jonathan Sundqvist <https://github.com/jonathan-s>
- nolan-m <https://github.com/nolan-m>
- Kartik Singhal <https://github.com/k4rtik>
- William Vincent
- Mike Gleen

for their suggestions, comments and feedback, which have greatly improved the quality of this resource.

Leif and David

Indices and tables

- genindex
- modindex
- search